Applied Supervised Learning for Cyber Security Module 0: Introduction

Presented by Brian Genz, CISSP, GREM, GCIH, GNFA, GCFA
brian.genz@gtkcyber.com

Course Agenda

Day 1: Introduction & Data Engineering

- Intro: What is Data Science?
- What is Machine Learning?
    - Overview of Machine Learning & Cyber Applications
- Overview of Data Preparation and Feature Engineering
- Data Visualization

Day 2: Supervised Learning

- Overview of Supervised Learning techniques

Day 3: Machine Learning

- Feature Engineering
- Supervised Machine Learning
- Unsupervised Machine Learning
- Hacking Machine Learning Models

Day 4: Advanced Topics

- Introduction to Big Data Tools
- Anomaly Detection
    - PySpark, ELK & Kafka
- Hunting with Data Science

Expectations

- Please participate and **ask questions.**
- Please follow along and **TRY OUT** the examples yourself during the class
- All the answers are in the slide decks or GitHub repository, but please try to complete the exercises **without looking at the answers**.
- Join the conversation in slack!
- Have fun!

Introduction

Our Lawyers Make Us Say This

Brian Genz

Brian leads the Threat & Vulnerability Management team at Splunk and serves as a Senior Instructor with GTK Cyber.

Prior to leading the red team at Splunk, he led threat hunting and security orchestration, automation and response (SOAR) efforts with Splunk Phantom at a Fortune 100 financial company. He has experience in the defense intelligence, manufacturing, and financial sectors and has worked in both offensive and defensive security. Degrees and certifications include MBA, M.S. in Information Technology Management, CISSP, GREM, GNFA, GCFA, and GCIH.

Brian has taught with GTK Cyber at Black Hat USA and in Europe. He has presented at multiple industry conferences, including the DEFCON AI Village, the Conference on Applied Machine Learning for Information Security (CAMLIS), Derby Con, Circle City Con, Cyphercon, the ISSA International Conference, ISACA Milwaukee, Infragard, and others. He has served as adjunct faculty at Wayne State University and Milwaukee Area Technical College. He is also featured in the book, "Tribe of Hackers: Red Team Edition."

Who are you?

- Your name (or what you want us to call you)
- Your job role
- What you hope to get out of this class
- Your level of experience with coding

What is Data Science?

Data Science is the automated extraction of information from raw data.

Data Science is the art of turning data into actions. This is accomplished through the creation of data products, which provide actionable information without exposing decision makers to the underlying data or analytics

Booz Allen Hamilton, Field Guide to Data Science, Pg. 17

Analyst  Developer

Analyst ✚ Developer

What Data Science is Not

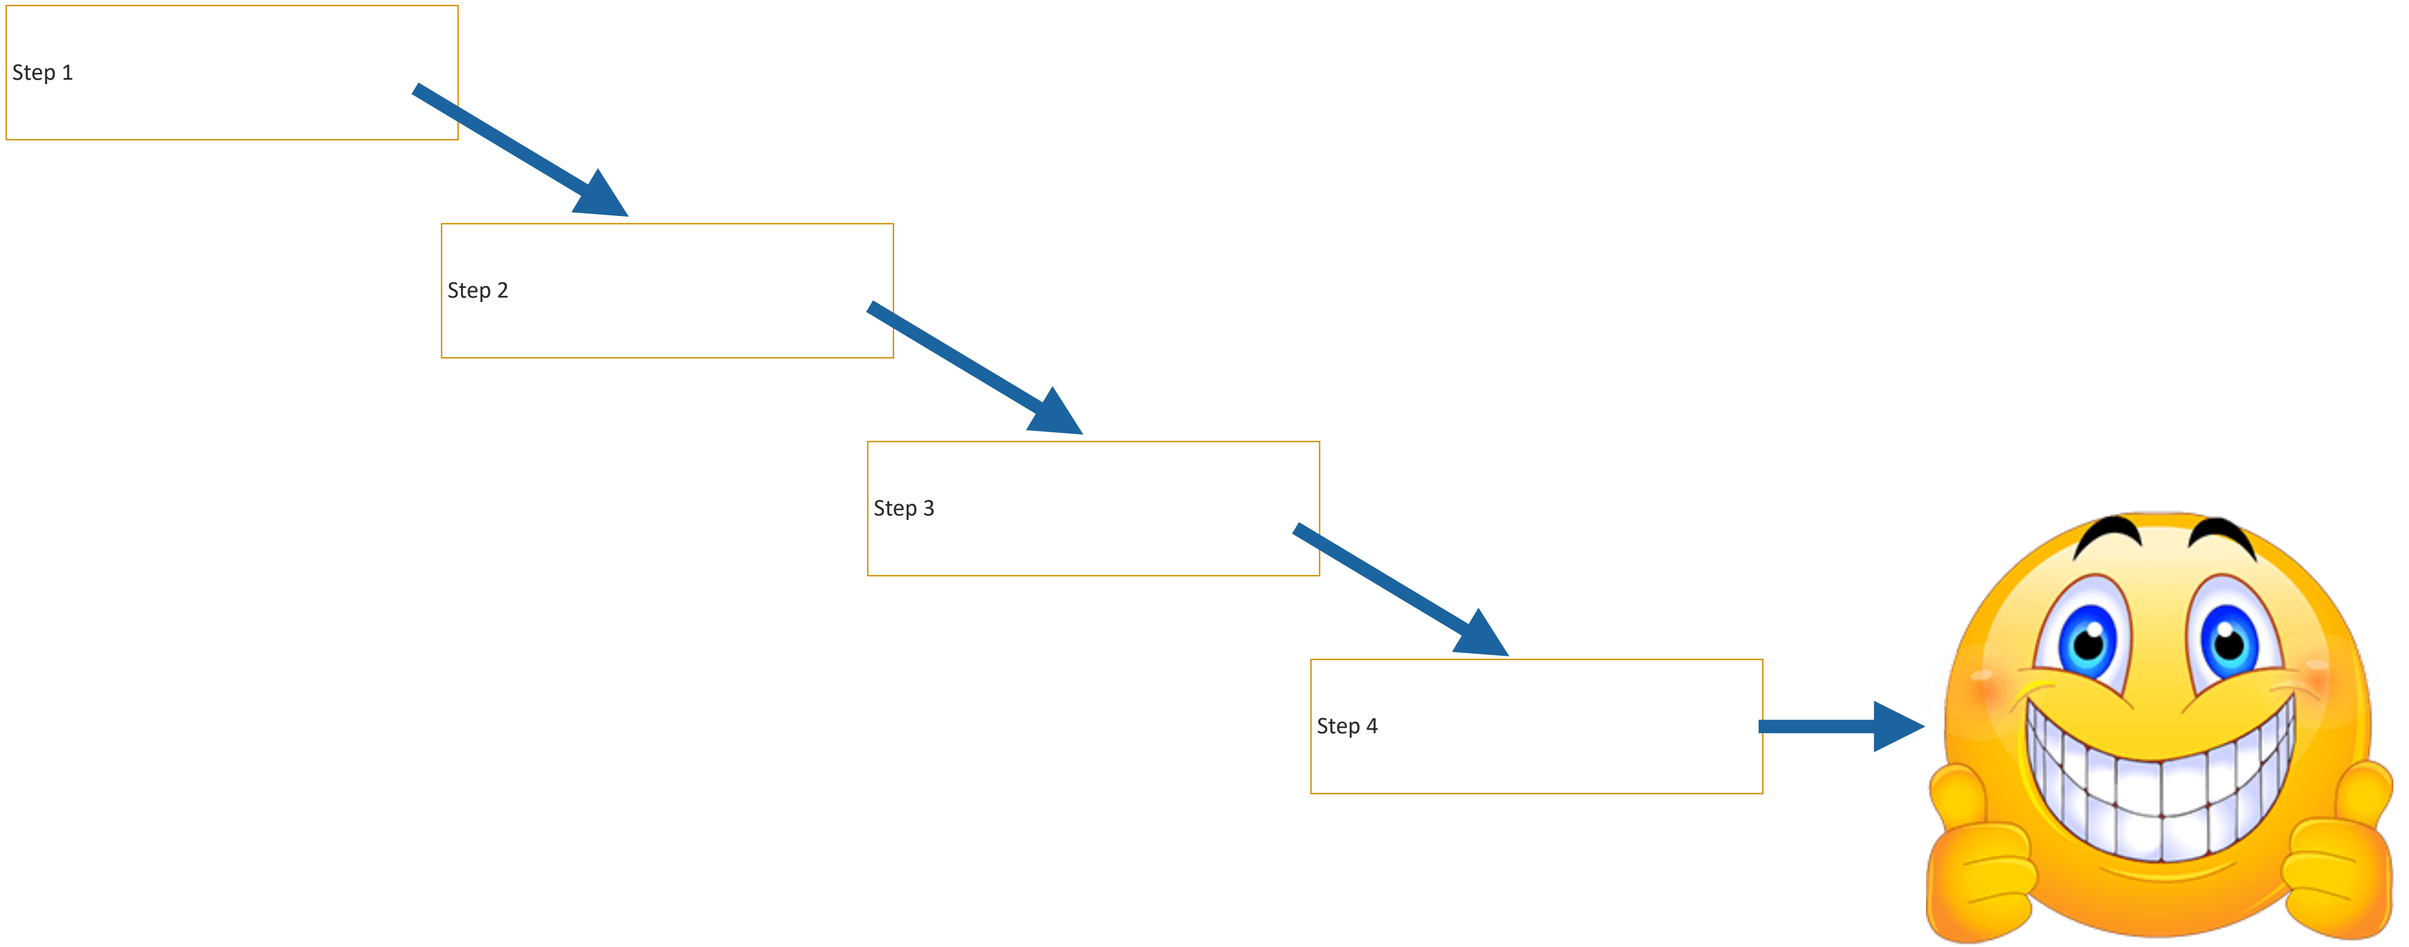

What Data Science Is

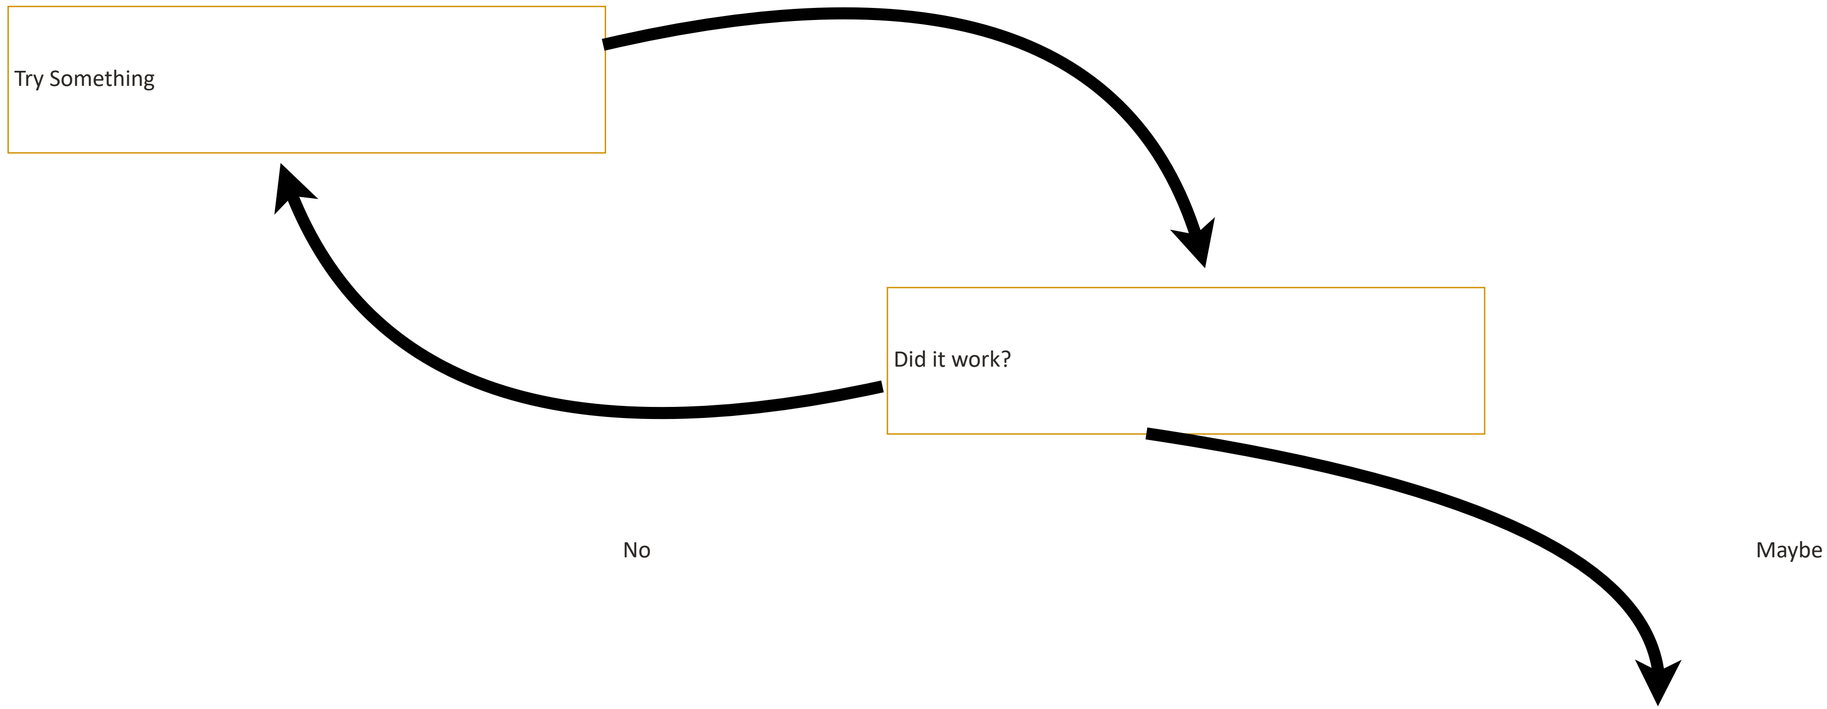

Research Process

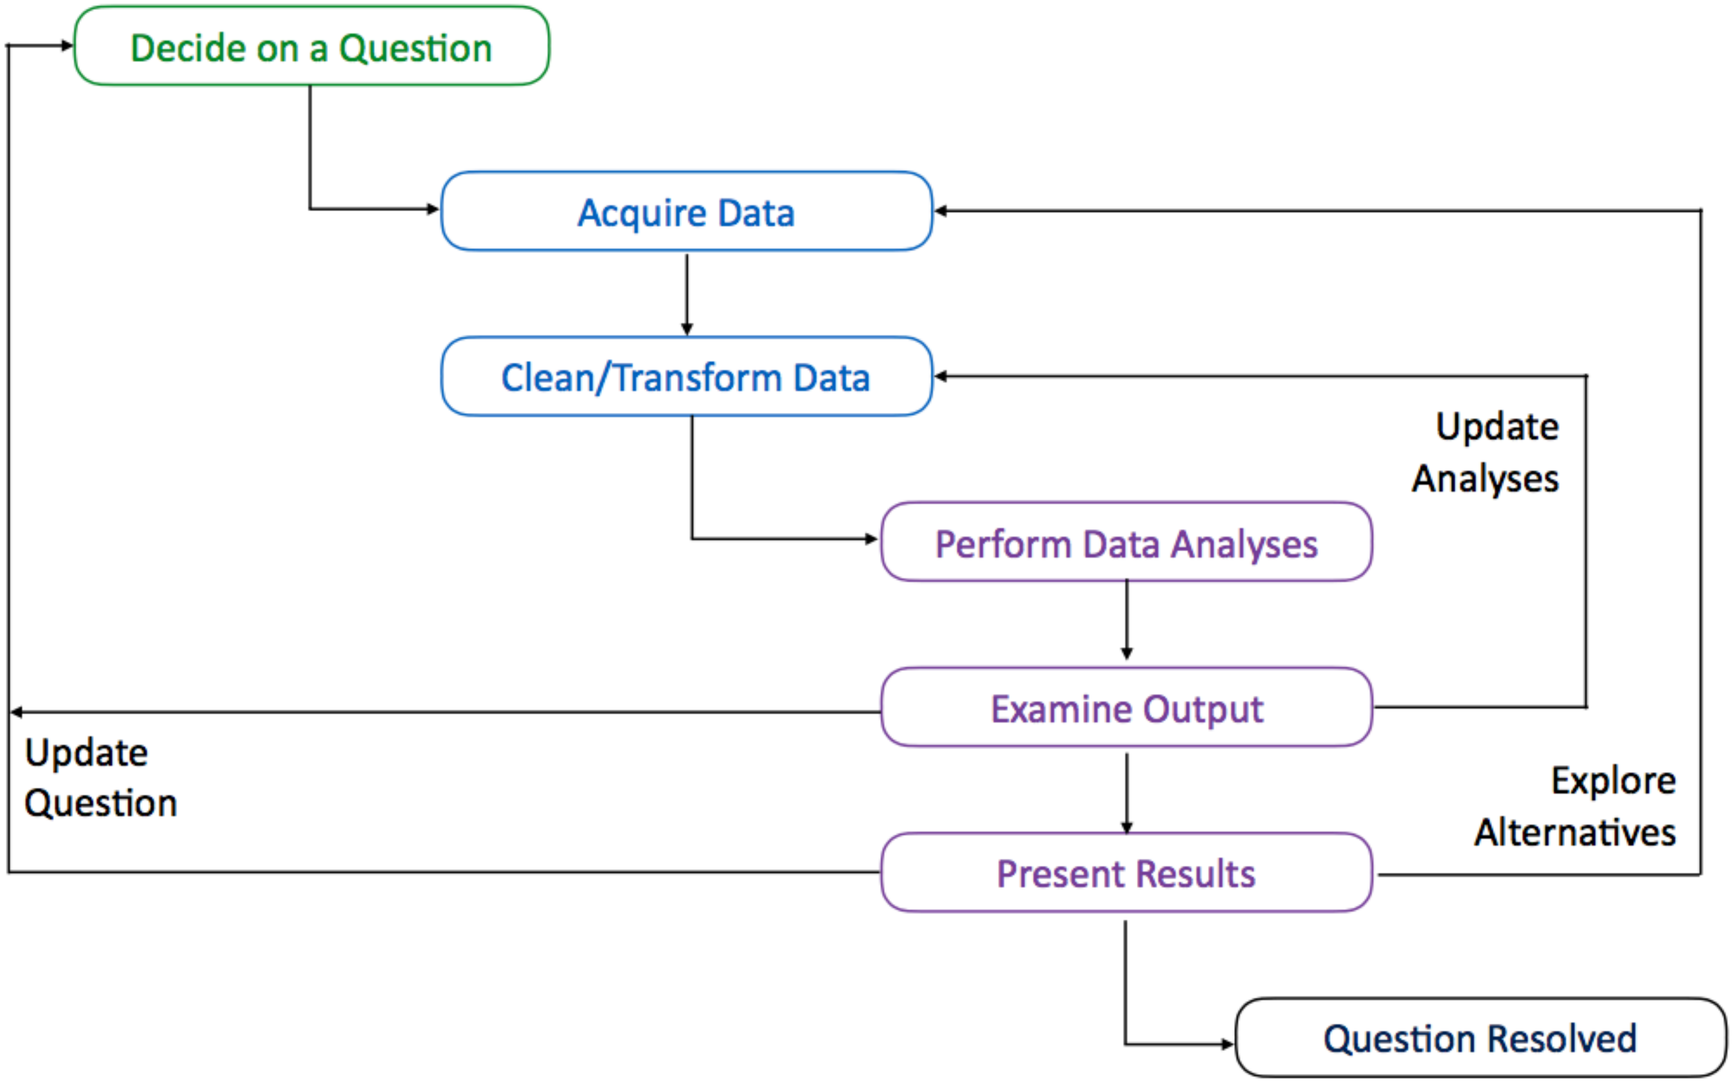

Research Process

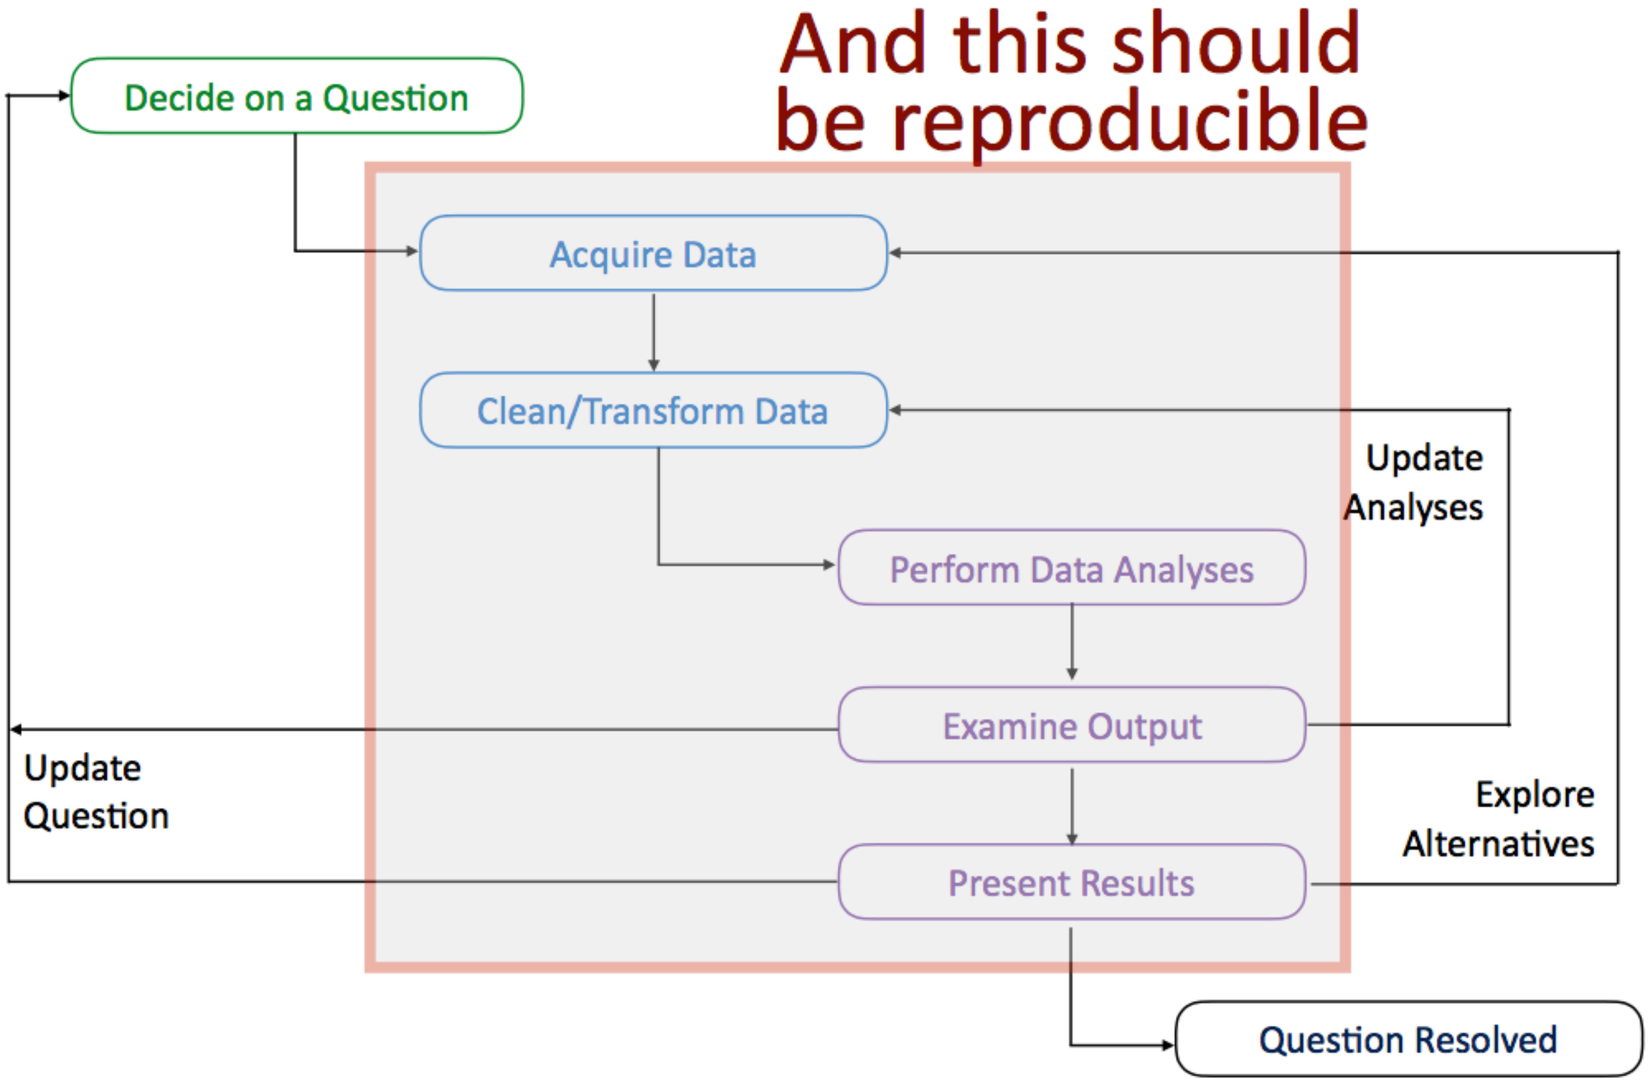

“The term "data scientist" will subside and may well sound dated five years from now. **The skills will become more commonplace and commoditized. When that happens, the real boom will begin**, because the technology will become widely adopted and thus more useful. … **Instead, we need self-service tools that empower smart and tenacious business people to perform Big Data analysis themselves.**

–Andrew Brust, “Data scientists don't scale”, http://www.zdnet.com/article/data-scientists-dont-scale/

Time to Insight

Time to Insight

Time = $$

Value of Insights Over Time

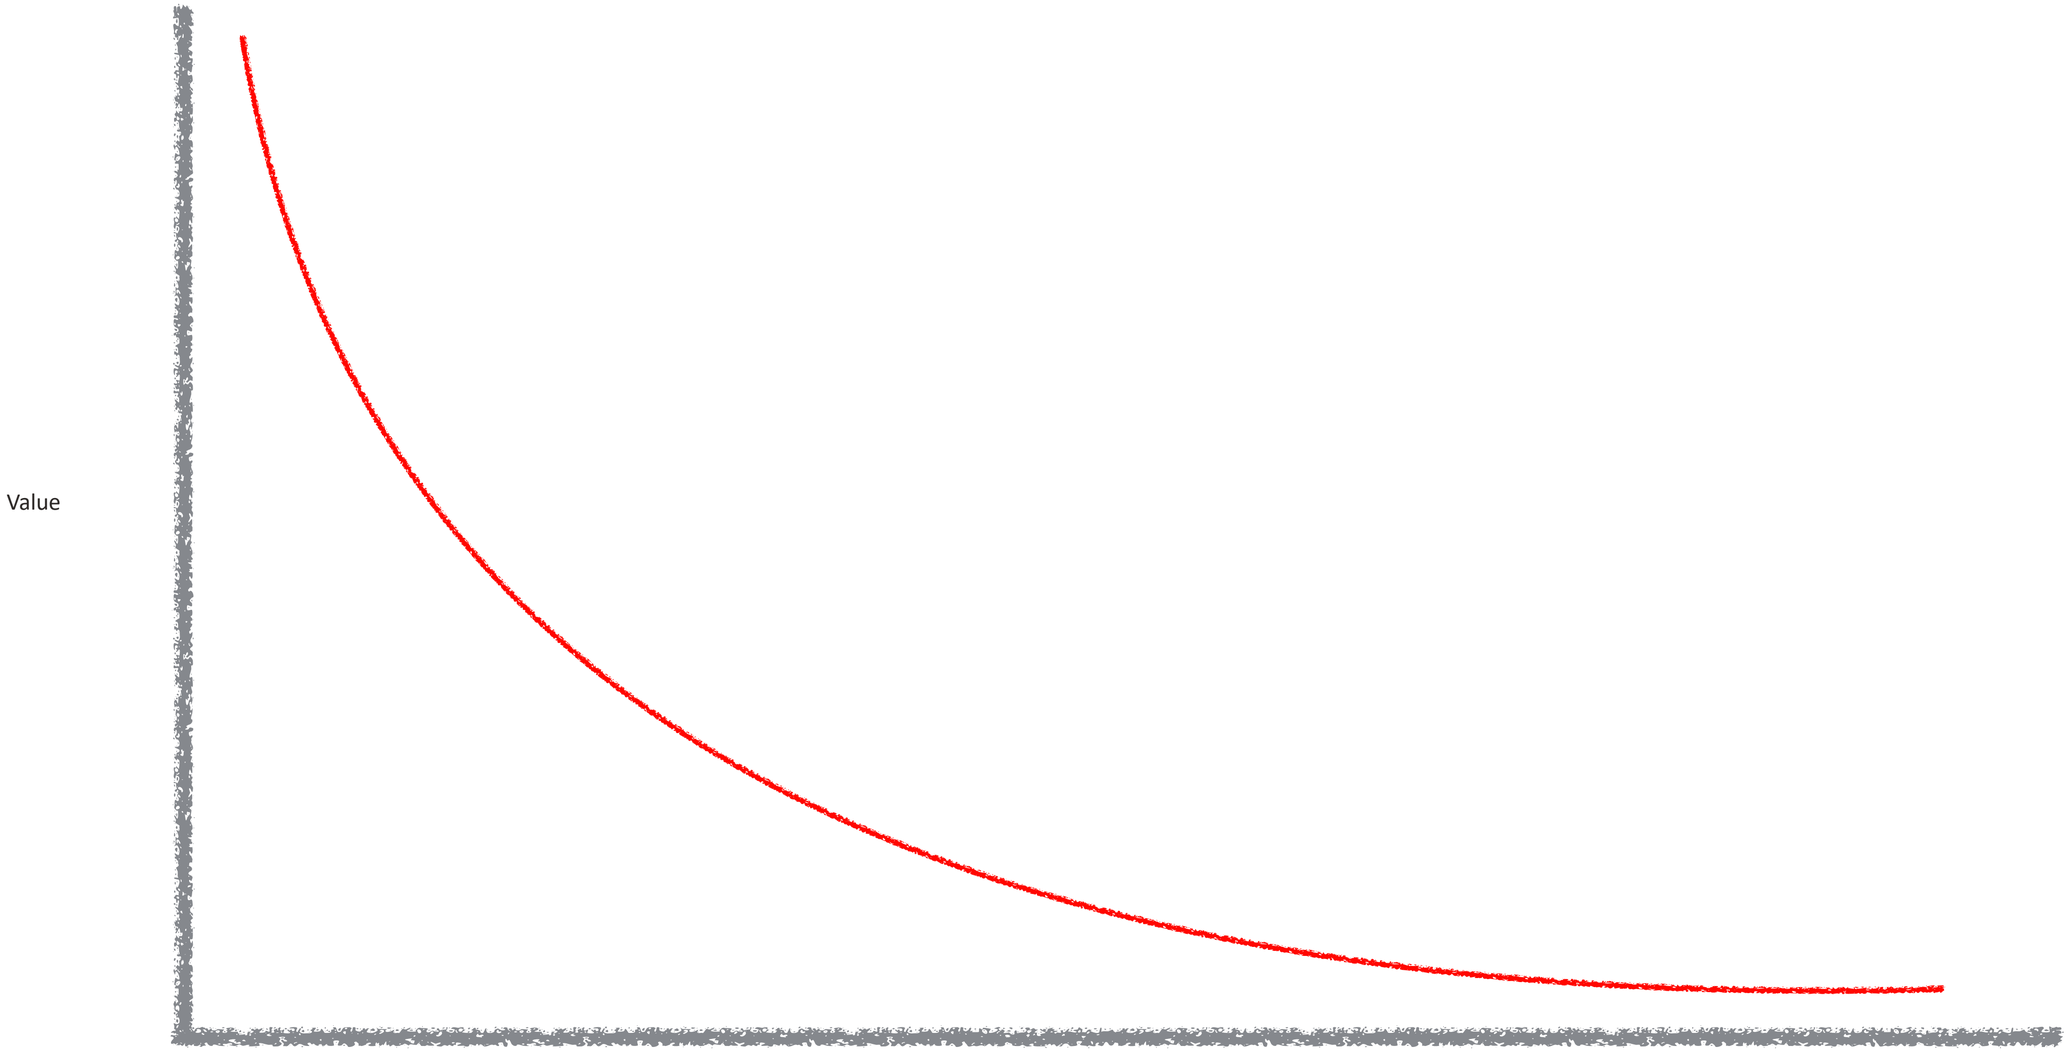

Costs of Insights Over Time

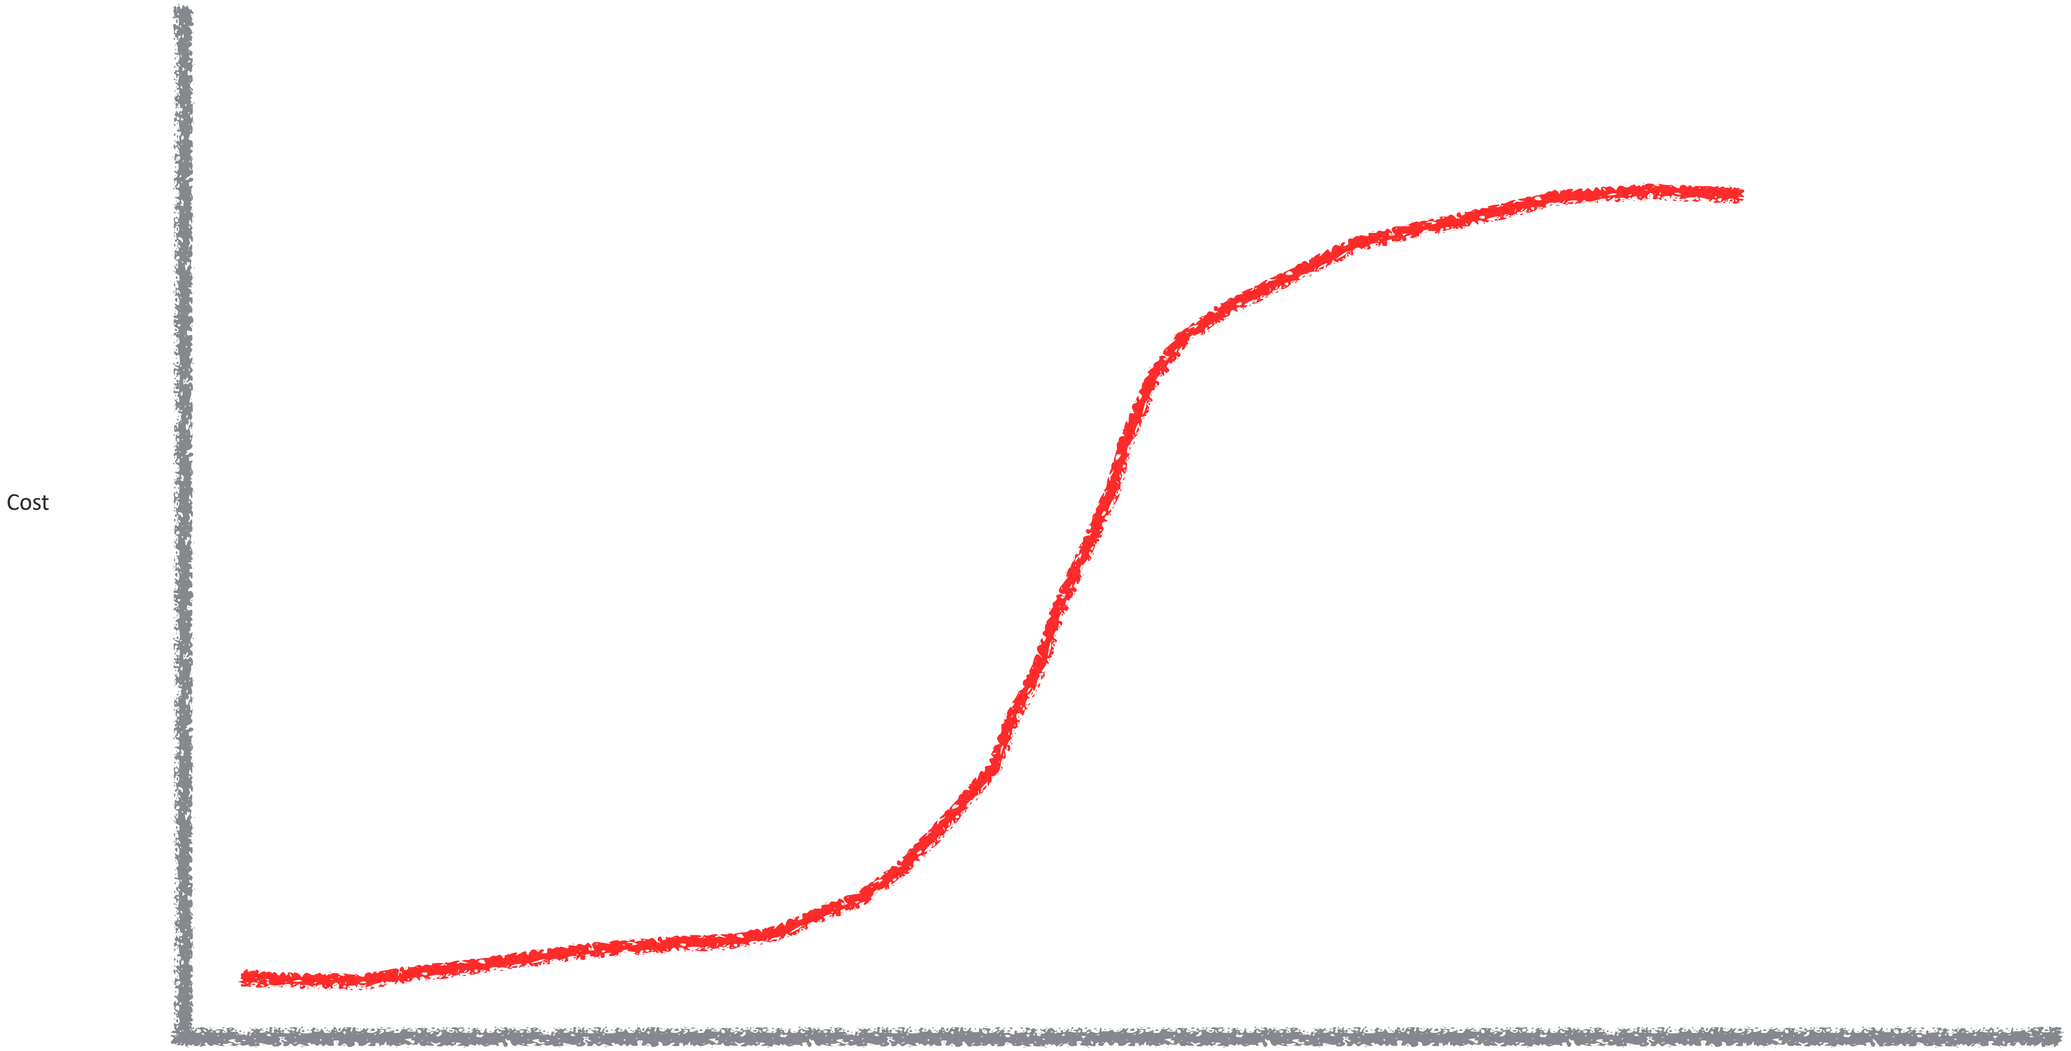

Cost of Insights Over Time

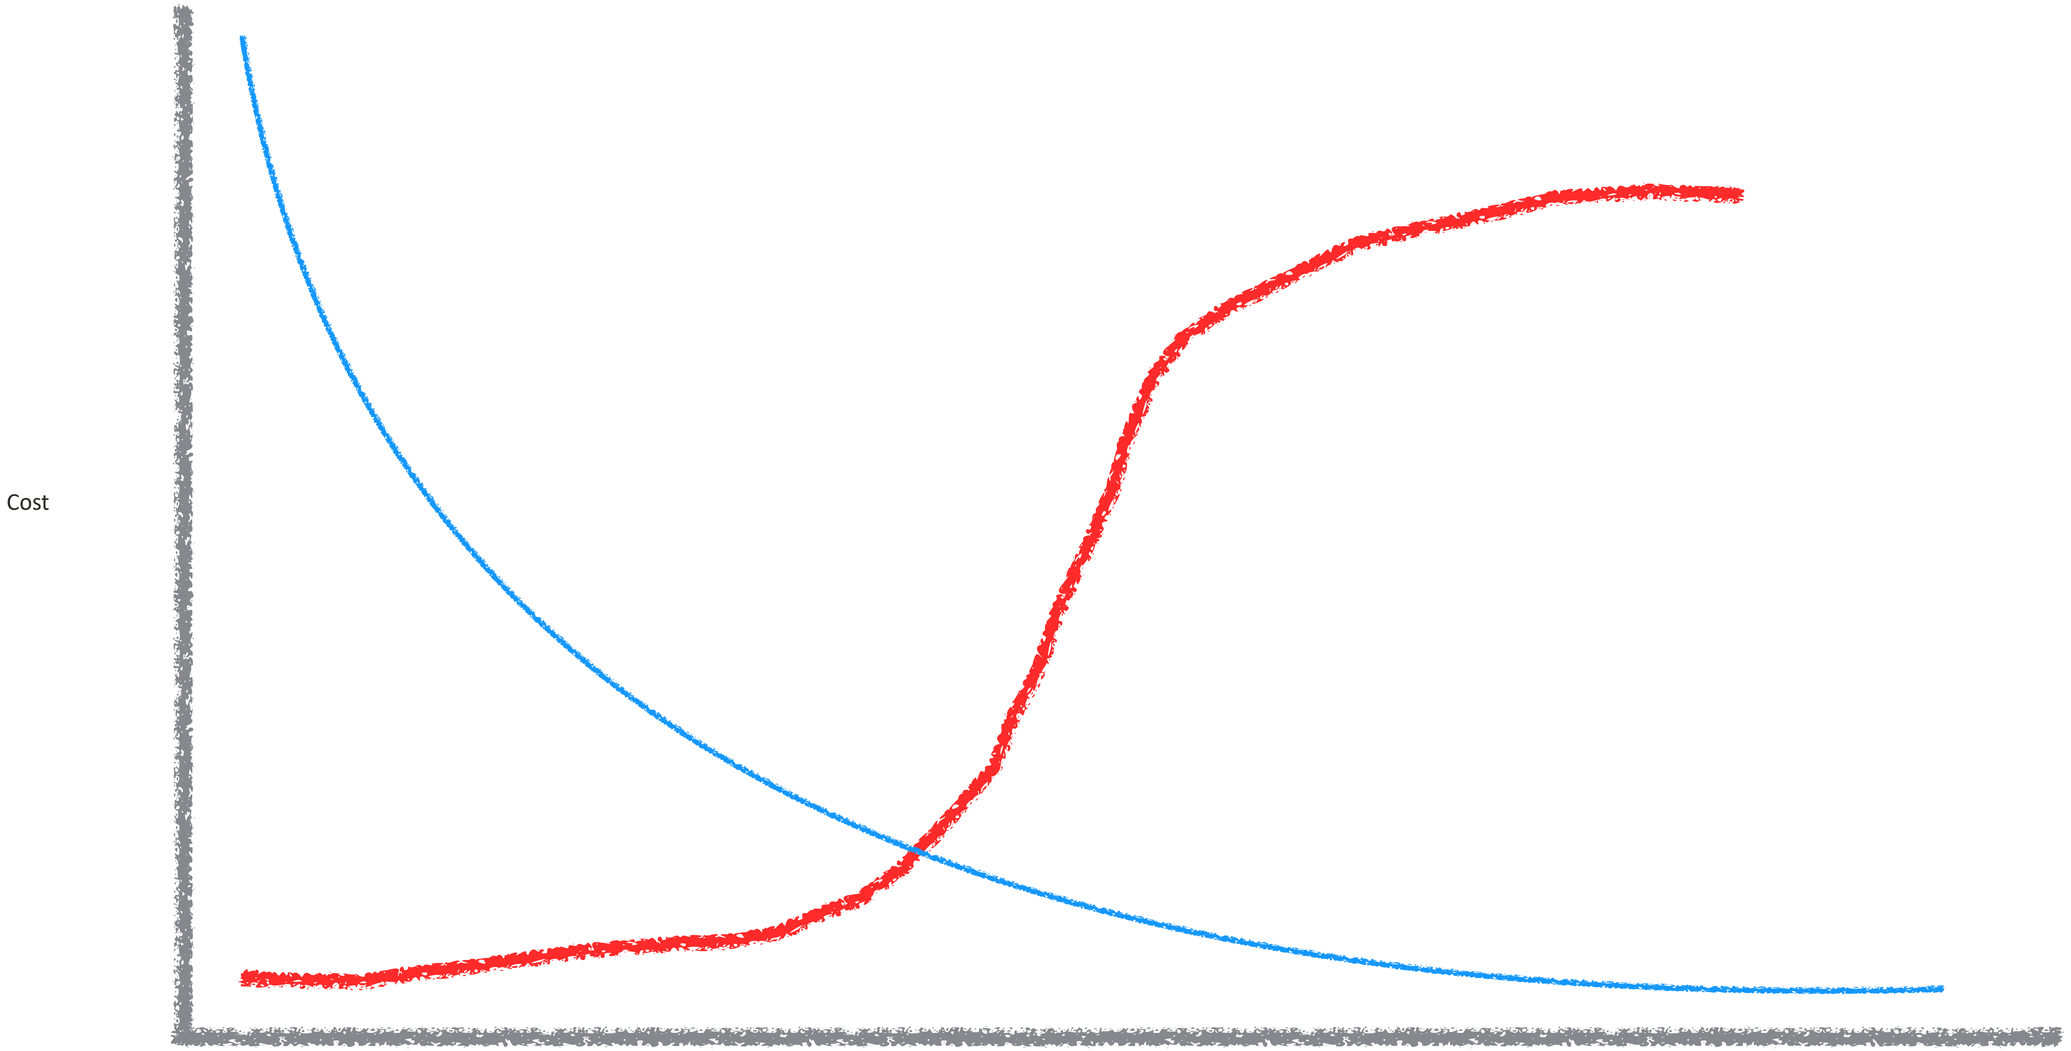

What Skills Does a Data Scientist Need?

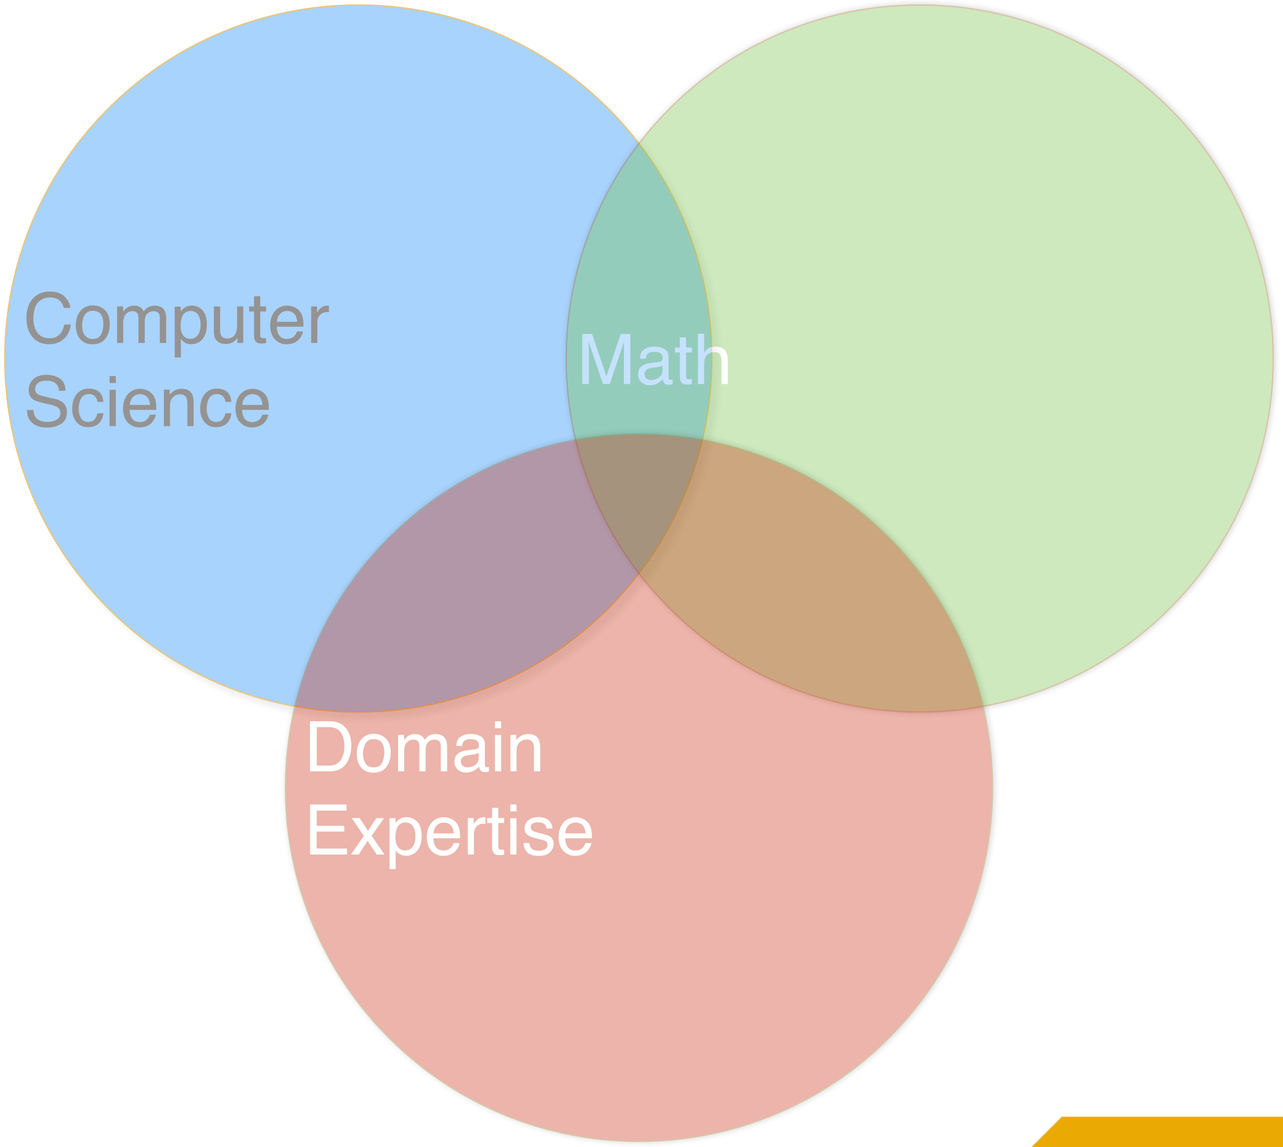
Computer
Science
Math
Domain
Expertise

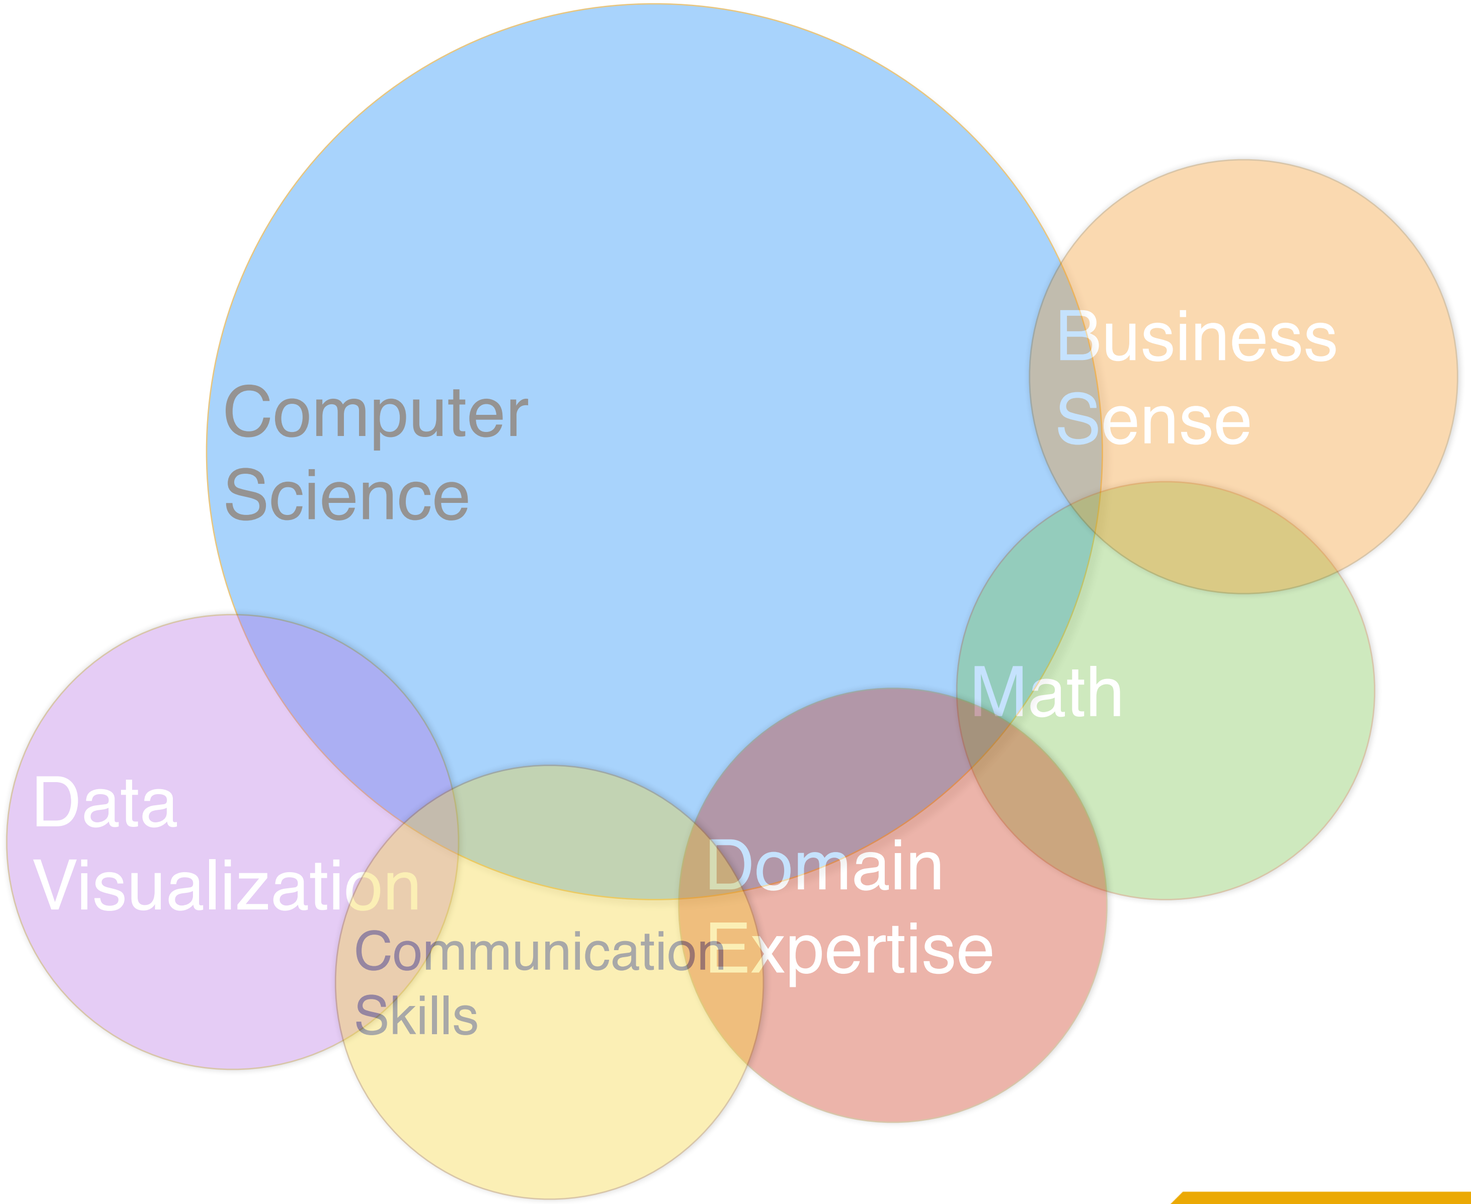
Computer Science
Business Sense
Math
Data Visualization
Communication Skills
Domain Expertise

The Data Scientist Venn Diagram

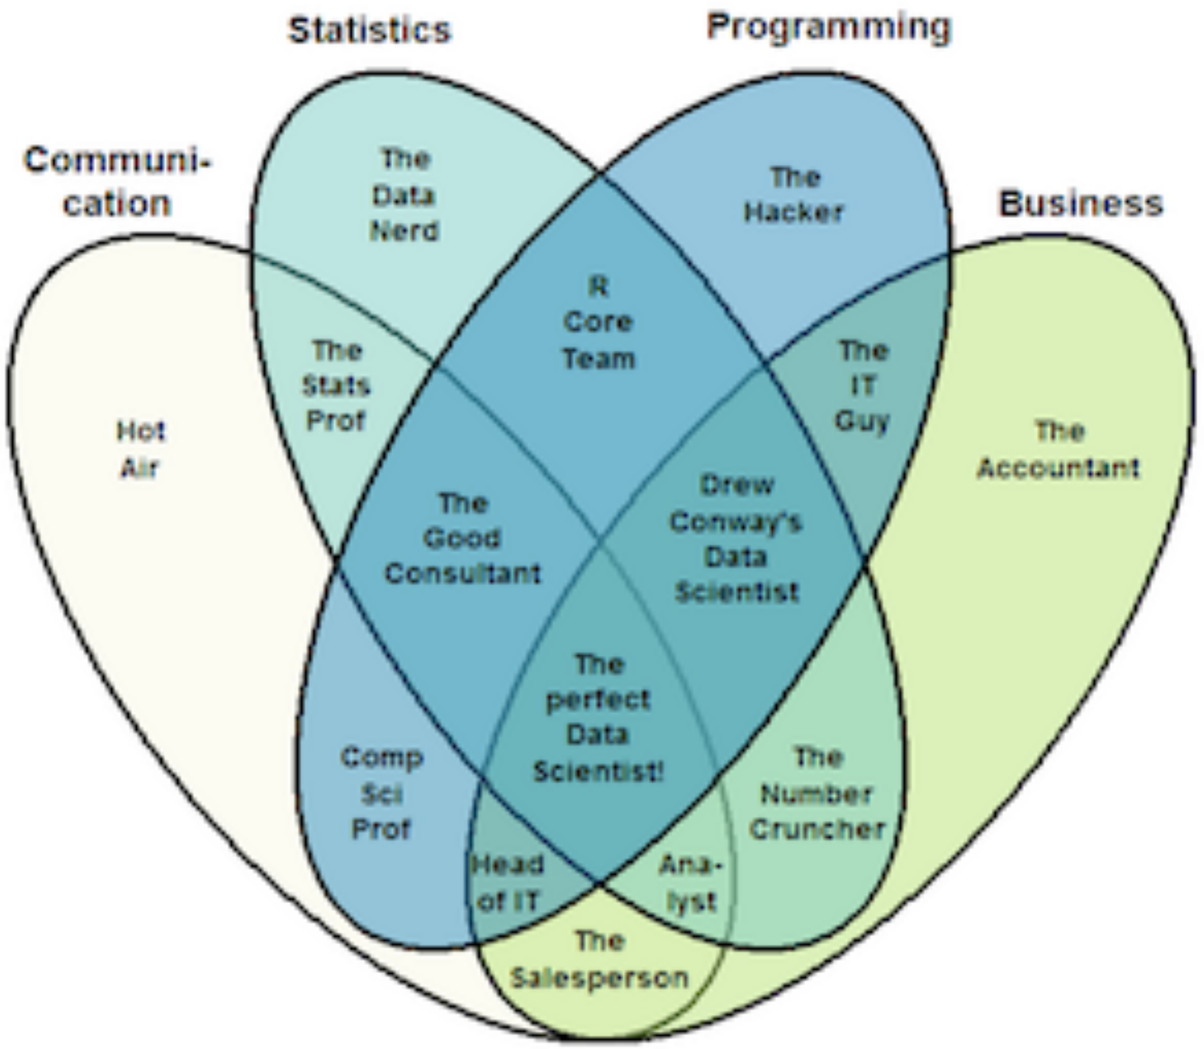

//www.kdnuggets.com/2017/03/data-science-data-scientist-do.html

Data Scientists spend 50-90% of their time being...

Data Janitors

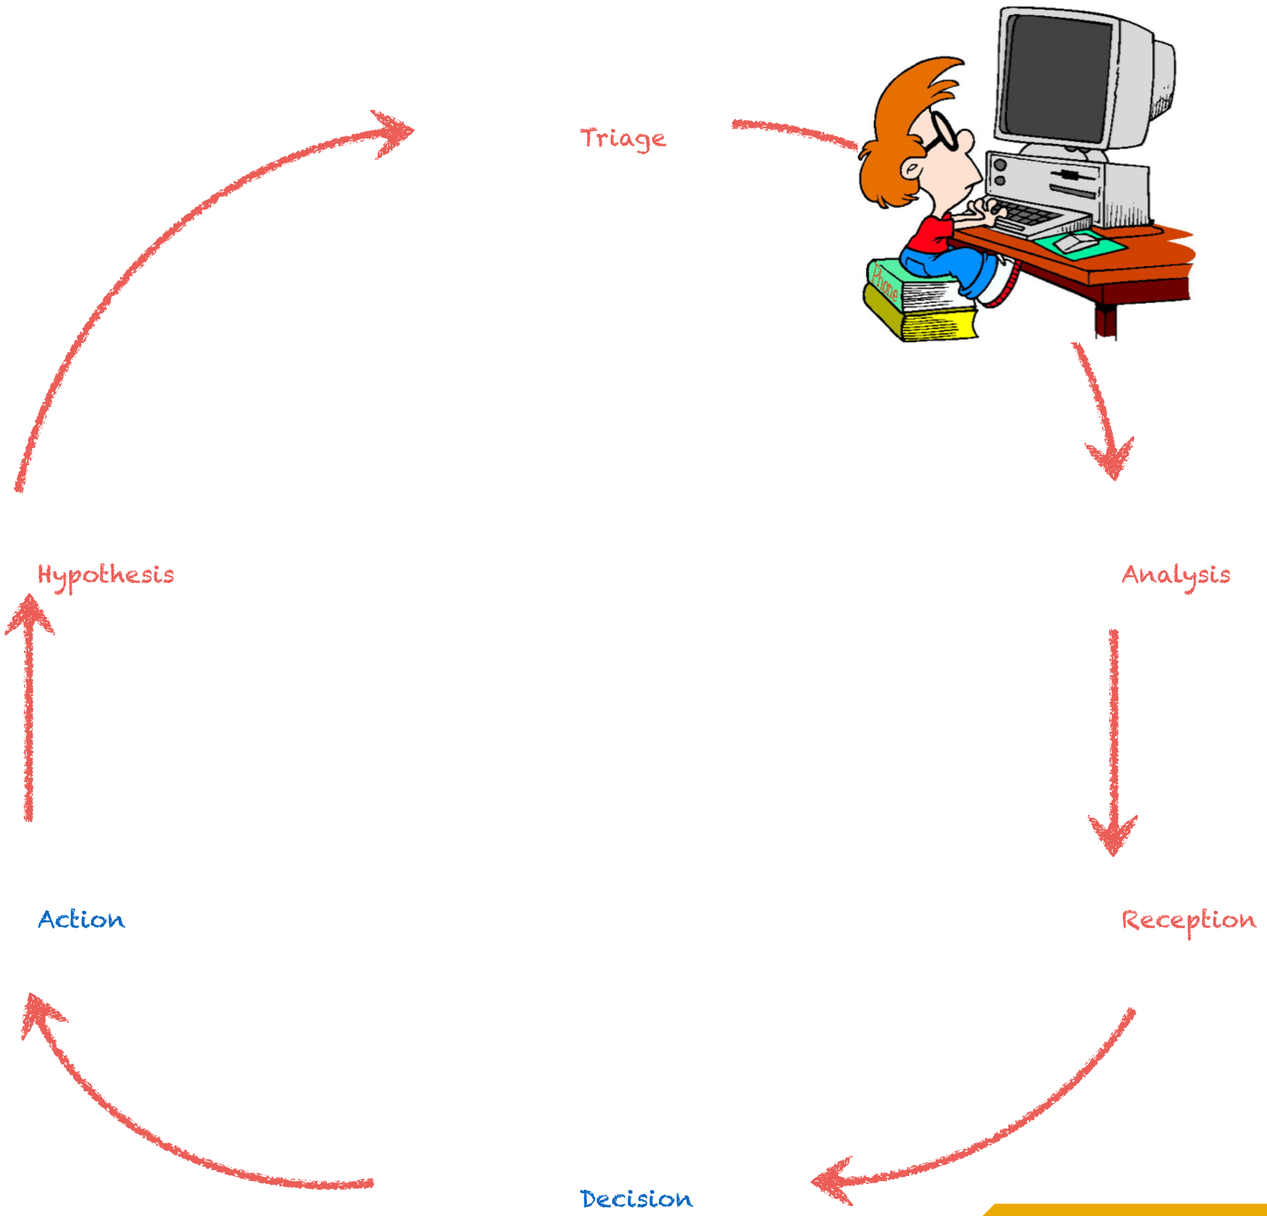
Triage
Analysis
Reception
Decision
Action
Hypothesis

Thoughts for Data Science Success

Data is a Strategic Asset... not a cost

Align Projects to Corporate Strategy

Align Projects to Corporate Strategy

Your:
Time
Money
Job?

Build the right team for Data Initiatives

Prioritize building appropriate data platform

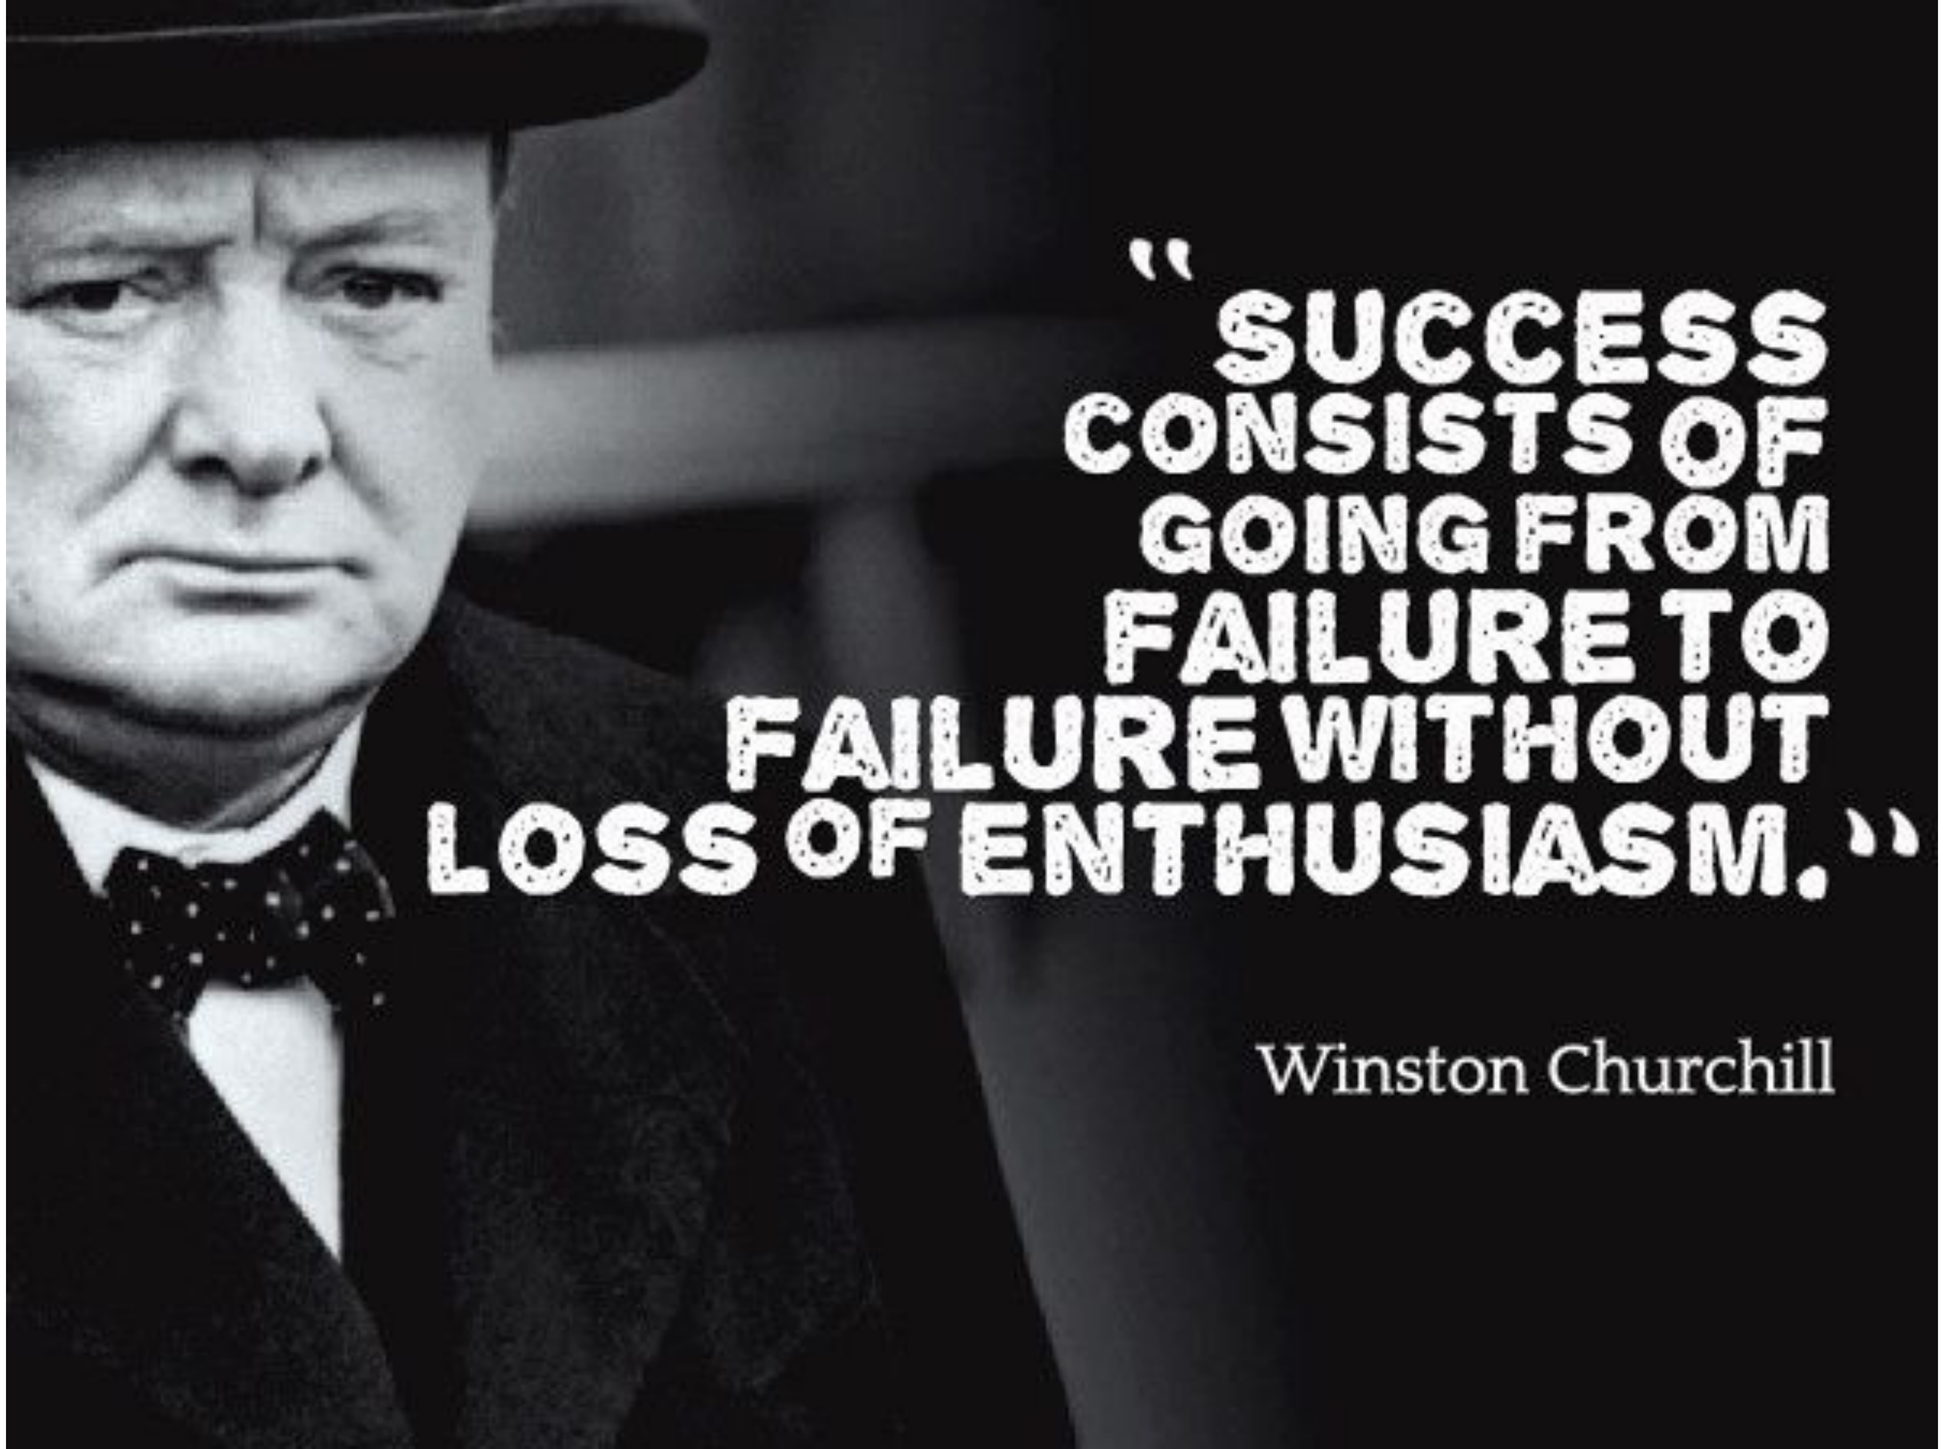
"SUCCESS CONSISTS OF GOING FROM FAILURE TO FAILURE WITHOUT LOSS OF ENTHUSIASM."
Winston Churchill

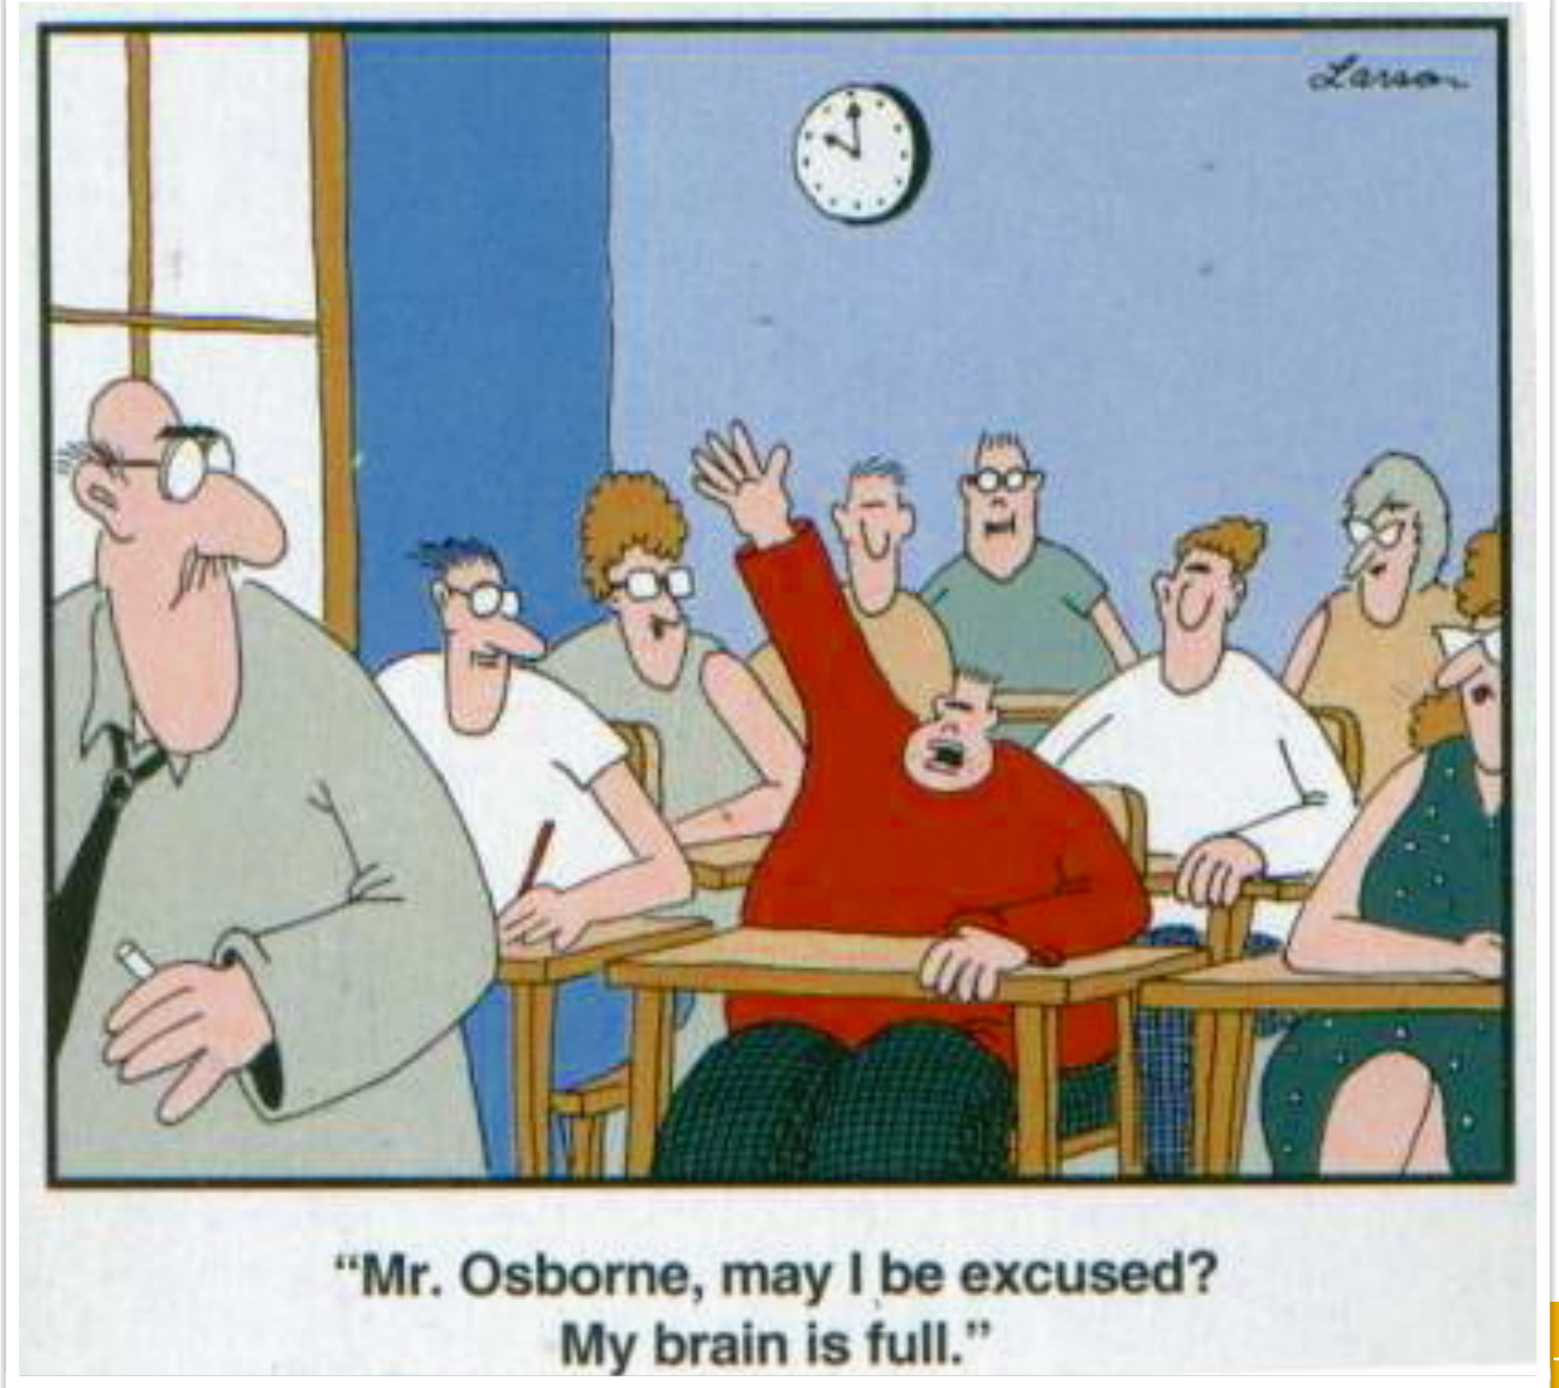
Larson
"Mr. Osborne, may I be excused?
My brain is full."

Building a Data Science Team

Building a Data Science Team

Programmer + Statistician + SME = DS Team?

Building a Data Science Team

Programmer + Statistician + SME = DS Team?

Sort of...

Building a Data Science Team

- Data Ambassador
- Data Scientist
- Data Storyteller
- Data Engineer

By The End of the Class, You Will Be Able To:

- Quickly and effectively prepare data for analysis
- Apply machine learning techniques to enhance security


I THINK WE SHOULD BUILD AN SQL DATABASE.
UH-OH
E-mail: SCOTTADAMS@AOL.COM
DOES HE UNDERSTAND WHAT HE SAID OR IS IT SOMETHING HE SAW IN A TRADE MAGAZINE AD?
© 1996 United Feature Syndicate, Inc.(NYC)
WHAT COLOR DO YOU WANT THAT DATABASE?
I THINK MAUVE HAS THE MOST RAM.


Buzzword

Recommended Reading

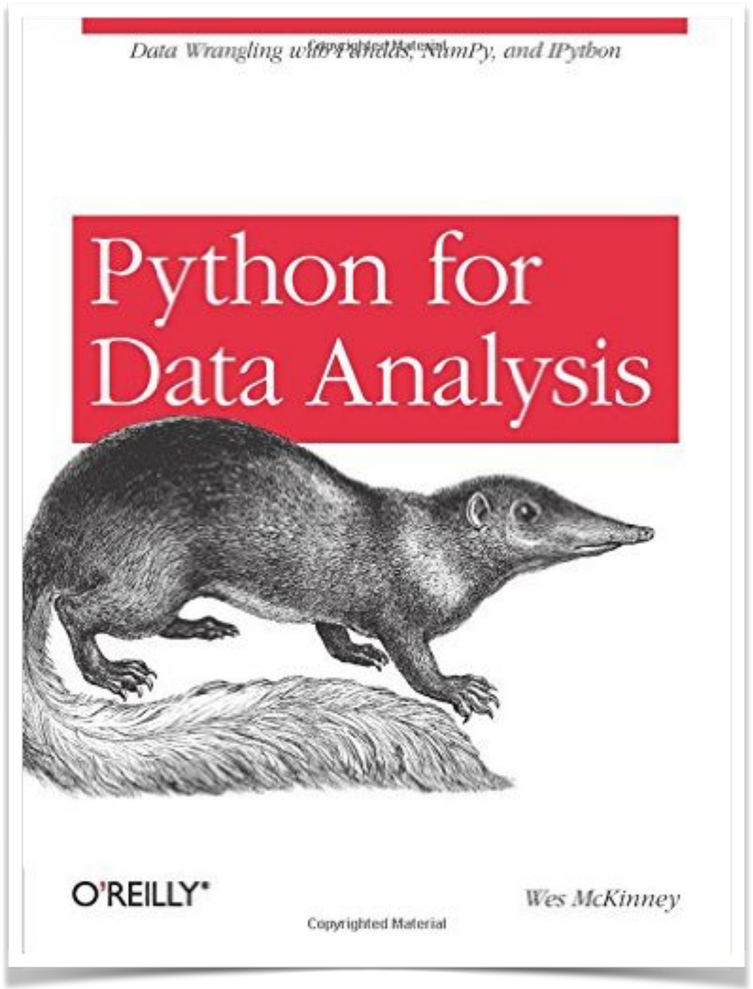

Recommended Reading

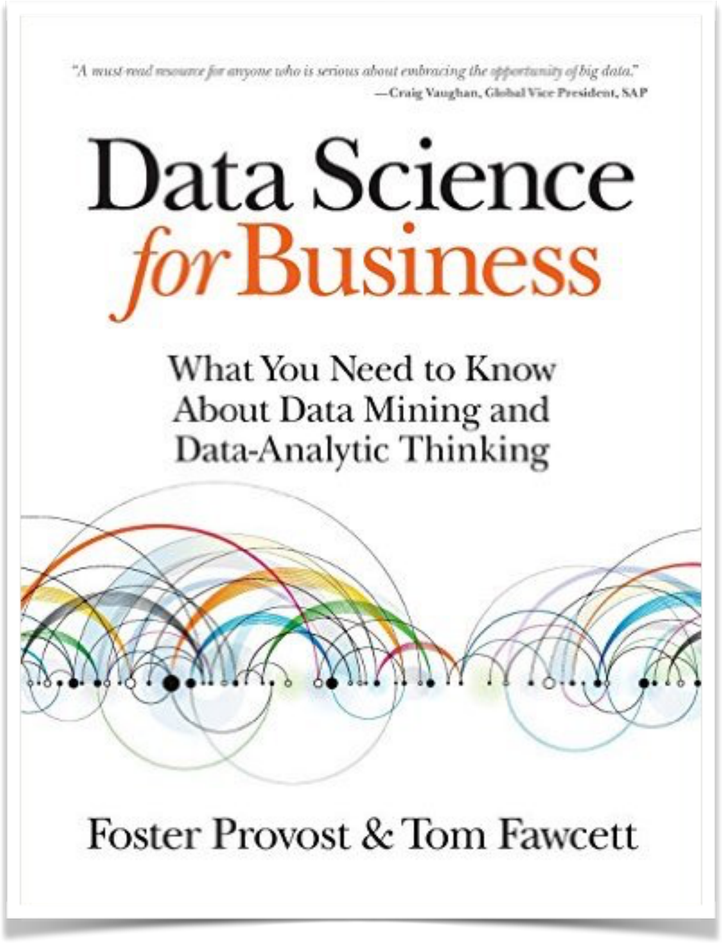

Recommended Reading

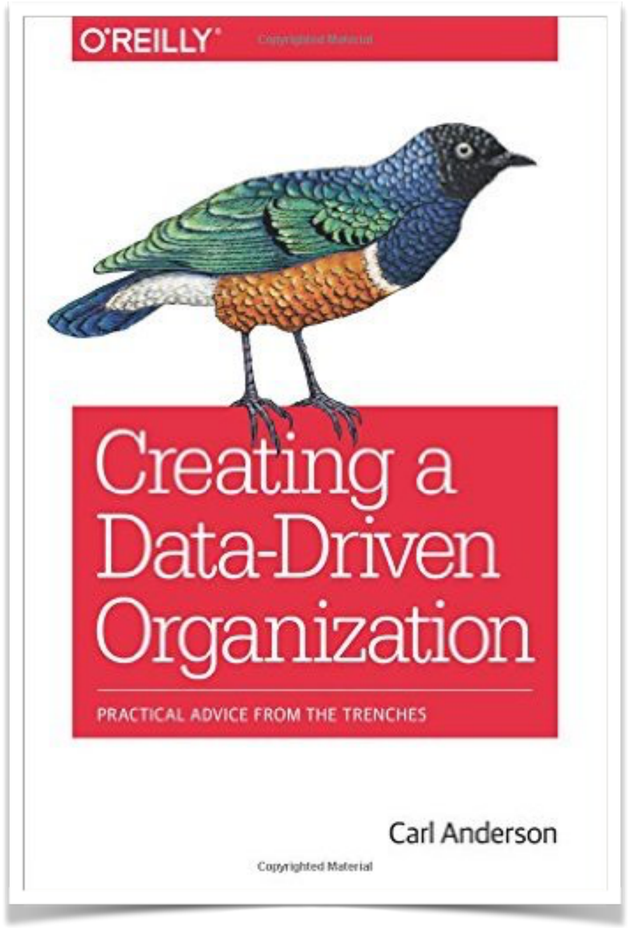

Recommended Reading

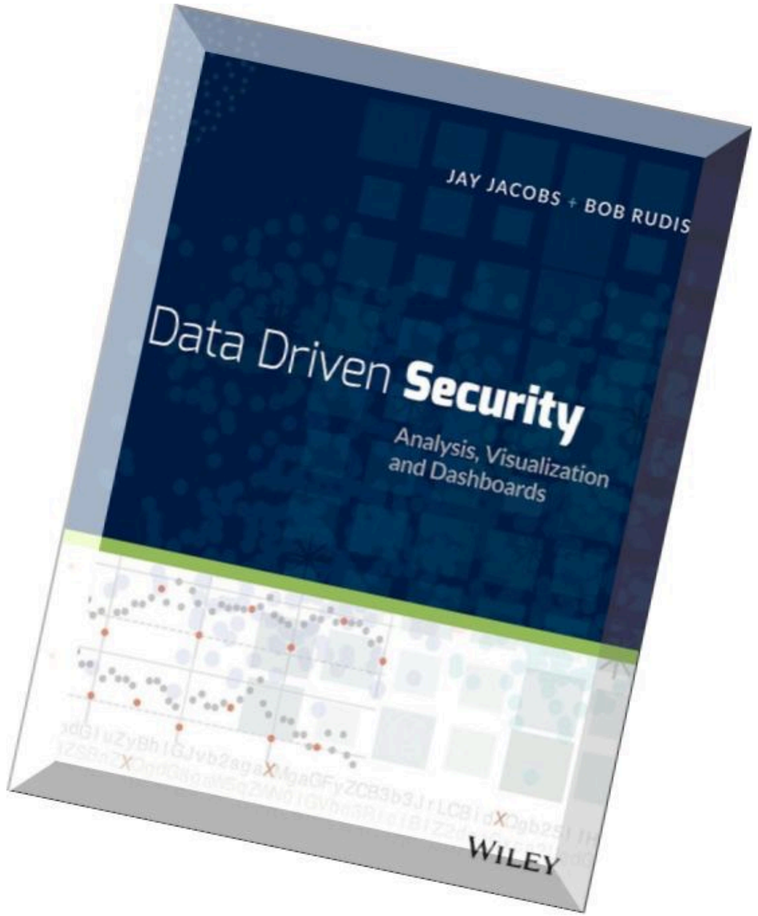

http://datadrivensecurity.info

Machine Learning

"Type a Quote Here"

-- Bill Gates

What is Machine Learning (ML) Artificial Intelligence (AI)

“Machine Learning is the science of getting computers to act without being explicitly programmed.”

– https://www.coursera.org/course/ml

“A computer program is said to learn from experience E with respect to some task T and some performance measure P, if its performance on T, as measured by P, improves with experience E.”

–Tom Mitchell, Carnegie Mellon University

“Machine learning explores the construction and study of algorithms that can learn from and **make predictions on data.** Such algorithms operate by building a model from example inputs in order to make data-driven predictions or decisions, **rather than following strictly static program instructions.**”

–https://en.wikipedia.org/wiki/Machine_learning


SPAM®
CHOPPED PORK AND HAM
LICENCED BY:
Hormel
Foods
Serving Suggestion


SPAM
CHOPPED PORK AND HAM
LICENCED BY:
Hormel Foods
Serving Suggestion


SPAM
TERIYAKI
Glorious SPAM!


SPAM
LITE
Glorious SPAM!


SPAM
Oven Roasted TURKEY
Hormel Foods

- Blacklists
- Simple keyword matching
- Naive Bayesian Classifiers
- Deep Learning

Artificial Intelligence

Any technique which enables computers to mimic human behavior.

Machine Learning

Subset of AI techniques which use statistical methods to enable machines to improve with experiences.

Deep Learning

Subset of ML which make the computation of multi-layer neural networks feasible.

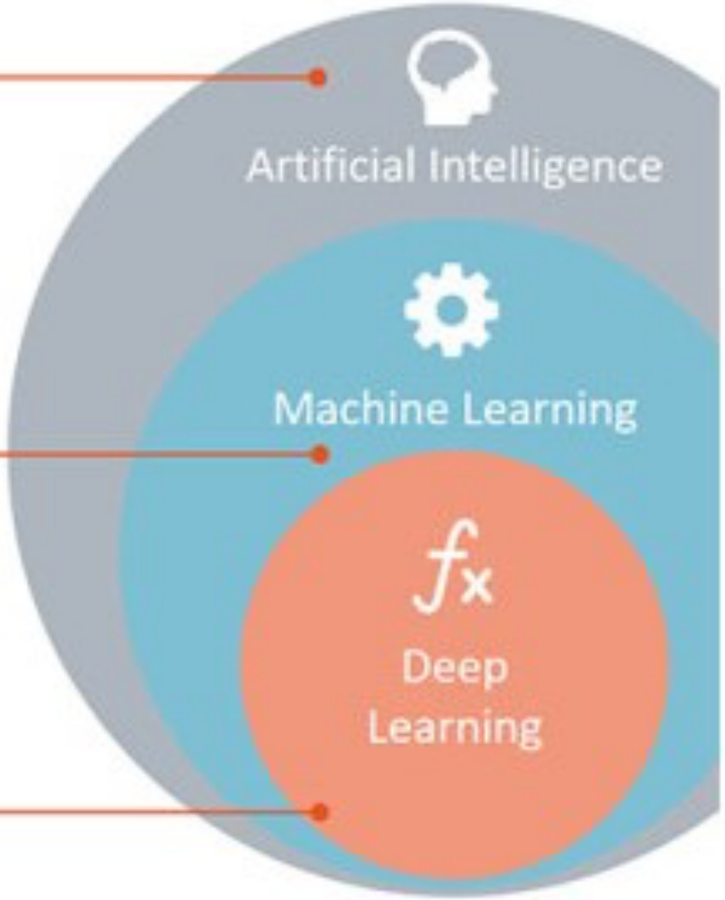

https://www.kdnuggets.com/2017/07/rapidminer-ai-machine-learning-deep-learning.html

@katherinebailey Because marketing? Every time someone calls simple linear regression “AI” Gauss turns over in his grave.

Machine Learning Problems

- **Supervised Learning:** Supervised Learning is a class of Machine Learning in which a model is "trained" using a set of pre-existing labeled data.
- **Unsupervised Learning:** A class of Machine Learning algorithms in which a model is built without the use of labeled data.

Machine Learning Problem Types

- **Classification**: Assigning or predicting a observation's membership in discrete class
- **Regression:** Predicting a continuous value based on the observations' features
- **Clustering:** Identifying groupings within a dataset
- **Dimensionality Reduction:** Reducing the number of variables in a feature set

What Problem am I solving?

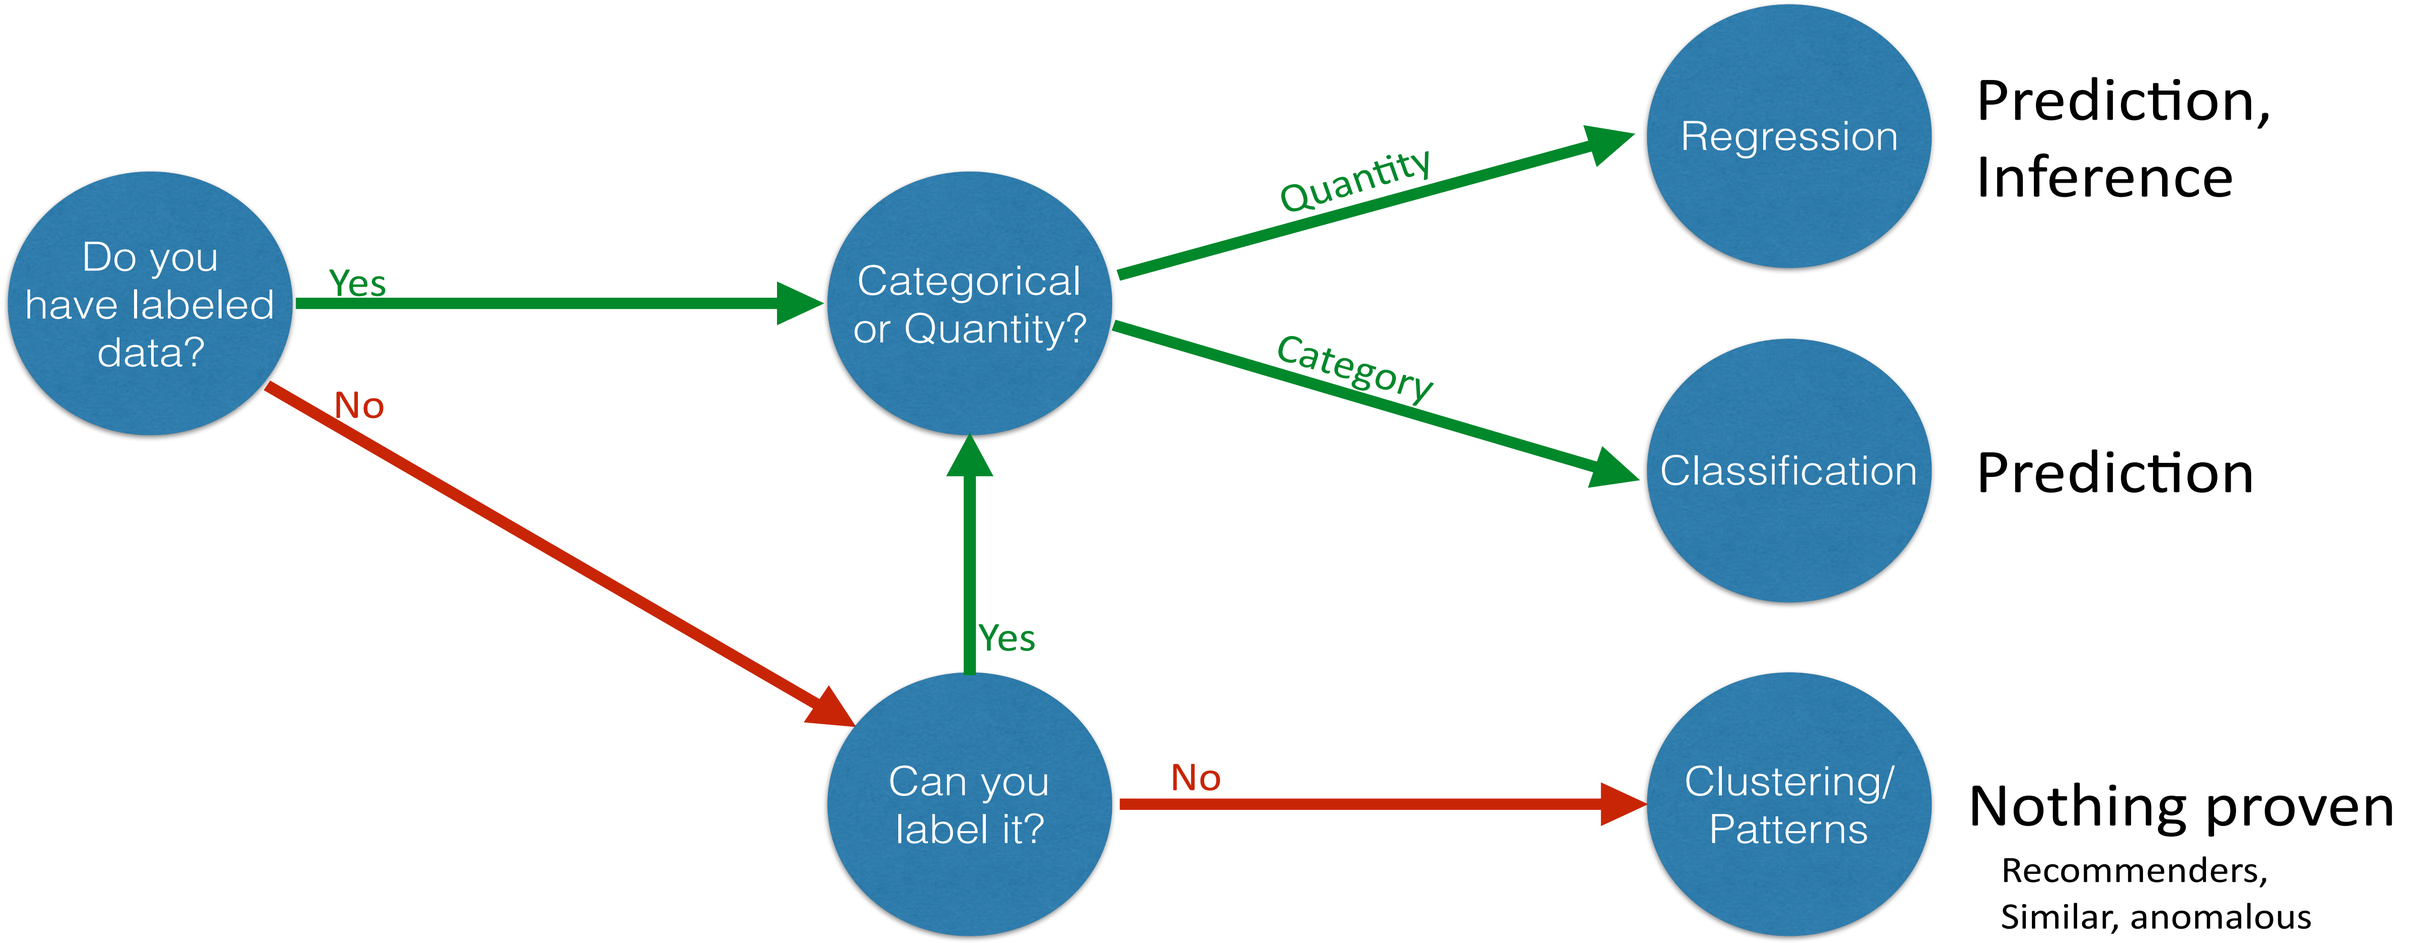

What it is Not

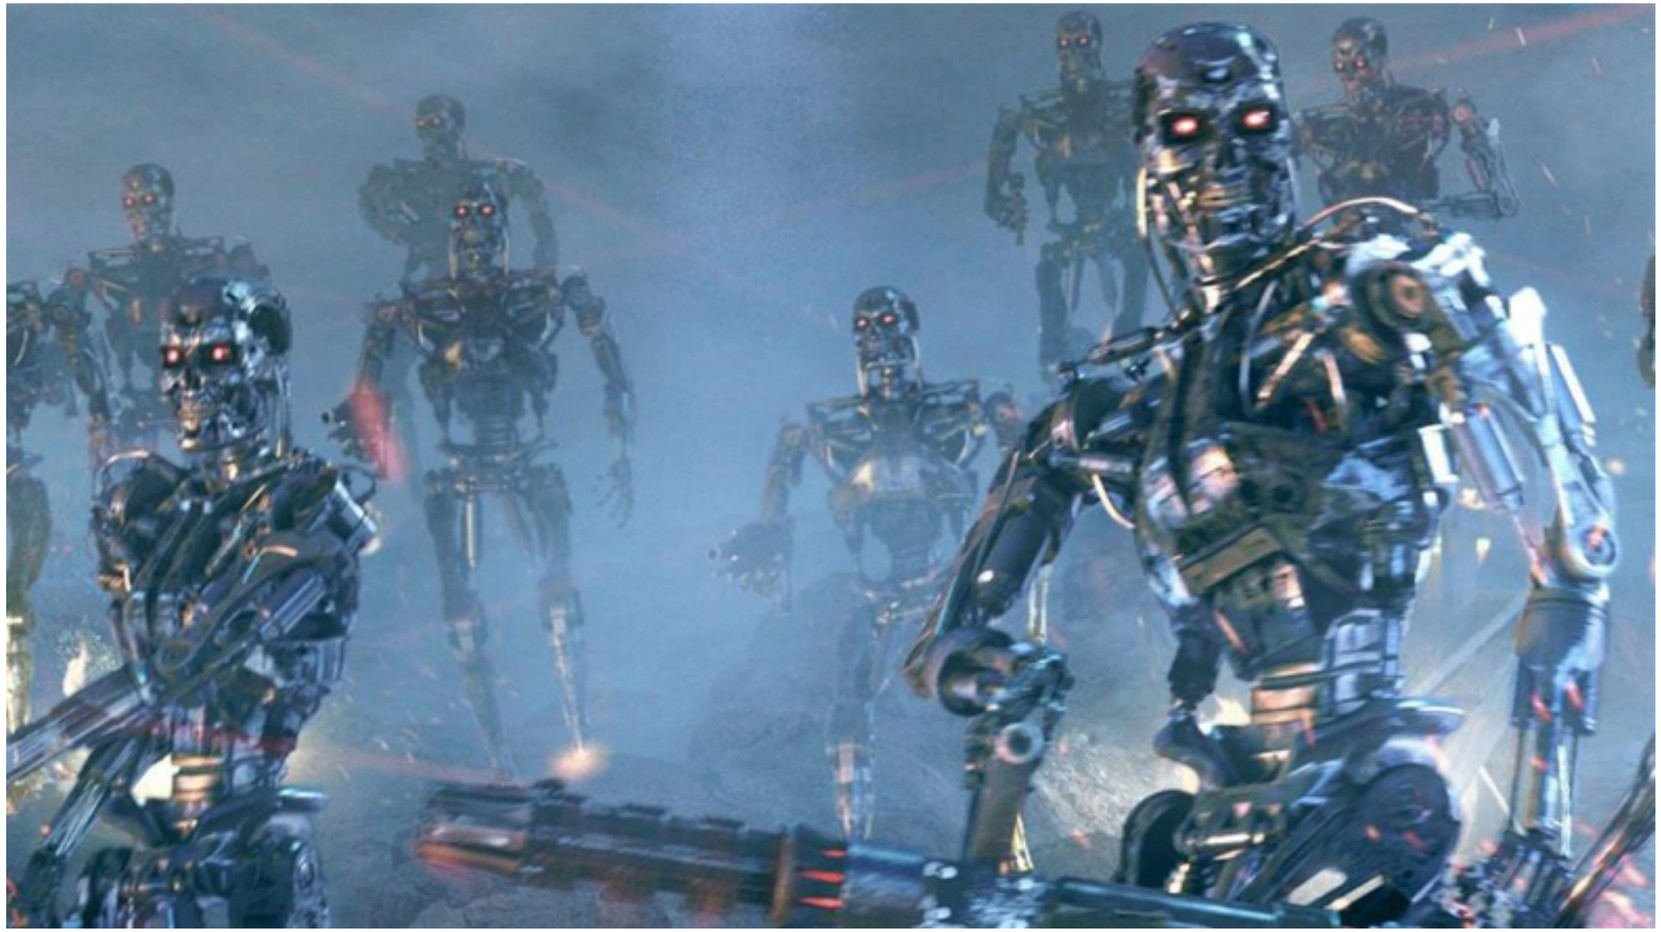

Applications to Security

Regression Example

Server Capacity Prediction: Regression analysis can be used to predict a server's capacity (or CPU usage) based on the server's historical performance.

https://www.researchgate.net/publication/256645877_LiRCUP_Linear_Regression_based_CPU_Usage_Prediction_Algorithm_for_Live_Migration_of_Virtual_Machines_in_Data_Centers

https://jgreenemi.com/predicting-capacity-with-linear-regression-ml/

Clustering Example

Anomaly Detection: Clustering techniques can be used to detect
ıything really.

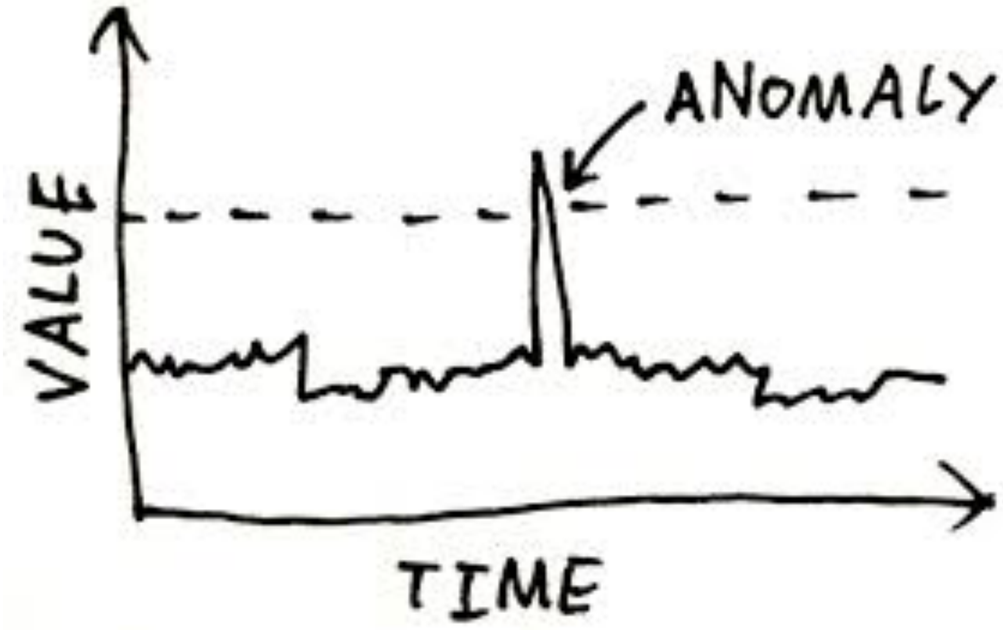

Network-Based Intrusion Detection

- Derive Features from Network Traffic Captures “pcap” at packet level or NetFlow level (tools: tshark, tcpdump, bro…)
- Example Features based on header information: 2s-windowing of connections > duration, protocol, src and dat bytes, service.
- Get data sets: http://www.netresec.com/?page=PcapFiles, https://maccdc.org/, http://www.westpoint.edu/crc/SitePages/DataSets.aspx http://www.unb.ca/cic/research/datasets/,

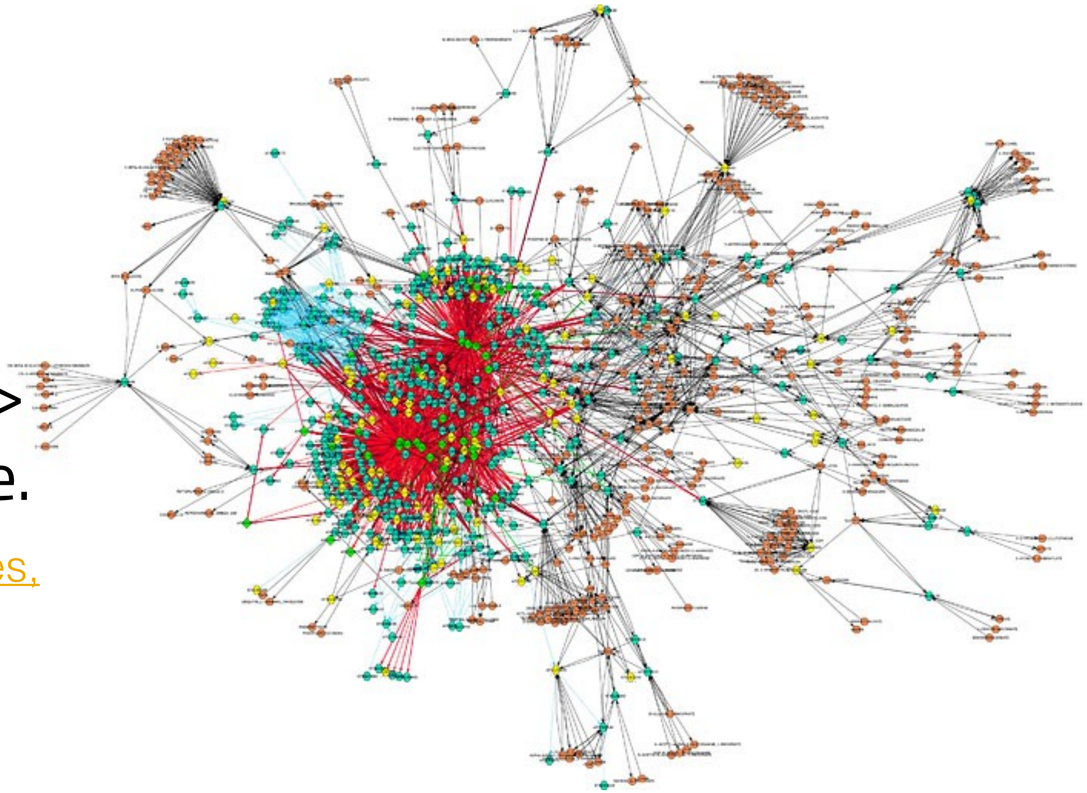

Malware Detection/Classification

- Derive Features from Binary Content and metadata manifest (function calls, string obtained from IDA Disassembler)
- Example Features: opcode count (n-grams), segment count, asm pixel intensity, n-gramming of bytes, function name.
- Featureless Deep Learning with word2vec embedding
- Get open source malware samples: Vx Heaven, Virus Share, Maltrieve, Open Malware

Security Applications of Machine Learning

- Domain Generation Algorithm (DGA) Detection (Classification)
- Malicious URL Detection (Classification)
- Network Traffic: Beaconing Detection (Classification/Clustering)
- Detection of new classes of malware (Classification/Clustering)
- General Network Traffic Anomaly Detection (Classification/Clustering)
- Log Analysis - Anomaly Detection (Classification/Clustering)
- Phishing Detection (Classification)
- Identifying SQL Injection (Classification)
- Identifying XSS cross-site scripting (Classification)
- DOS/DDOS Detection (Classification)
- Authentication (Classification)

Data Science Hunting Funnel

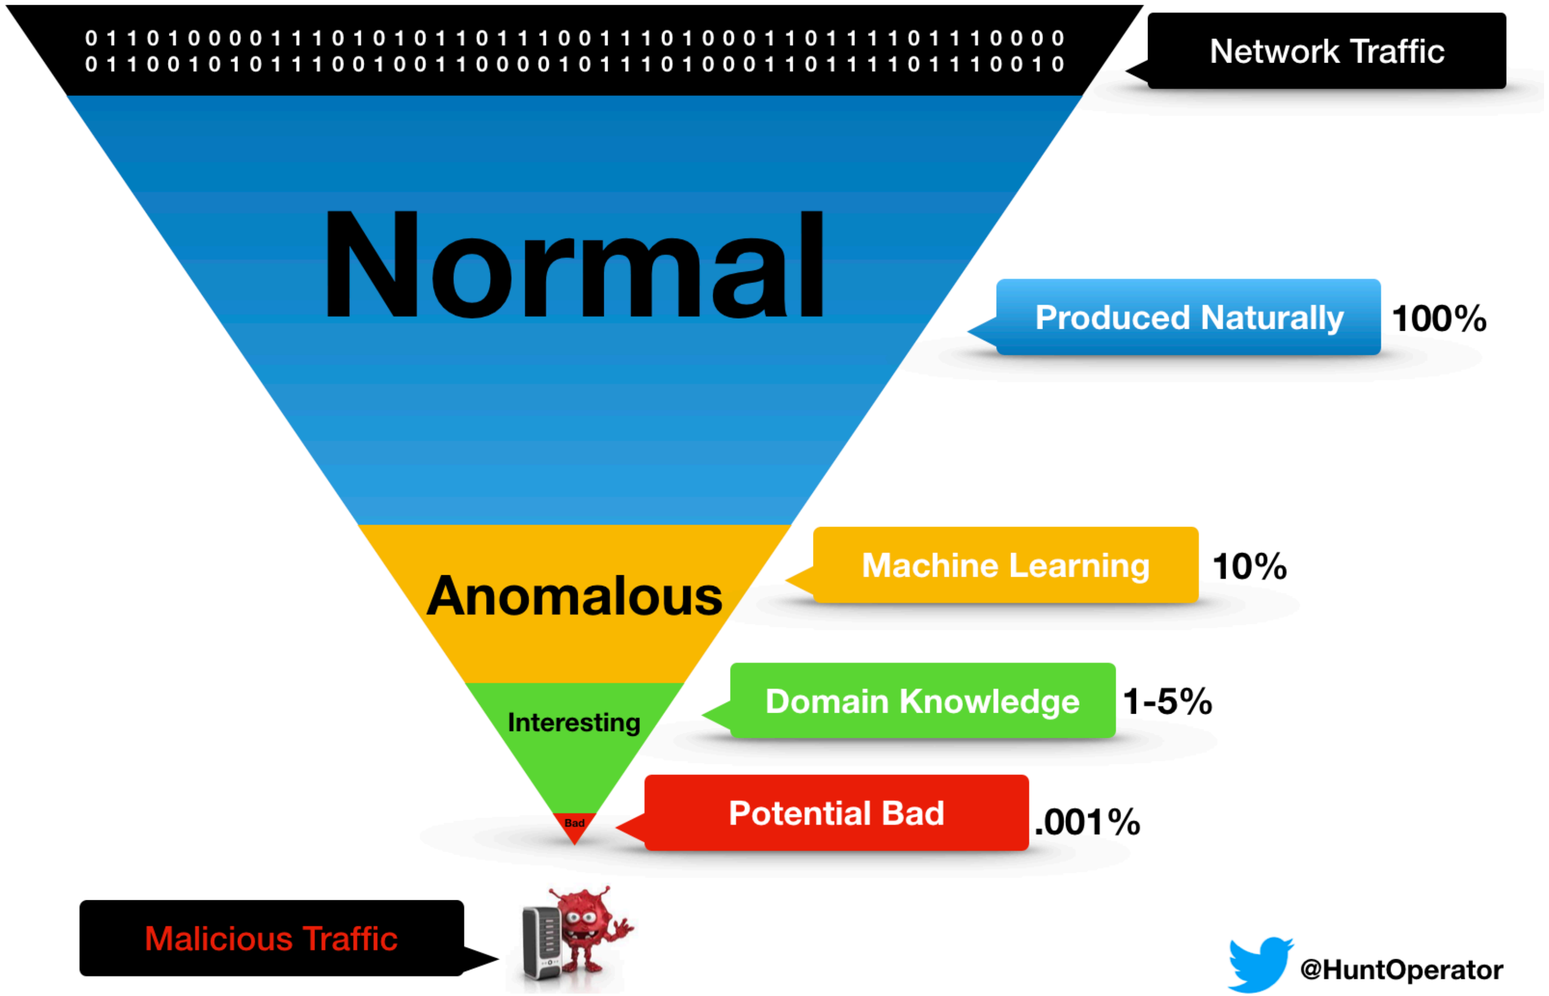

@HuntOperator

http://www.austintaylor.io/network/traffic/threat/data/science/hunting/funnel/machine/learning/domain

The Machine Learning Process

Supervised Machine Learning Process

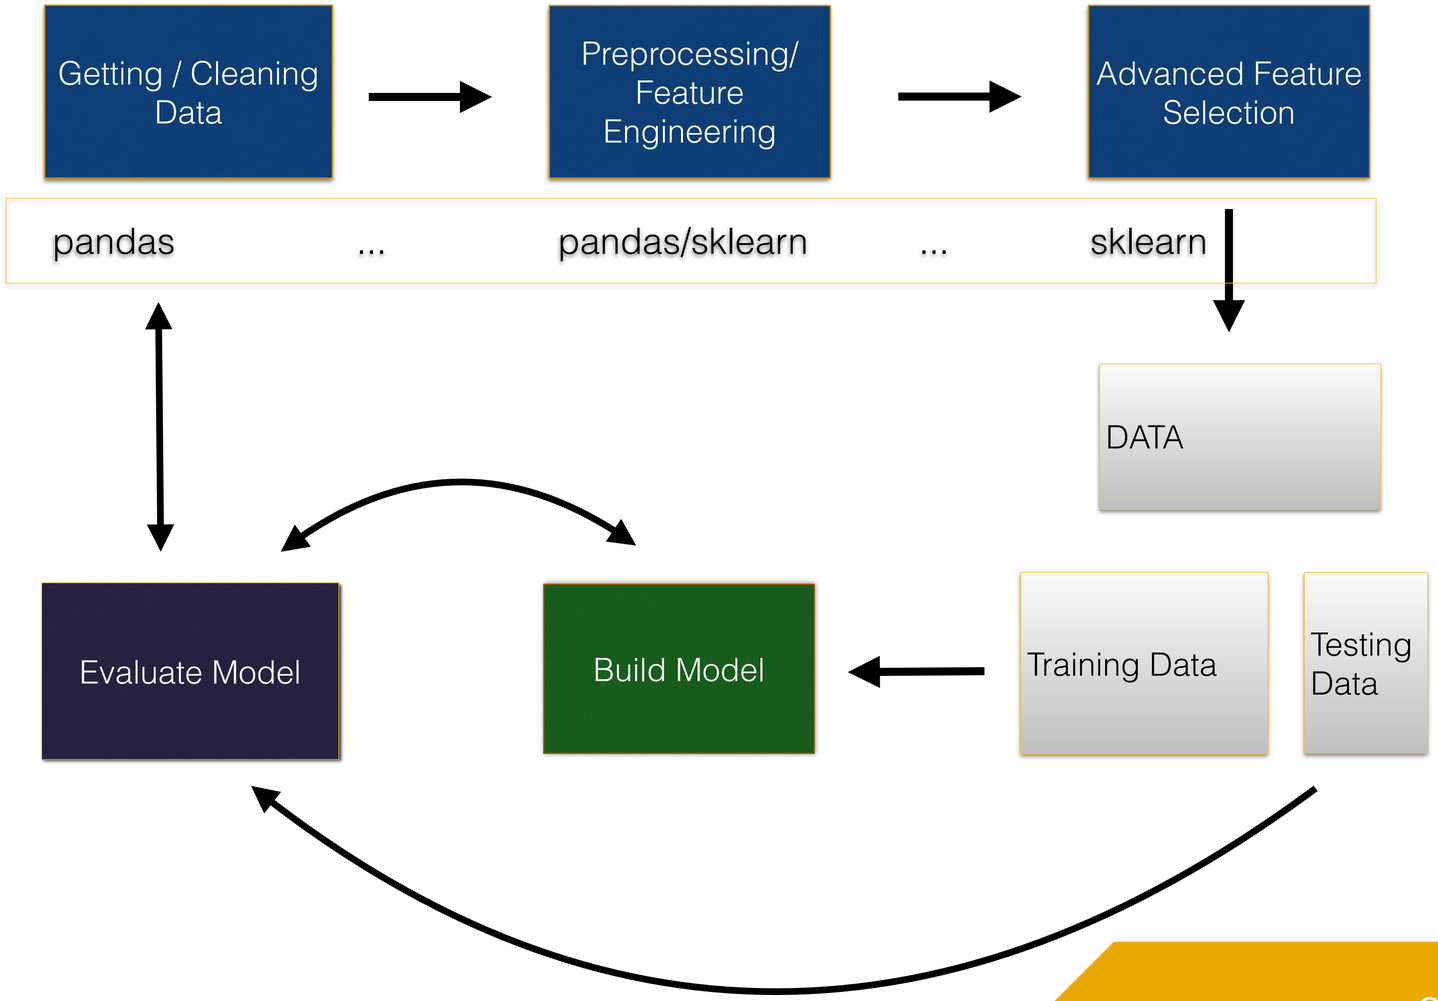

Unsupervised Machine Learning Process

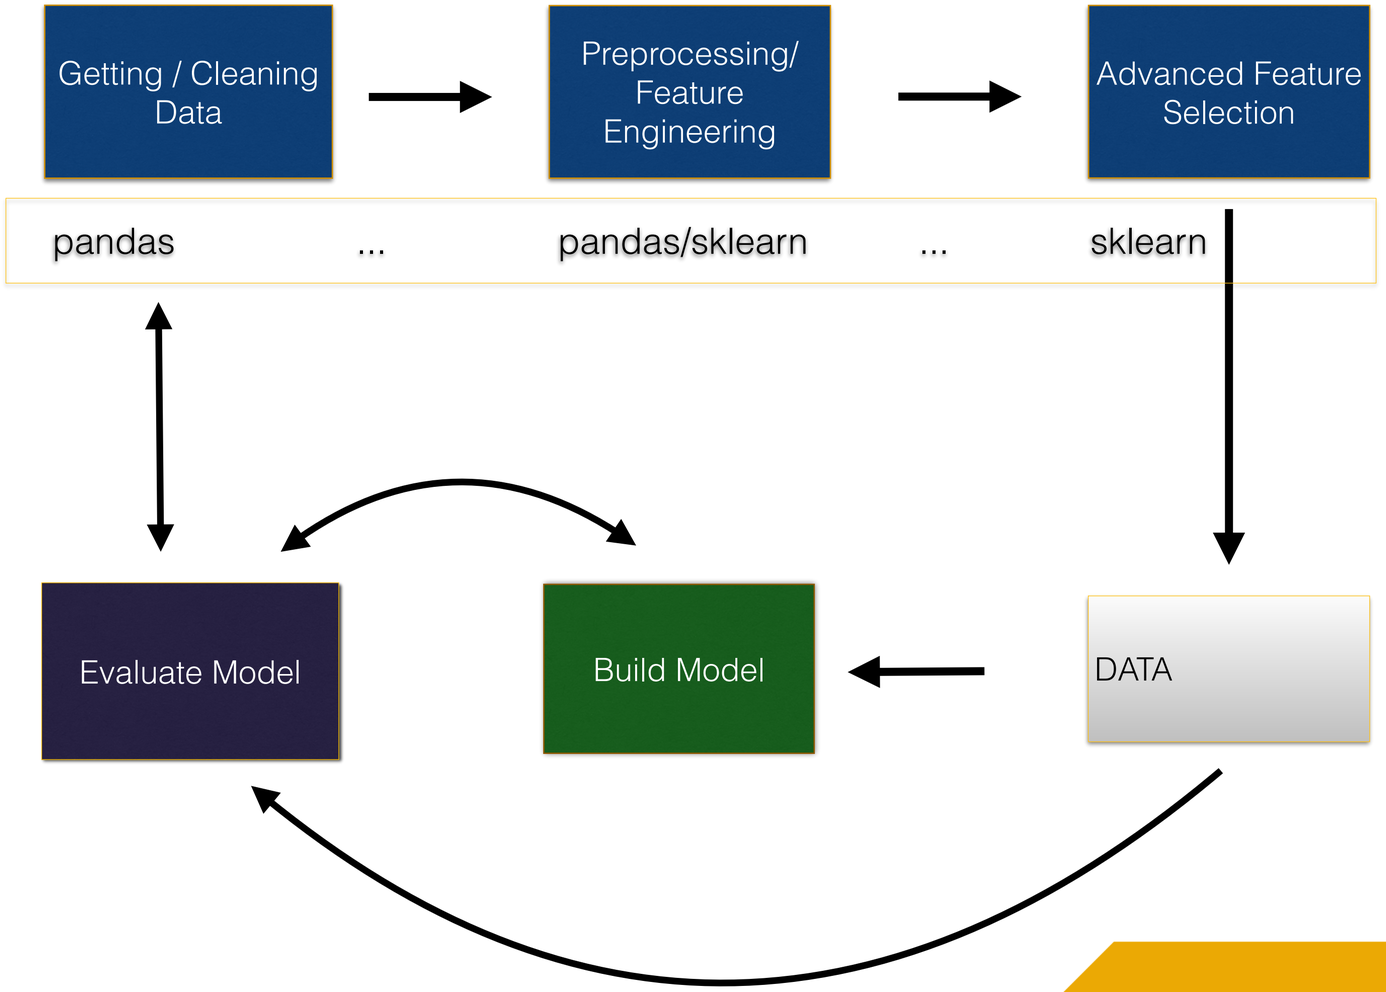

First, define your analytic question.

What are you trying to do?

How do you define success? What are you measuring?

Choose data sources

- What is available?
- Is it enough?
- Is the data reliable/clean/consistent?
- What other data could you use?

Other Considerations

- Policies
- Legal contraints
- Biases in Data
- Latency
- Data size

Gather and Explore Your Data

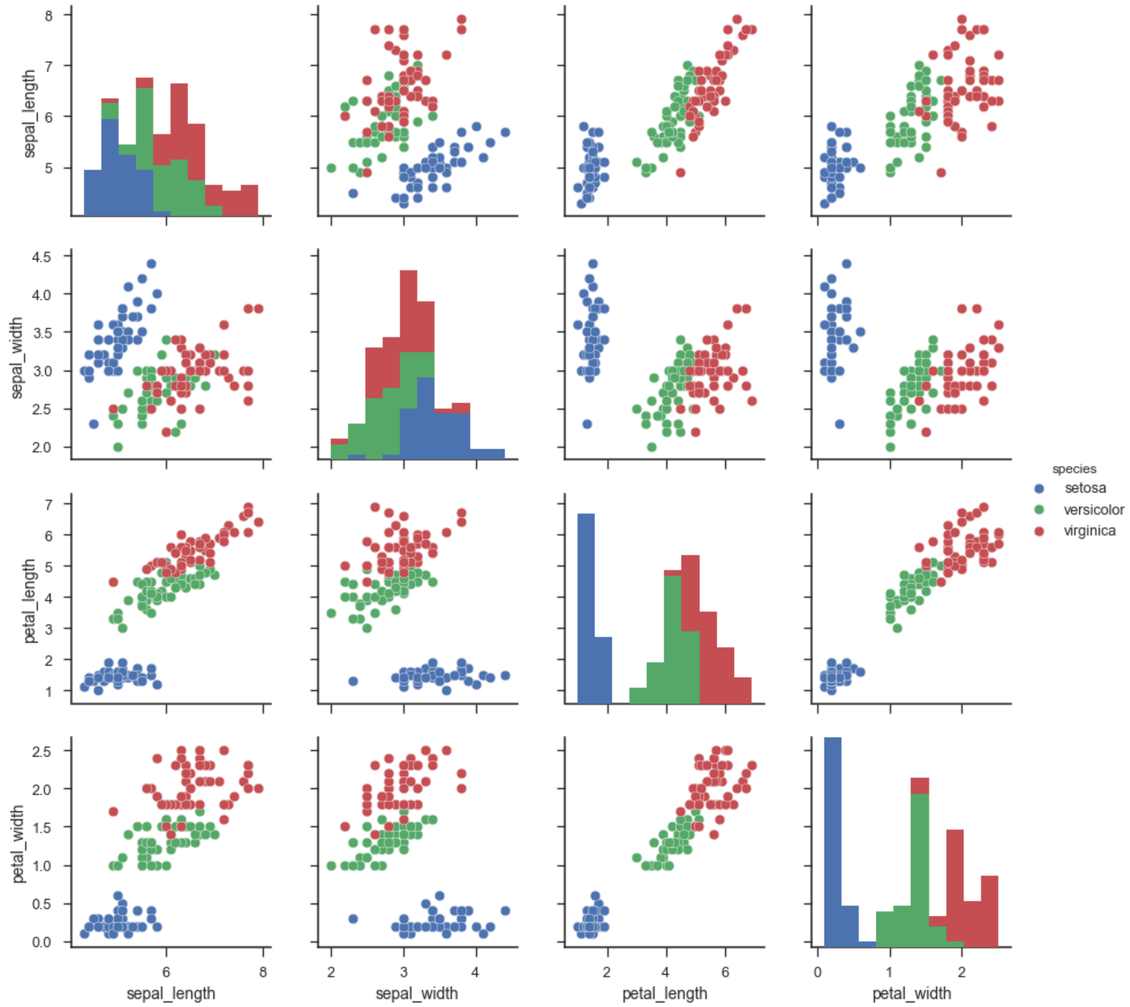

Is the data good enough?

What are the rules governing its use?

Do I have enough?

Do problems or biases exist in the data that could cause problems?

Feature Engineering

- Define what you are trying to measure. These will become the observations or rows of your final dataset
- Define how you will mathematically represent your data. This will be come the features or columns of your final dataset.

Feature Engineering

Feature	Value
Color	Gray
Fins	7
Predator	TRUE

Feature Engineering

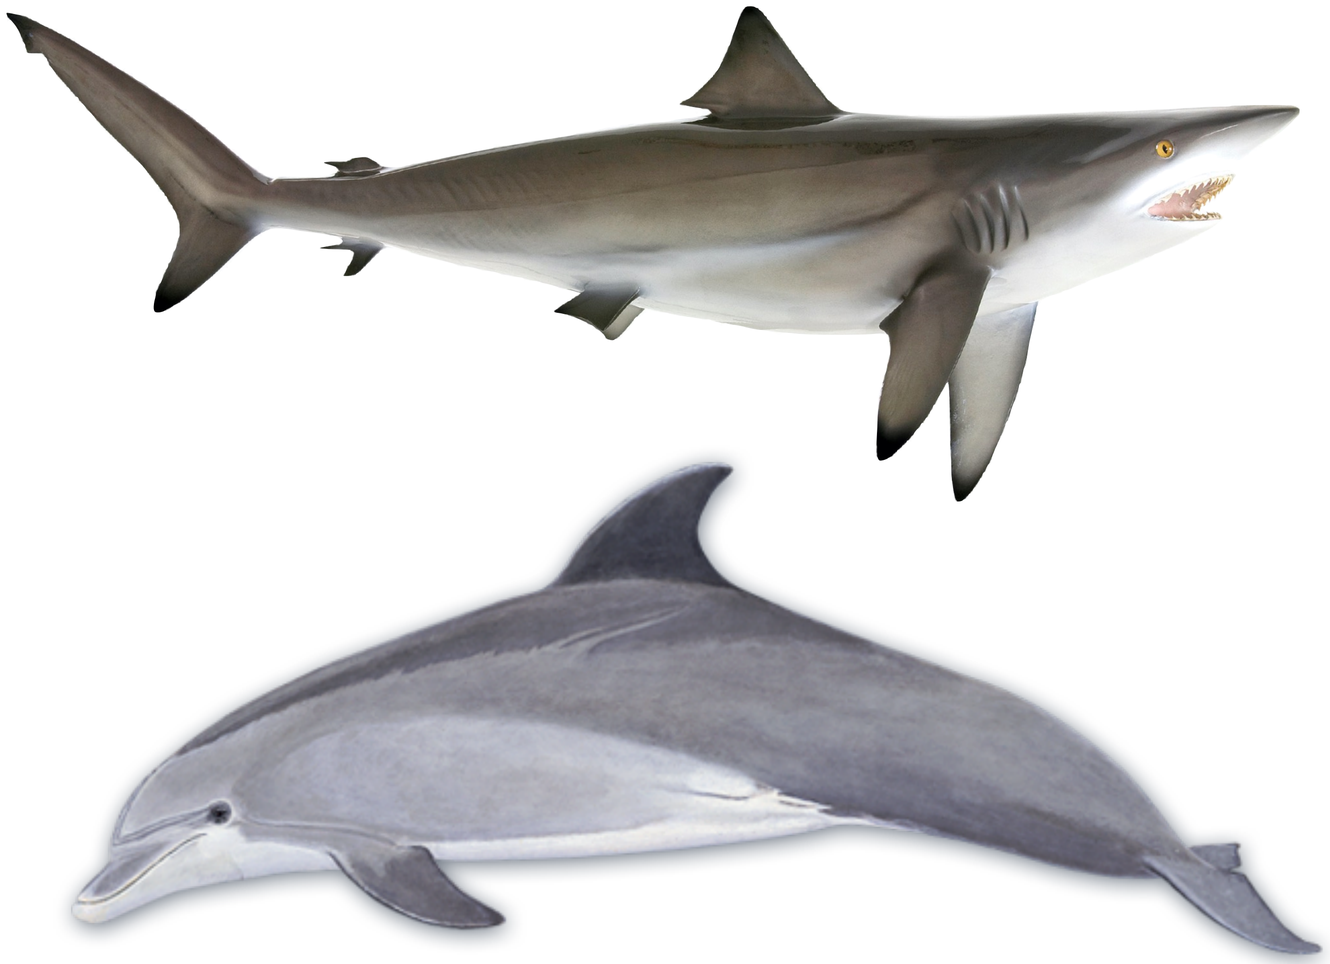

Feature	Value
Color	Gray
Fins	7
Predator	TRUE

Feature Engineering

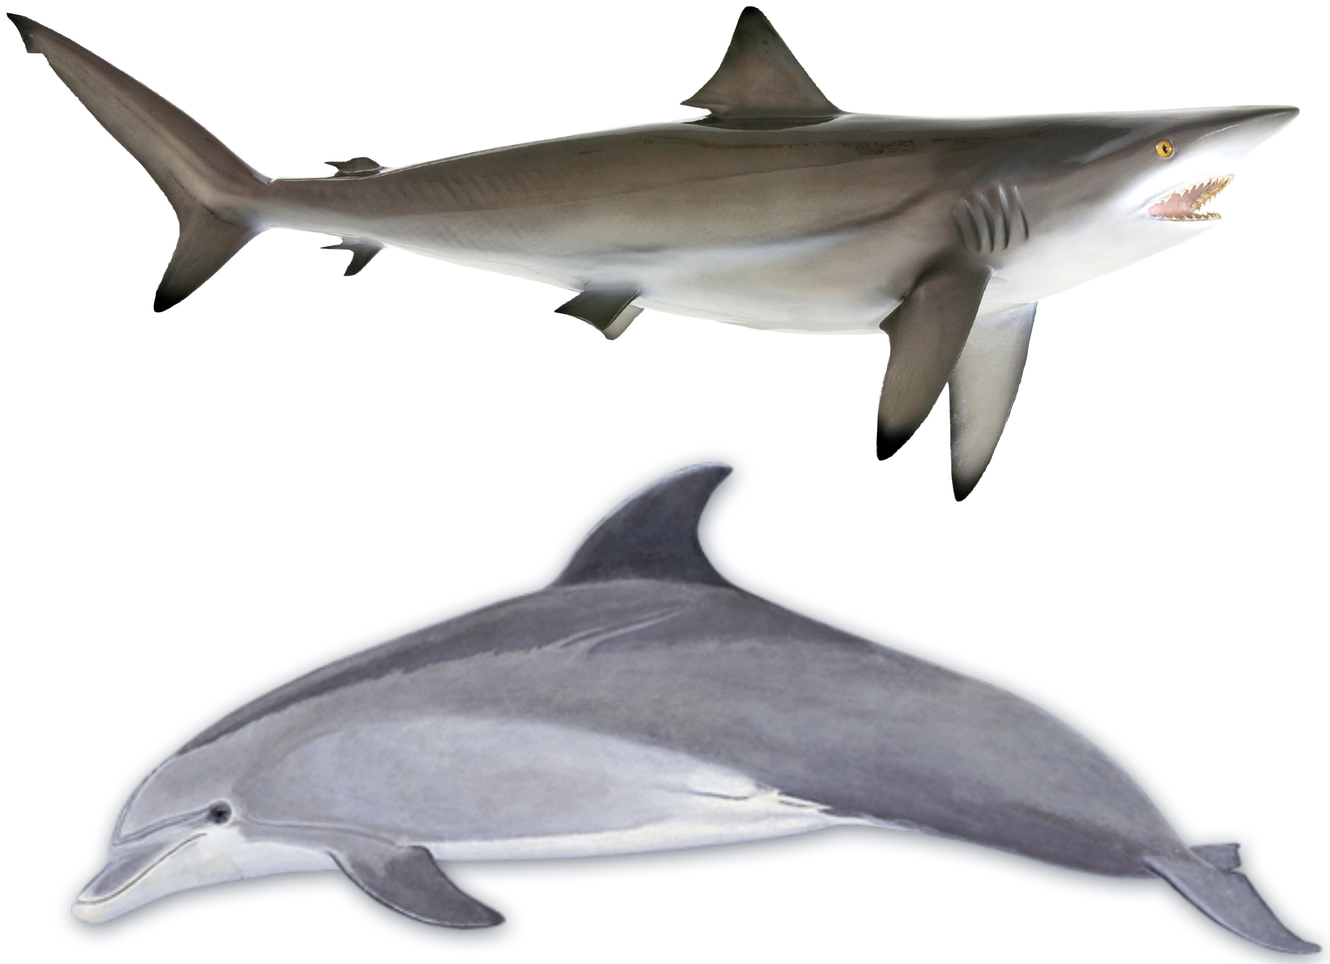

Feature	Value
Color	Gray
Fins	7
Predator	TRUE
Mammal	**TRUE**

Build and Tune your Model

- Believe it or not, this is the easy part.
- Most of this is **done using libraries** like scikit-learn, mllib, tensorflow, caret or keras, and **many steps can be automated**.
- You can even do it in Splunk or Elasticsearch.

Evaluate Performance

- Use various scoring methods, or write your own to determine model performance.
- Go back to step 1 and repeat! (Do not pass go, do not collect $200)

Group Discussion

Consider that you are building a system to identify fraudulent credit card transactions. In your groups, try to answer the following questions:

1. What are some features that you would want to capture?
2. What data sets will you need?
3. What legal and policy challenges might you face?
4. What other challenges you could foresee in this problem?
5. How will you define success?
6. How can you articulate the value of this model to stakeholders?

The Python Data Science Ecosystem

Machine Learning Ecosystem

- **Data Gathering:** Pandas, Drill, BeautifulSoup, PyDBAPI, PyDAL, Boto3
- **Feature Extraction:** Pandas, NumPy, Featuretools
- **Machine Learning**
  - **"Regular" ML:** Scikit-learn (sklearn), h2o, mllib (PySpark)
  - **Deep Learning:** Tensorflow, Keras, Theano, Caffe, PyTorch
- Visualization: Matplotlib, Seaborn, Yellowbrick, LIME, ggplot, plot.ly,

Recommended Reading

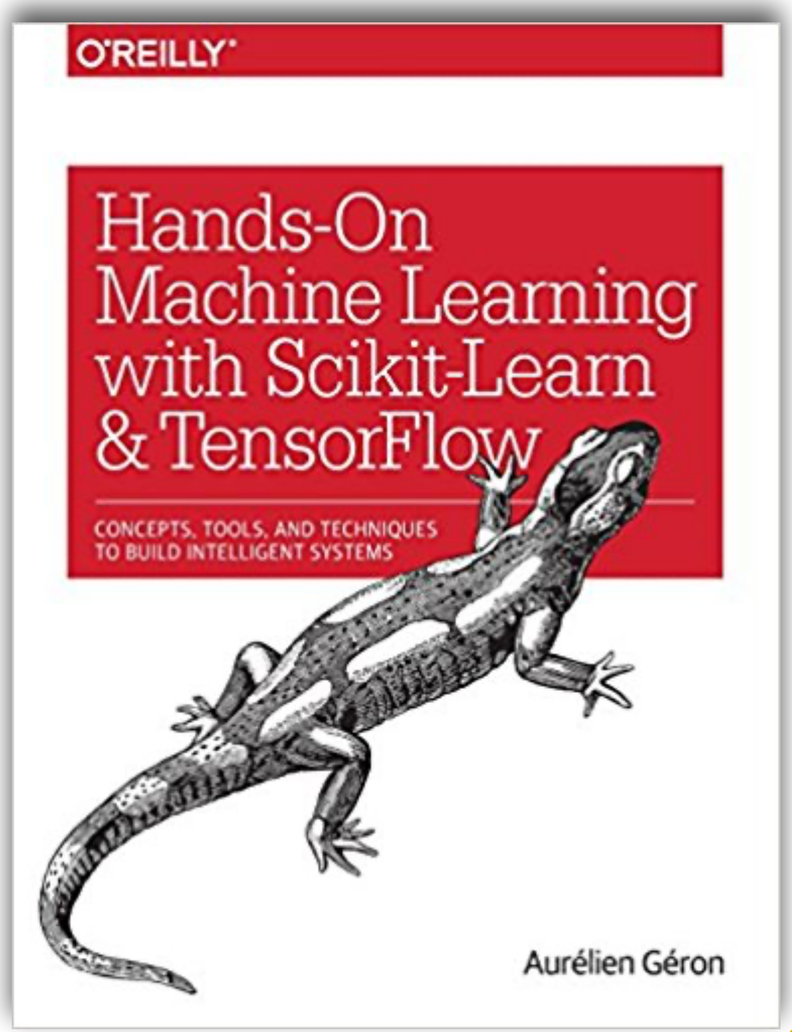

Recommended Reading

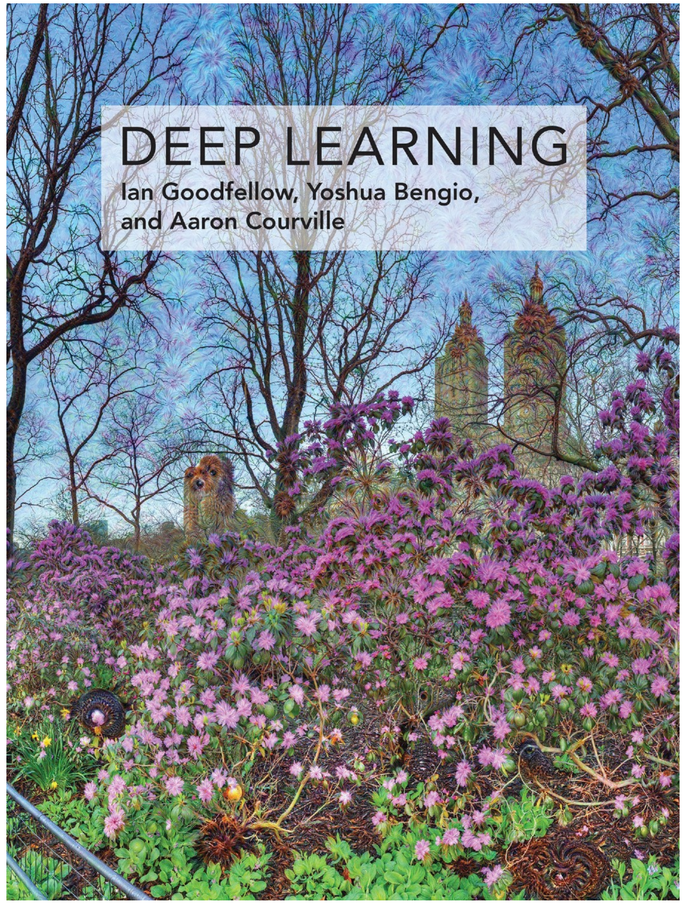

http://www.deeplearningbook.org/

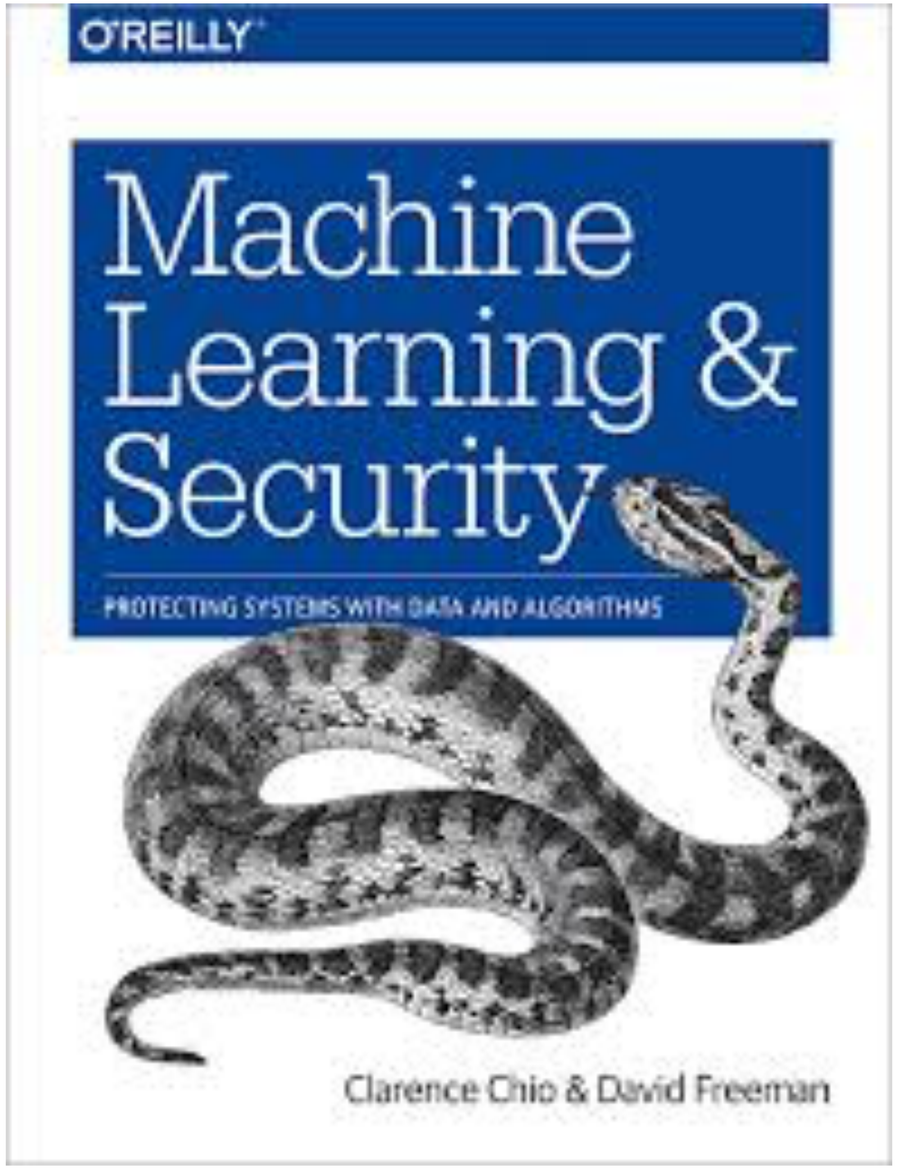
O'REILLY
Machine Learning & Security
PROTECTING SYSTEMS WITH DATA AND ALGORITHMS
Clarence Chio & David Freeman

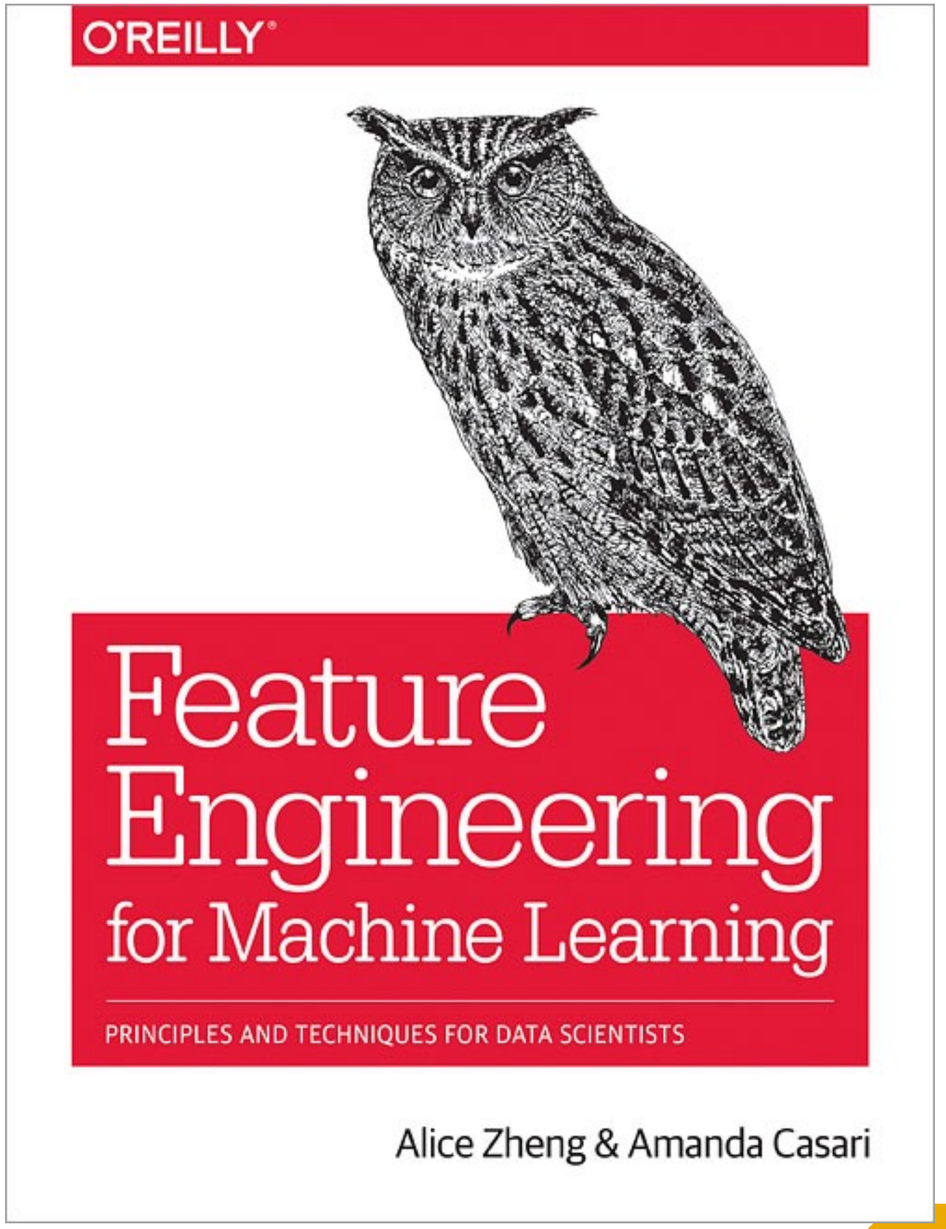
O'REILLY
Feature Engineering for Machine Learning
PRINCIPLES AND TECHNIQUES FOR DATA SCIENTISTS
Alice Zheng & Amanda Casari

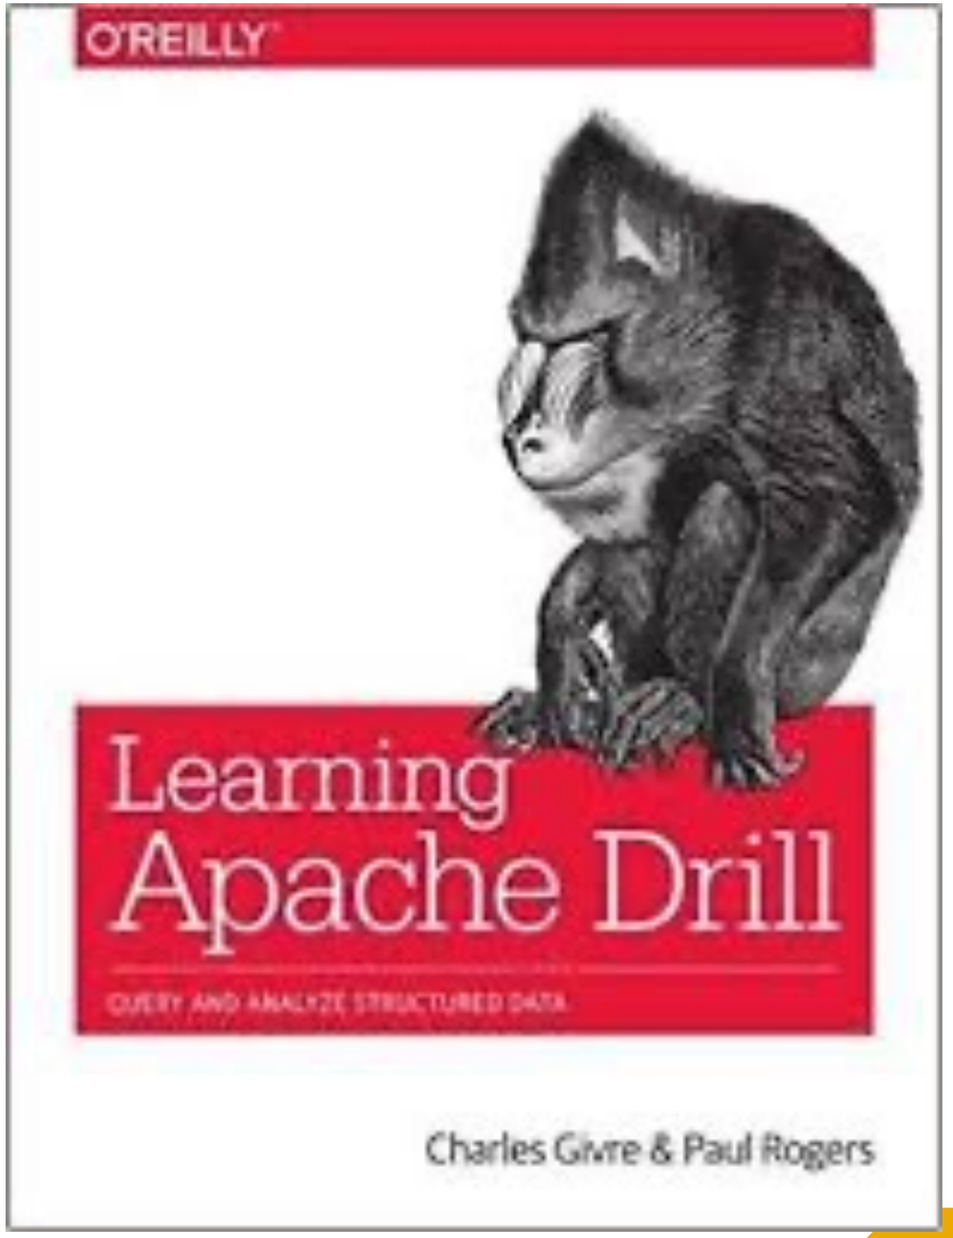
O'REILLY
Learning
Apache Drill
QUERY AND ANALYZE STRUCTURED DATA
Charles Givre & Paul Rogers

The Virtual Machine: Griffon

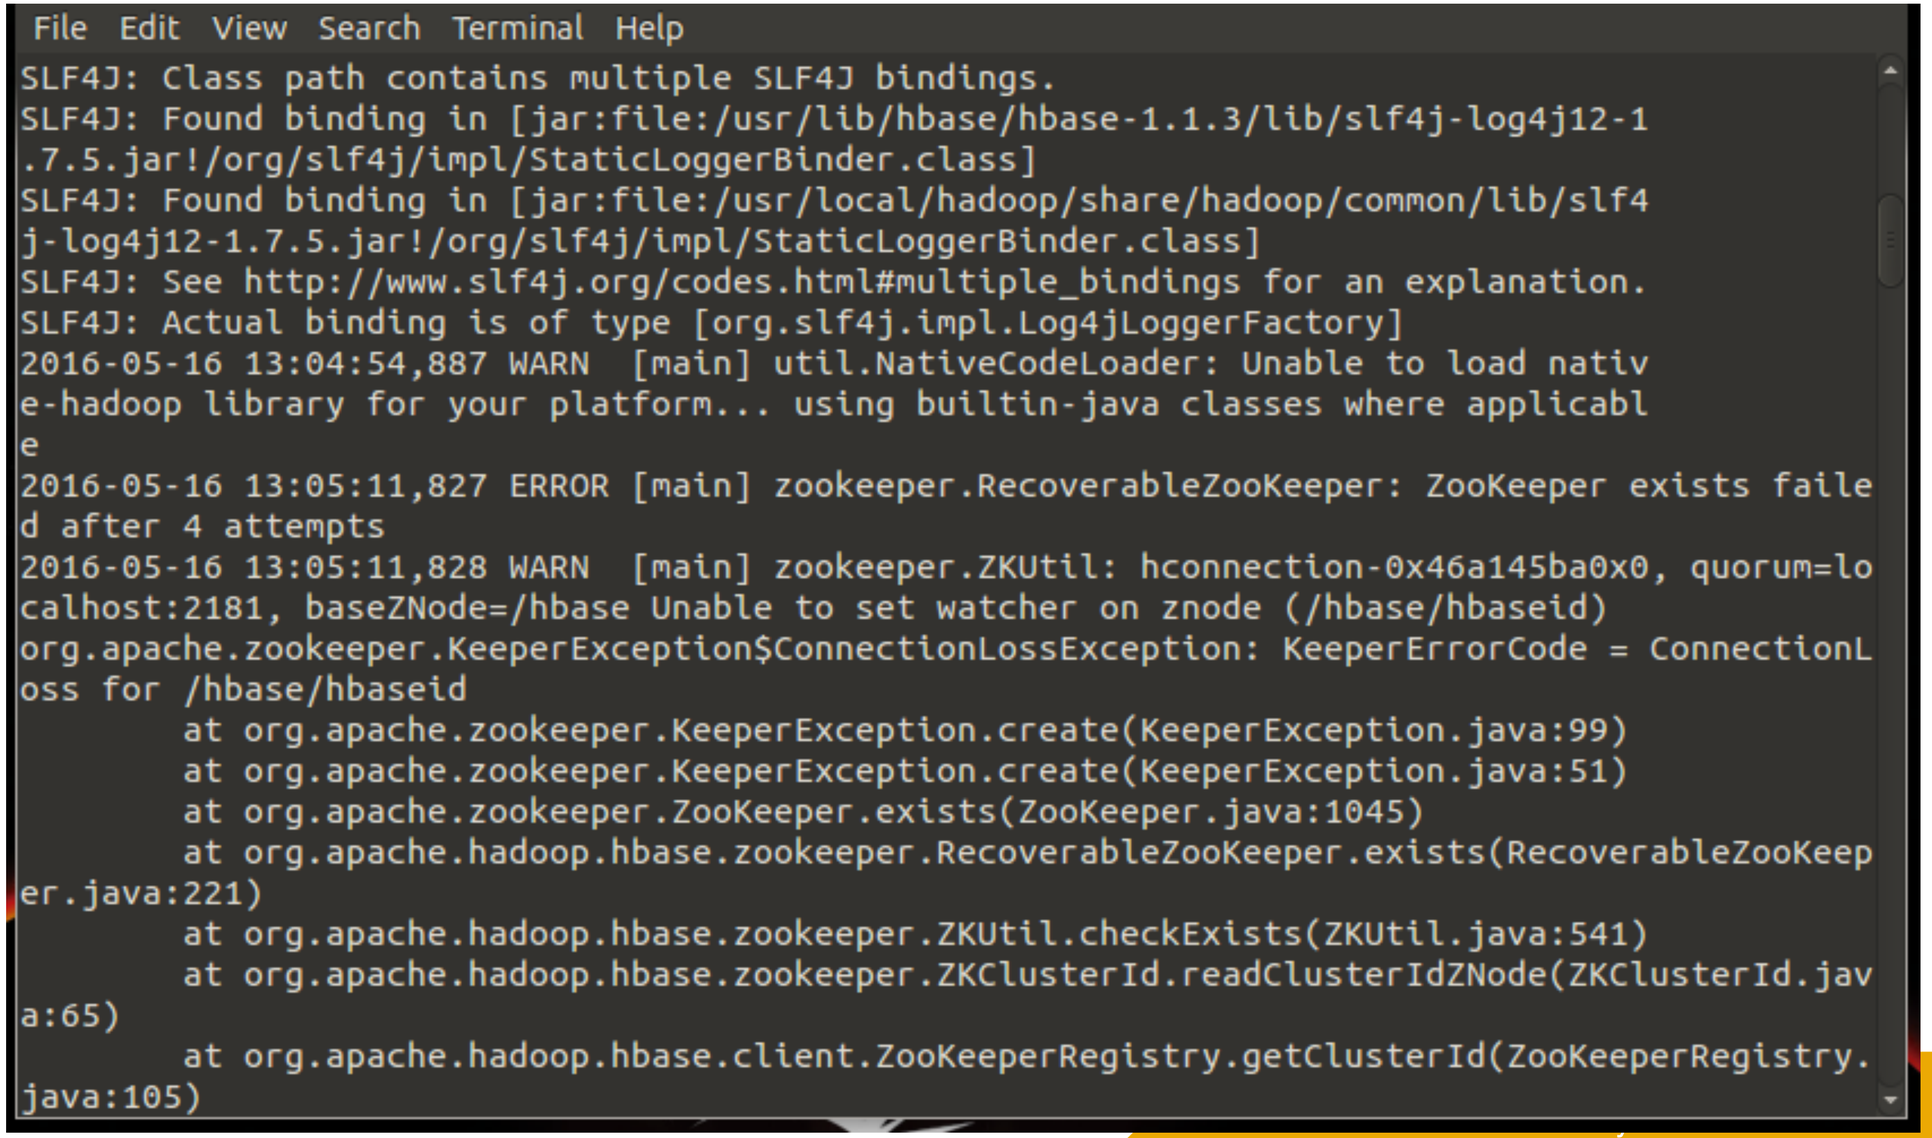
File Edit View Search Terminal Help
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/lib/hbase/hbase-1.1.3/lib/slf4j-log4j12-1
.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4
j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
2016-05-16 13:04:54,887 WARN [main] util.NativeCodeLoader: Unable to load nativ
e-hadoop library for your platform... using builtin-java classes where applicabl
e
2016-05-16 13:05:11,827 ERROR [main] zookeeper.RecoverableZooKeeper: ZooKeeper exists faile
d after 4 attempts
2016-05-16 13:05:11,828 WARN [main] zookeeper.ZKUtil: hconnection-0x46a145ba0x0, quorum=lo
calhost:2181, baseZNode=/hbase Unable to set watcher on znode (/hbase/hbaseid)
org.apache.zookeeper.KeeperException$ConnectionLossException: KeeperErrorCode = ConnectionL
oss for /hbase/hbaseid
at org.apache.zookeeper.KeeperException.create(KeeperException.java:99)
at org.apache.zookeeper.KeeperException.create(KeeperException.java:51)
at org.apache.zookeeper.ZooKeeper.exists(ZooKeeper.java:1045)
at org.apache.hadoop.hbase.zookeeper.RecoverableZooKeeper.exists(RecoverableZooKeep
er.java:221)
at org.apache.hadoop.hbase.zookeeper.ZKUtil.checkExists(ZKUtil.java:541)
at org.apache.hadoop.hbase.zookeeper.ZKClusterId.readClusterIdZNode(ZKClusterId.jav
a:65)
at org.apache.hadoop.hbase.client.ZooKeeperRegistry.getClusterId(ZooKeeperRegistry.
java:105)

Do Data Science, Not Sysadmin

Built on Ubuntu MATE

Easy to Use

Programming Languages

- Languages
- Libraries
- Editors and Notebooks
- Databases + Administrative tools
- Big Data Tools
- Machine Learning Libraries
- Data Visualization

Scripting Languages


Griffon RC3 [Running]
Applications
Places
System
Accessories
Big Data
Databases
Data Visualization
Education
Graphics
Internet
Office
Other
Programming
Sound & Video
System Tools
Universal Access
passwords.txt
Shells
Atom
DBeaver CE
H2O.ai
IntelliJ IDEA Community Edition
Jupyter Notebook
PyCharm Community Edition
Rodeo
RStudio
Spyder
Julia
NodeJS
PHP Shell
Pig Shell
Python (v2.7)
Python (v3.5)
R
Ruby
Scala

Editors & Notebooks

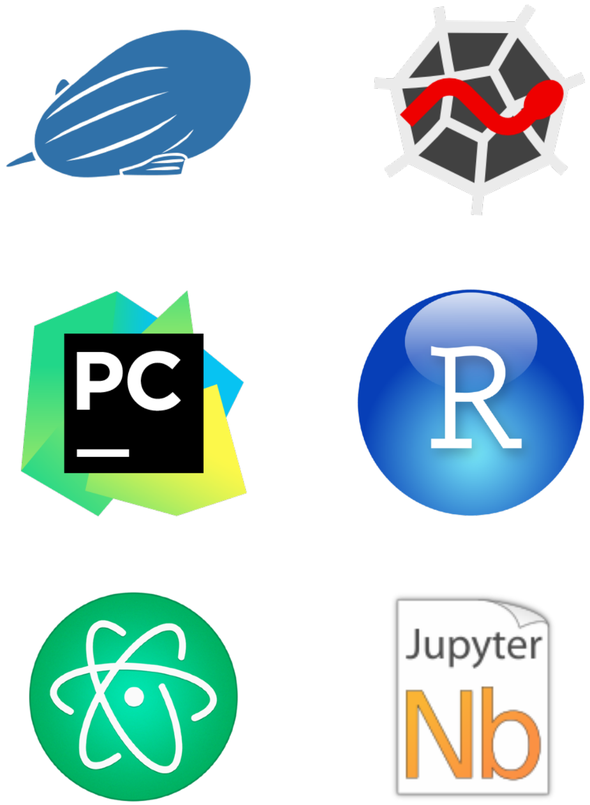
Jupyter
Nb

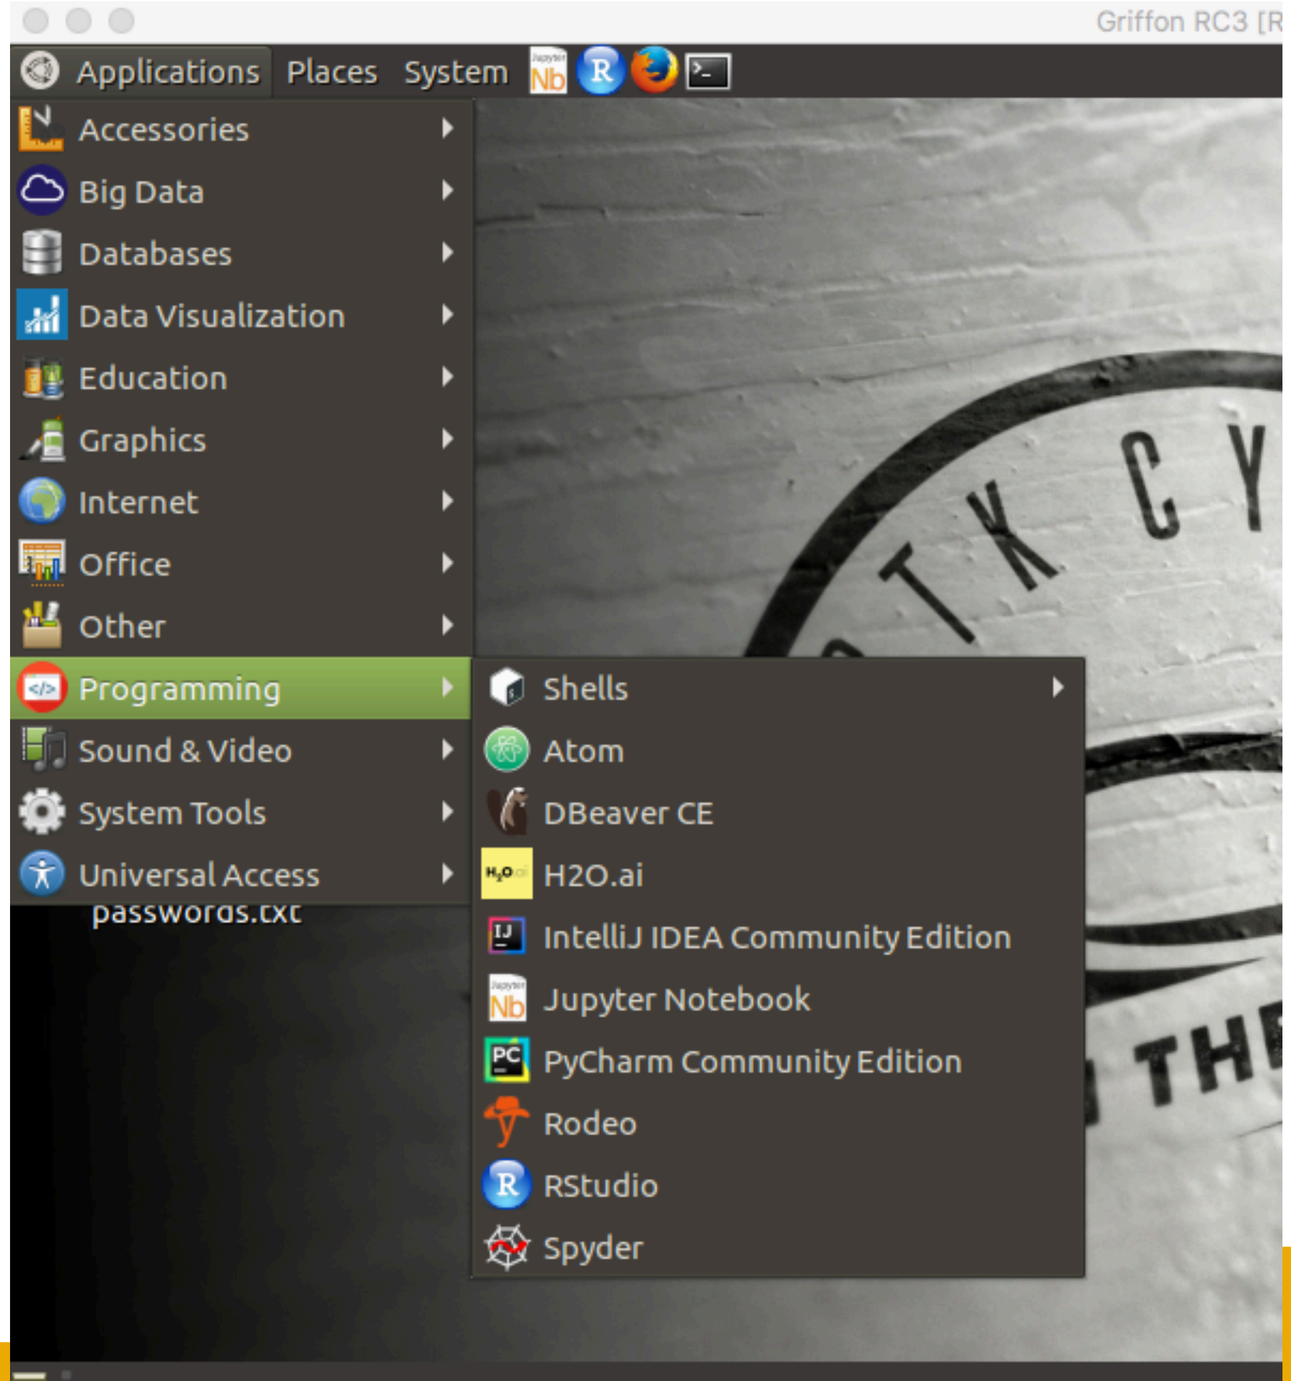
Griffon RC3 [R
Applications
Places
System
Accessories
Big Data
Databases
Data Visualization
Education
Graphics
Internet
Office
Other
Programming
Sound & Video
System Tools
Universal Access
passwords.txt
Shells
Atom
DBeaver CE
H2O.ai
IntelliJ IDEA Community Edition
Jupyter Notebook
PyCharm Community Edition
Rodeo
RStudio
Spyder

Atom

Project Merlin Beta 3 (Small updates) [Running]

Applications Places System Mon Aug 15, 12:26

python_demo.py — /usr/share/atom/resources — Atom

File Edit View Selection Find Packages Help

resources

python_demo.py

```
1  import pandas as pd
2  df = pd.DataFrame( [2,3,4,5,6,7,8])
```


Jupyter
Nb

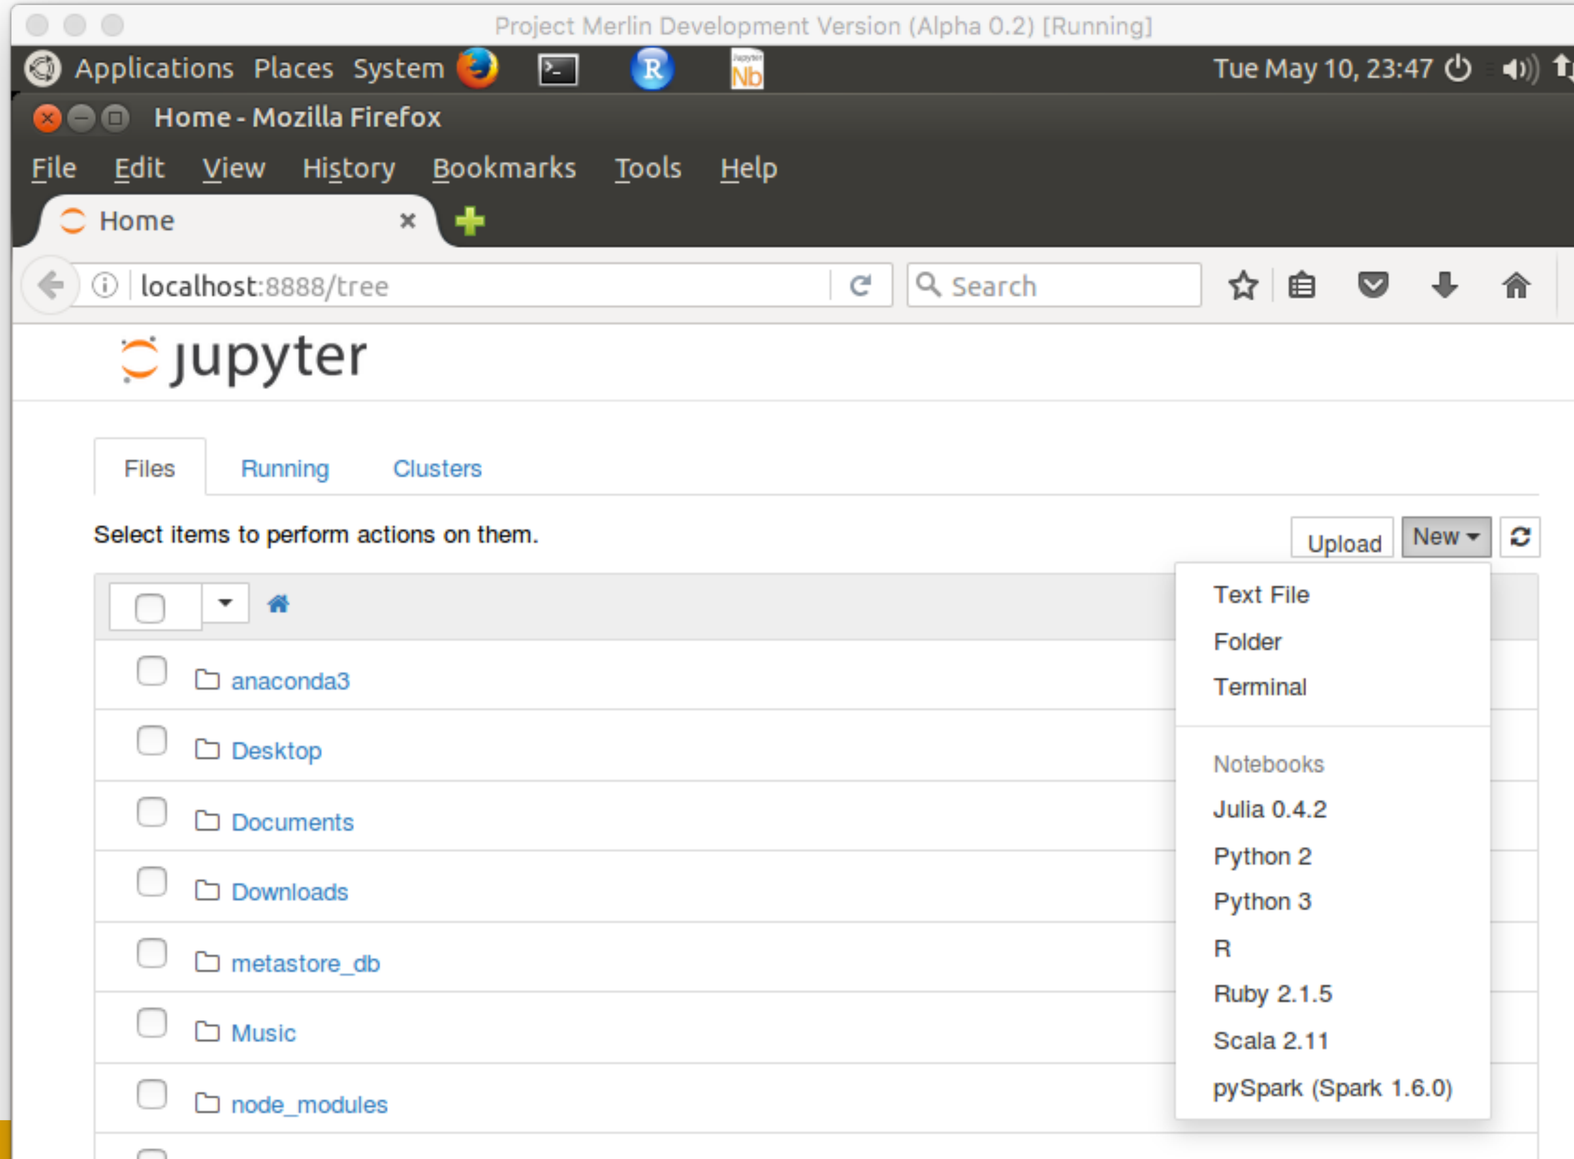
Project Merlin Development Version (Alpha 0.2) [Running]
Applications Places System
Tue May 10, 23:47
Home - Mozilla Firefox
File Edit View History Bookmarks Tools Help
Home
localhost:8888/tree
Search
jupyter
Files Running Clusters
Select items to perform actions on them.
Upload
New
Text File
Folder
Terminal
Notebooks
Julia 0.4.2
Python 2
Python 3
R
Ruby 2.1.5
Scala 2.11
pySpark (Spark 1.6.0)
anaconda3
Desktop
Documents
Downloads
metastore_db
Music
node_modules

Jupyter Notebook

- Python 2 & 3
- R
- Ruby
- Scala
- PySpark


MySQL®


mongoDB®


PostgreSQL


neo4j


SQLite

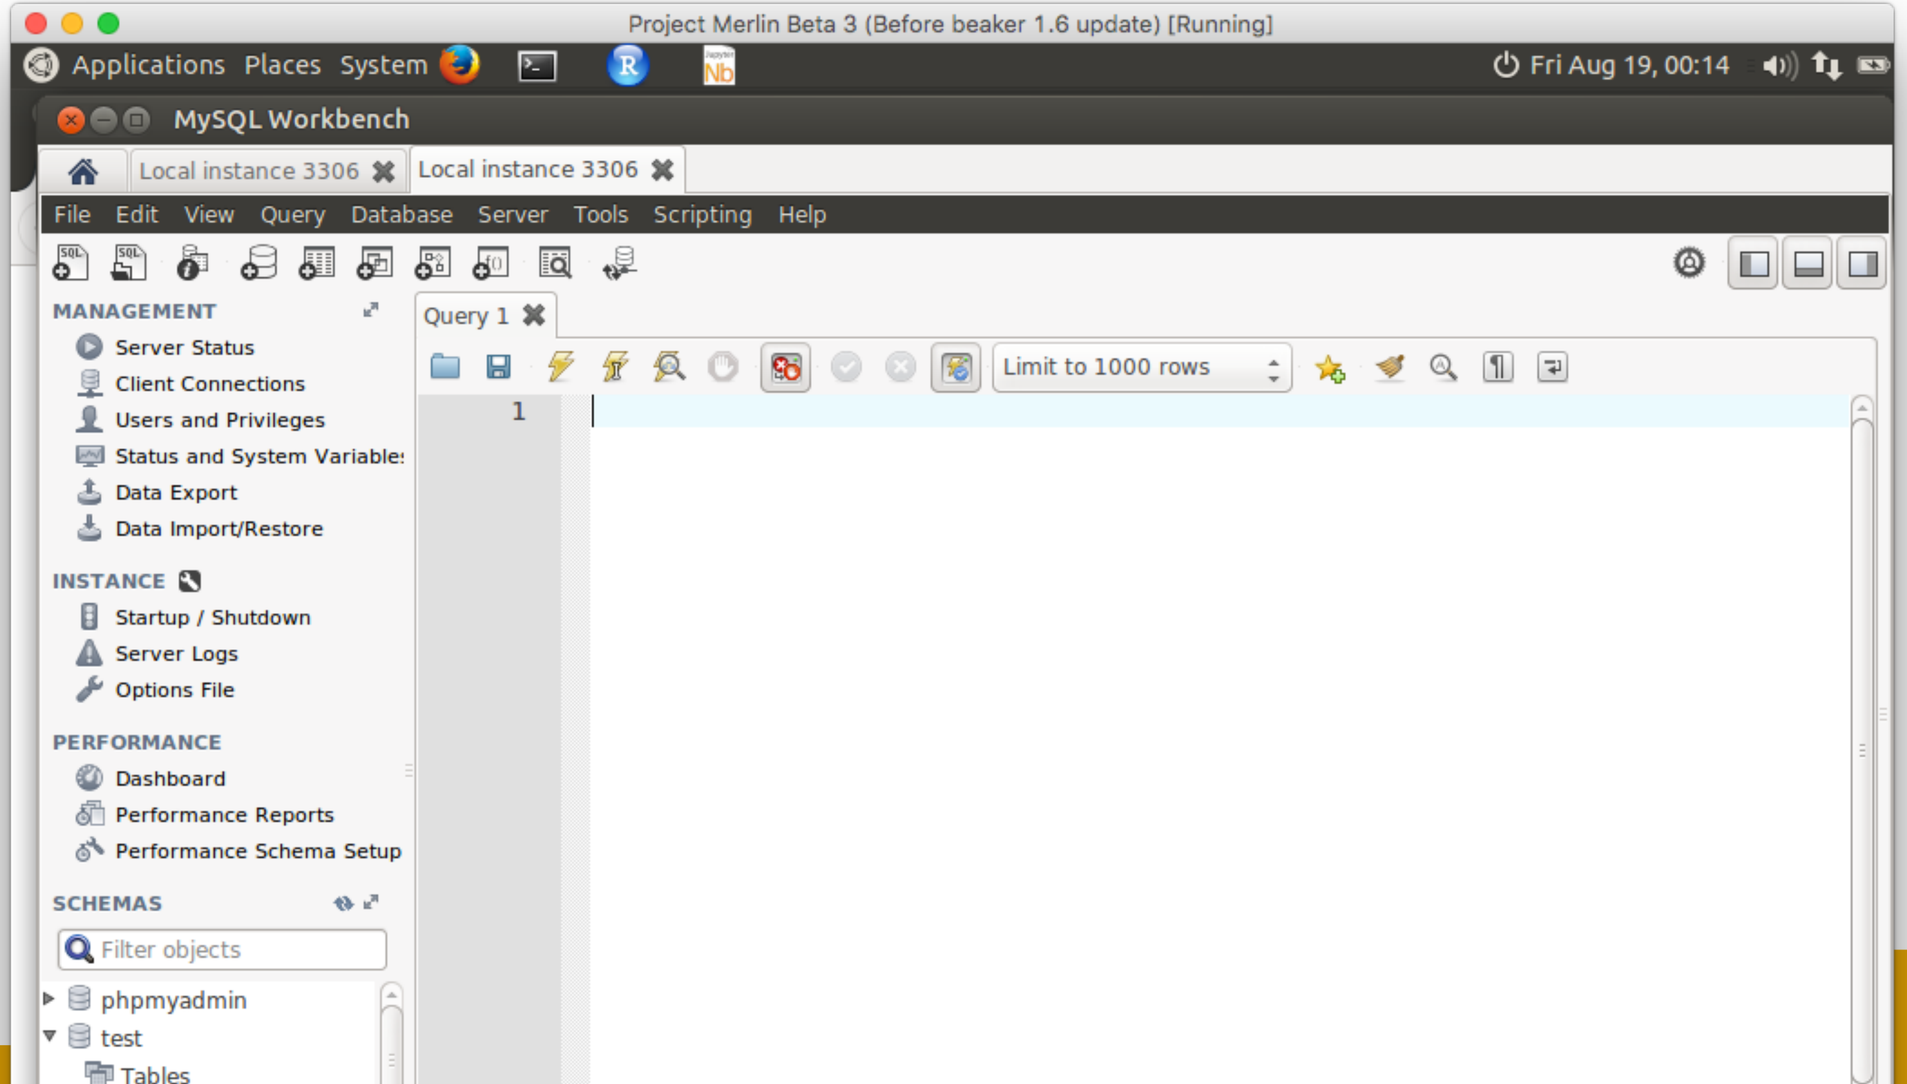
Project Merlin Beta 3 (Before beaker 1.6 update) [Running]
Applications Places System
Fri Aug 19, 00:14
MySQL Workbench
Local instance 3306
Local instance 3306
File Edit View Query Database Server Tools Scripting Help
MANAGEMENT
Server Status
Client Connections
Users and Privileges
Status and System Variables
Data Export
Data Import/Restore
INSTANCE
Startup / Shutdown
Server Logs
Options File
PERFORMANCE
Dashboard
Performance Reports
Performance Schema Setup
SCHEMAS
Filter objects
phpmyadmin
test
Tables
Query 1
Limit to 1000 rows
1

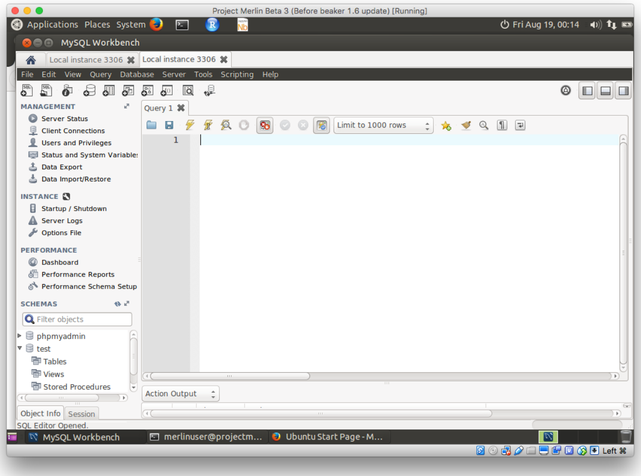

MySQL®

Project Merlin Beta 3 (Before beaker 1.6 update) [Running]

Applications Places System

Fri Aug 19, 00:30

localhost / localhost | phpMyAdmin 4.4.13.1deb1 - Mozilla Firefox

localhost / localhost...

localhost/phpmyadmin/index.php?token=62fb6ffb630fba16f23788ee21dc4e

Search

phpMyAdmin

Recent Favorites

- New
- information_schema
- mysql
- performance_schema
- phpmyadmin
- test

Server: localhost

Databases | SQL | Status | Users | Export | Import | Settings | More

General Settings

Change password

Server connection collation: utf8mb4_unicode_ci

Appearance Settings

Language: English

Theme: pmahomme

- Font size: 82%

More settings

Database server

- Server: Localhost via UNIX socket
- Server type: MySQL
- Server version: 5.6.31-0ubuntu0.15.10.1 - (Ubuntu)
- Protocol version: 10
- User: merlinuser@localhost
- Server charset: UTF-8 Unicode (utf8)

Web server

- nginx/1.10.1
- Database client version: libmysql - 5.6.31
- PHP extension: mysqli
- PHP version: 5.6.11-1ubuntu3.4

phpMyAdmin


mongoDB.

mongoDB.

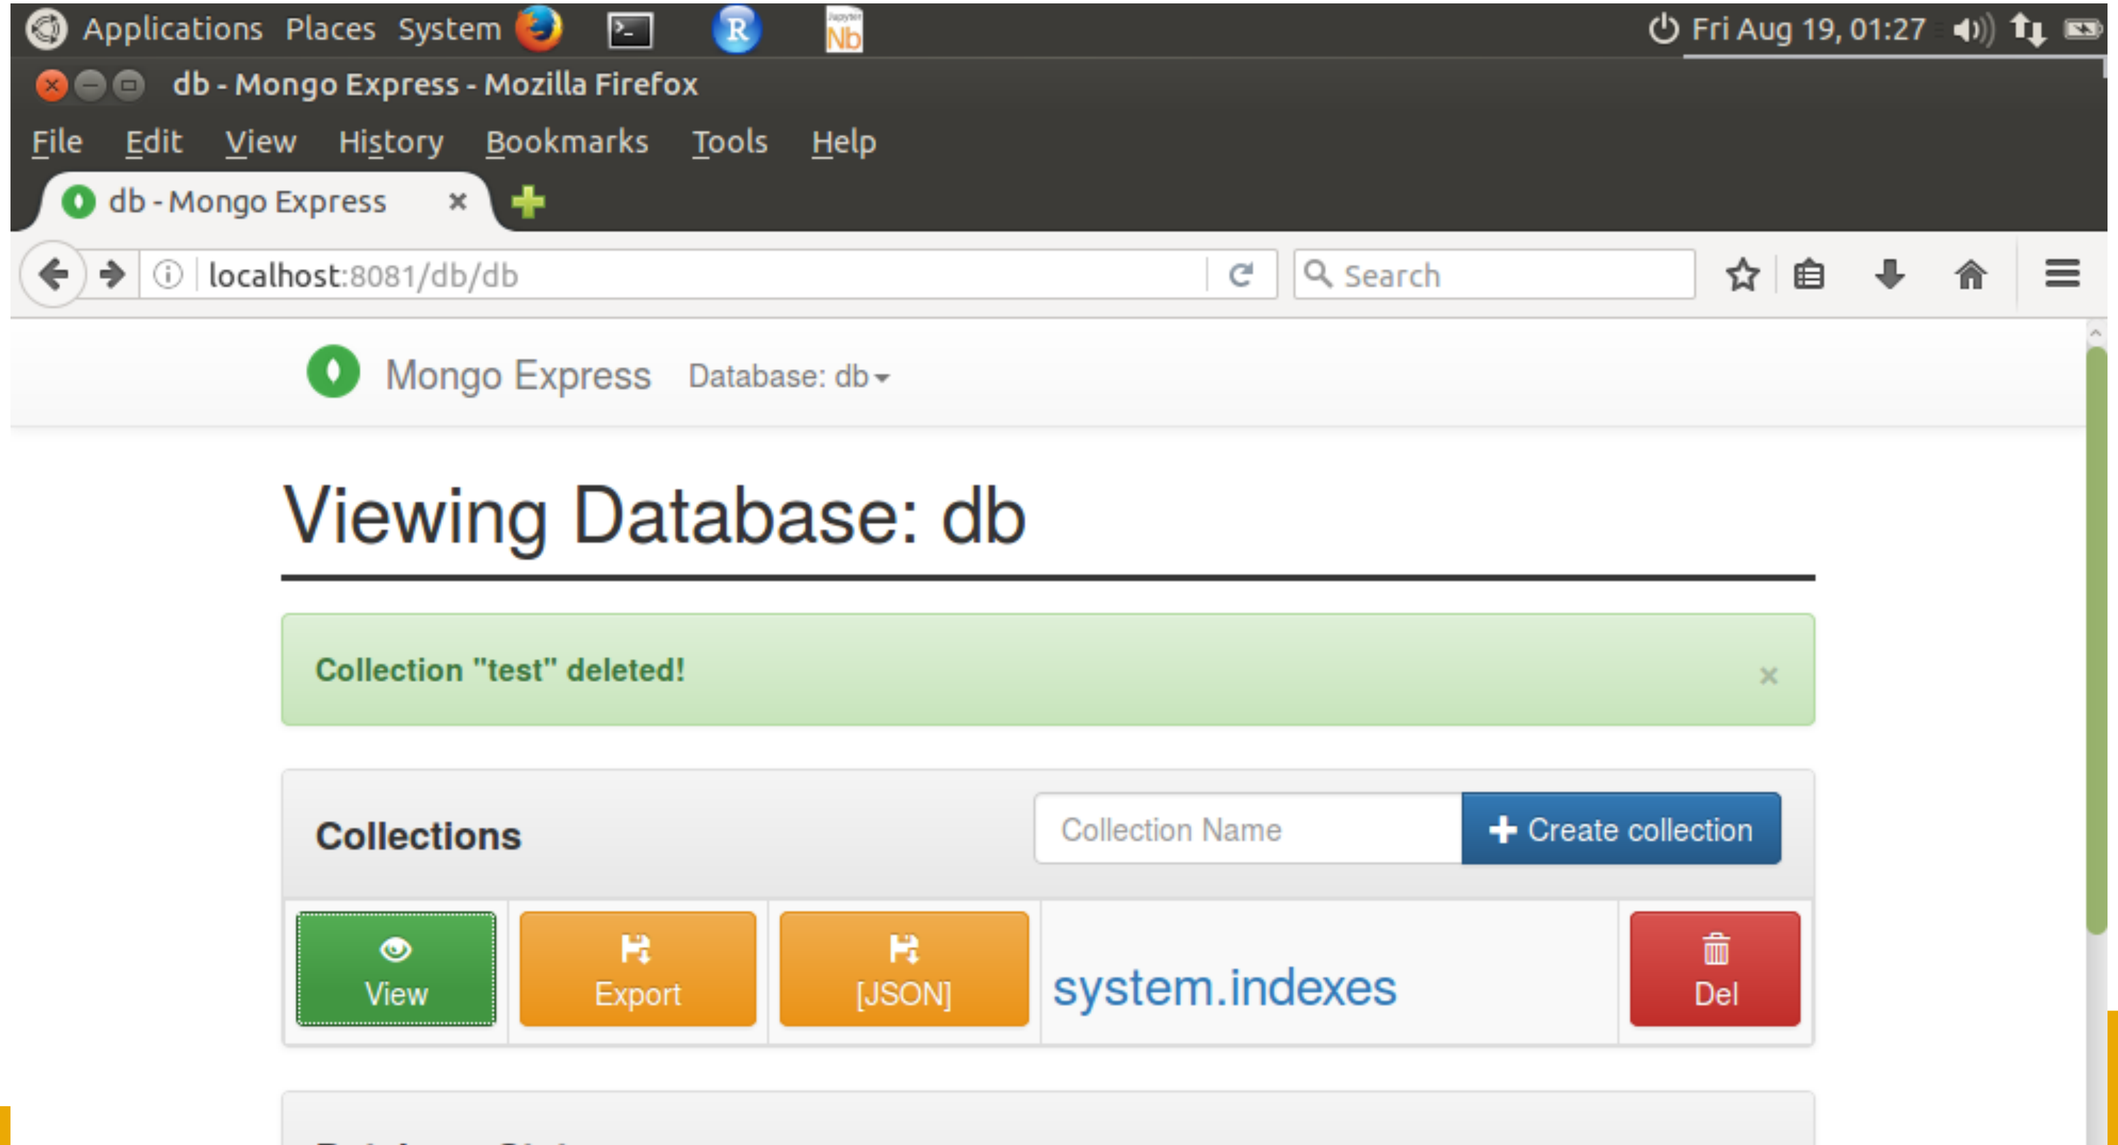
Applications Places System
Fri Aug 19, 01:27
db - Mongo Express - Mozilla Firefox
File Edit View History Bookmarks Tools Help
db - Mongo Express
localhost:8081/db/db
Search
Mongo Express
Database: db
Viewing Database: db
Collection "test" deleted!
Collections
Collection Name
Create collection
View
Export
[JSON]
system.indexes
Del

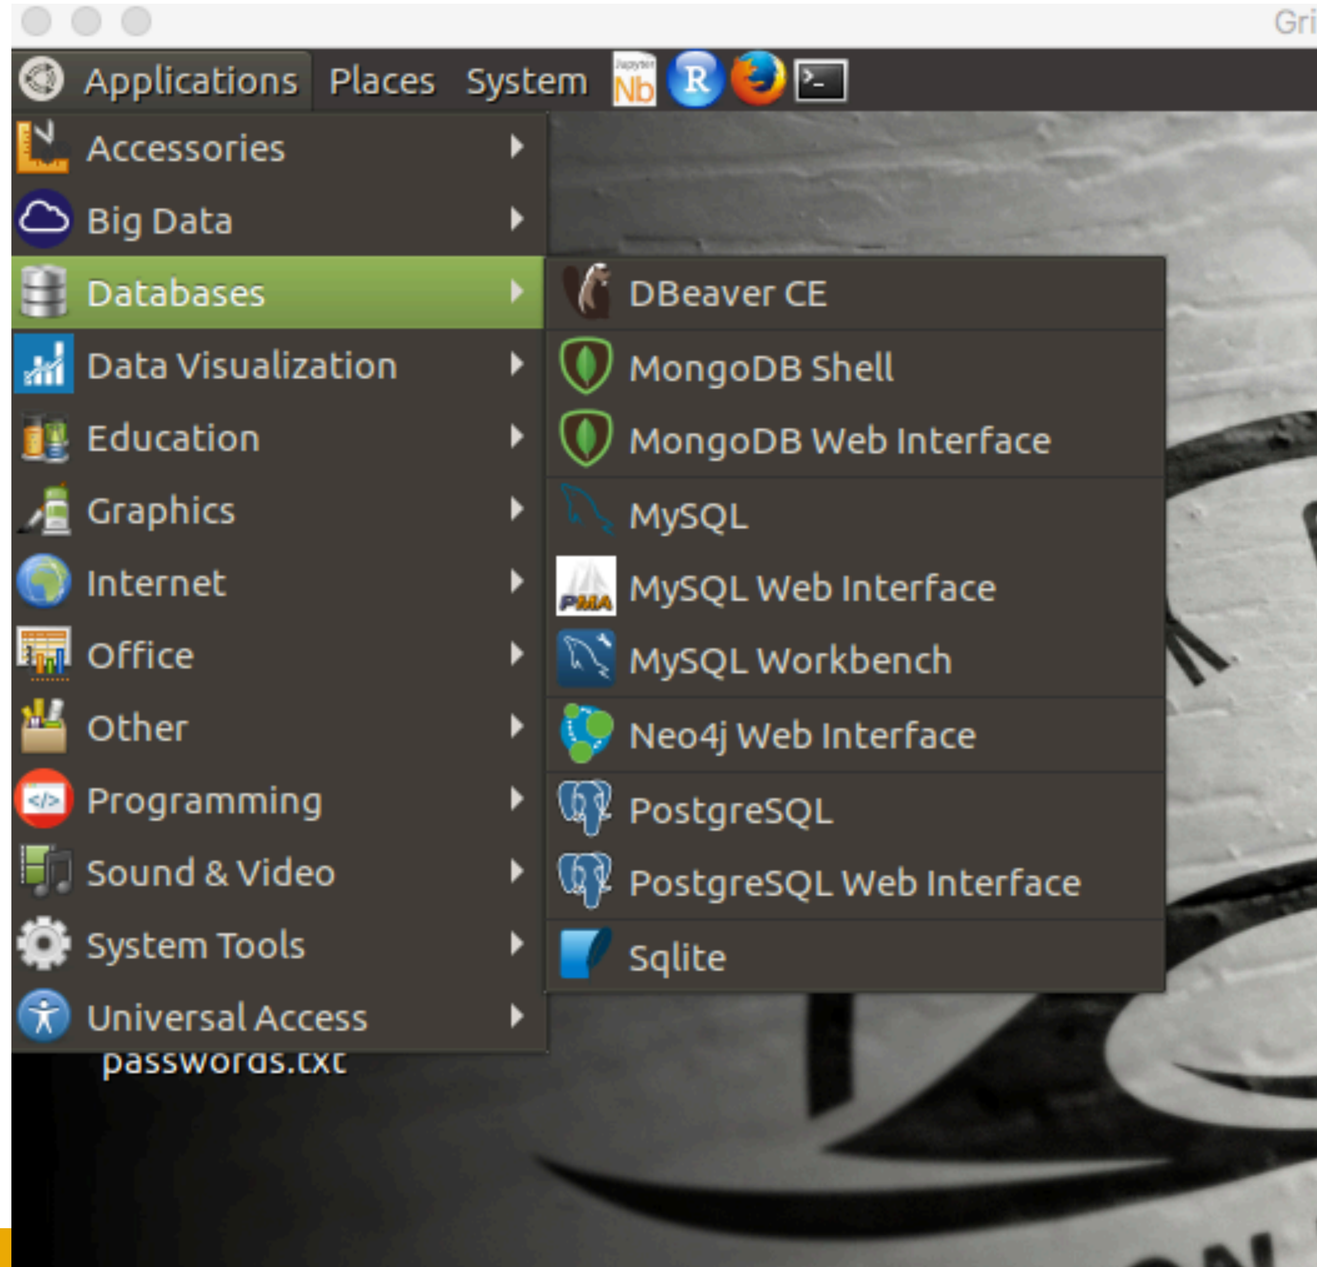
Applications
Places
System
Accessories
Big Data
Databases
Data Visualization
Education
Graphics
Internet
Office
Other
Programming
Sound & Video
System Tools
Universal Access
DBeaver CE
MongoDB Shell
MongoDB Web Interface
MySQL
MySQL Web Interface
MySQL Workbench
Neo4j Web Interface
PostgreSQL
PostgreSQL Web Interface
Sqlite
passwords.txt


projectmerlin: /usr/local/hadoop/sbin
Edit View
Terminal Help
:/usr/local/hadoop/sbin$ ./start-dfs.sh
16/05/11 23:23:39 WA
il.NativeCodeLoader: Unable to load native-hadoop libra
ry for your platform..
ng builtin-java classes where applicable
Starting namenodes on [lo
ost]
merlinuser@localhost's pas
merlinuser@localhost's passw
localhost: Permission denied, please try again.
merlinuser@localhost's password:
lhost: Permission denied, please try again.
localhost: Permission denied (publi
assword).
merlinuser@localhost's password:
merlinuser@localhost's password: localho
ermission denied, please try again.

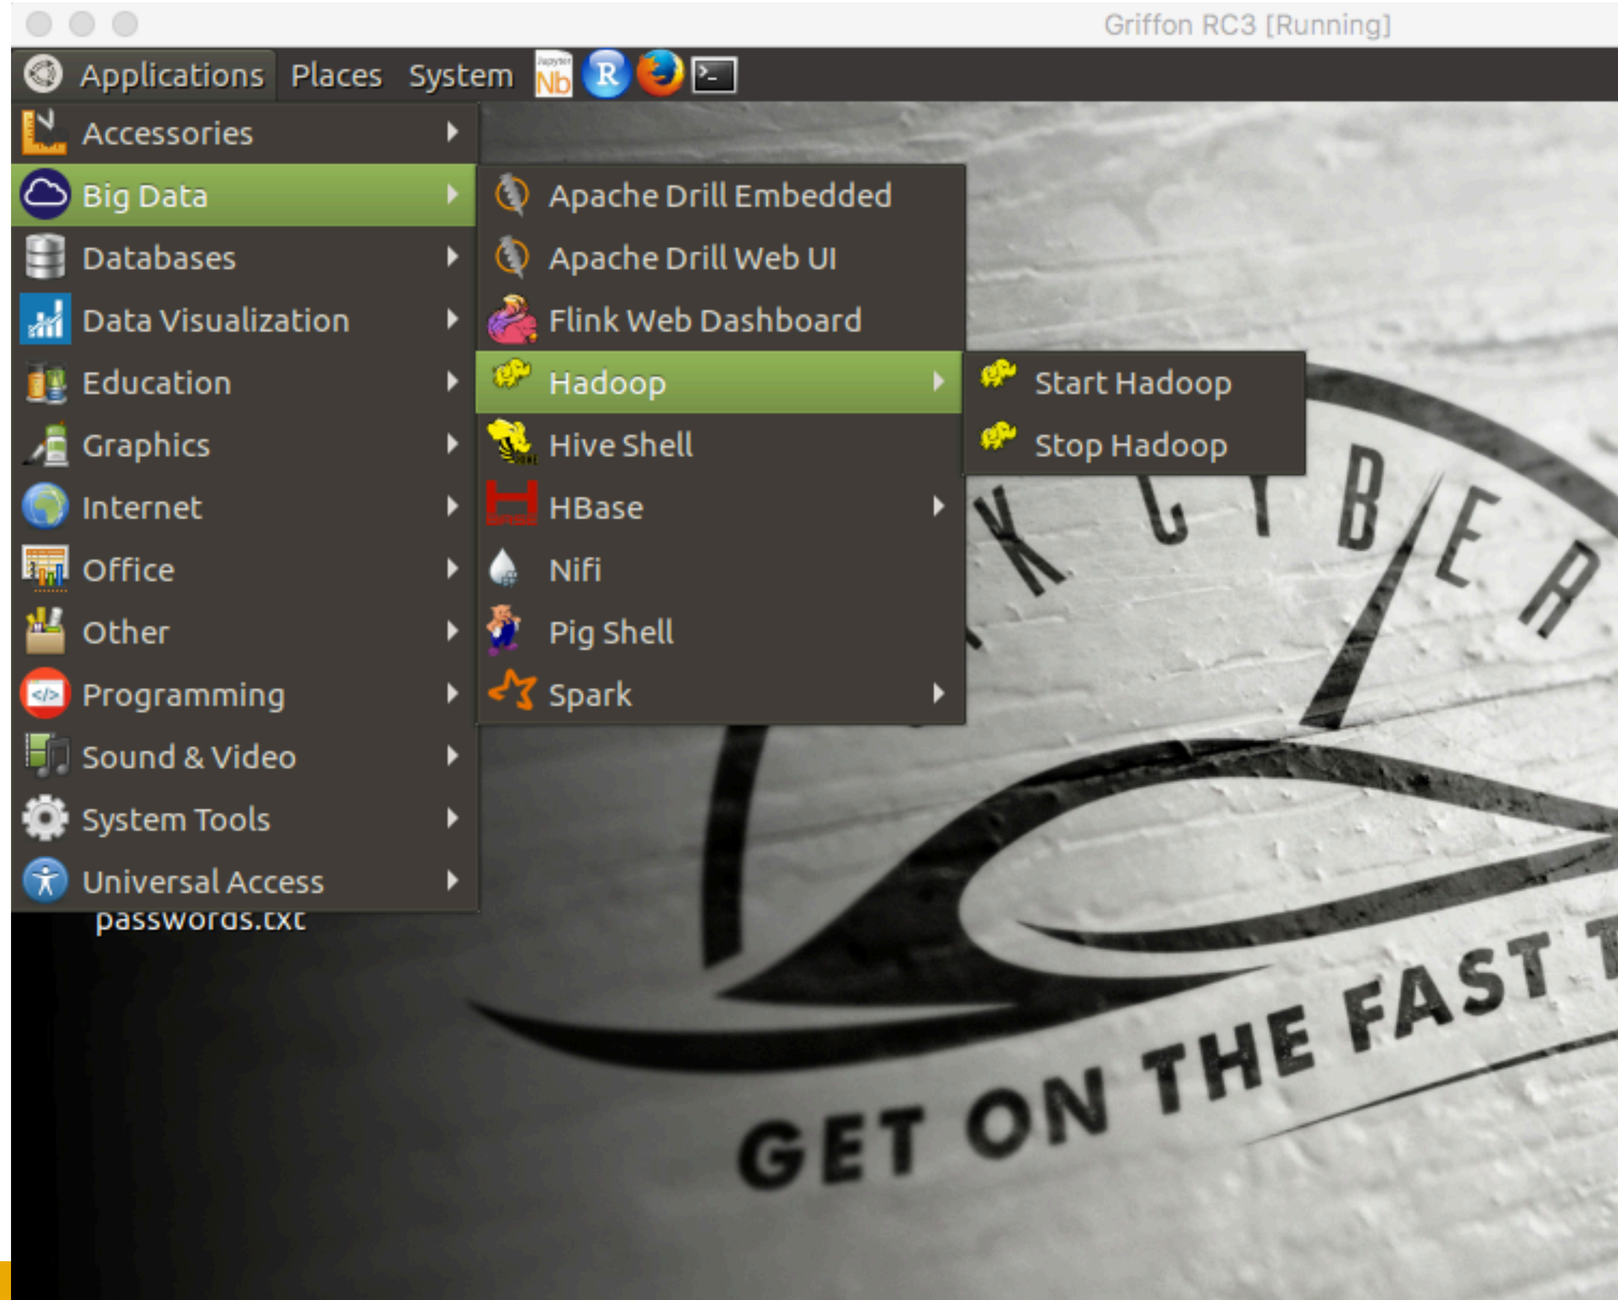
Griffon RC3 [Running]
Applications
Places
System
Accessories
Big Data
Databases
Data Visualization
Education
Graphics
Internet
Office
Other
Programming
Sound & Video
System Tools
Universal Access
Apache Drill Embedded
Apache Drill Web UI
Flink Web Dashboard
Hadoop
Hive Shell
HBase
Nifi
Pig Shell
Spark
Start Hadoop
Stop Hadoop
passwords.txt
GET ON THE FAST


plotly


orange
DATA MINING
FRUITFUL&FUN

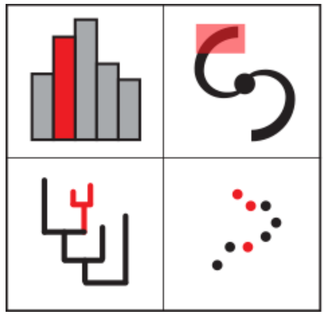


orange
DATA MINING
FRUITFUL&FUN

orange

DATA MINING
FRUITFUL&FUN

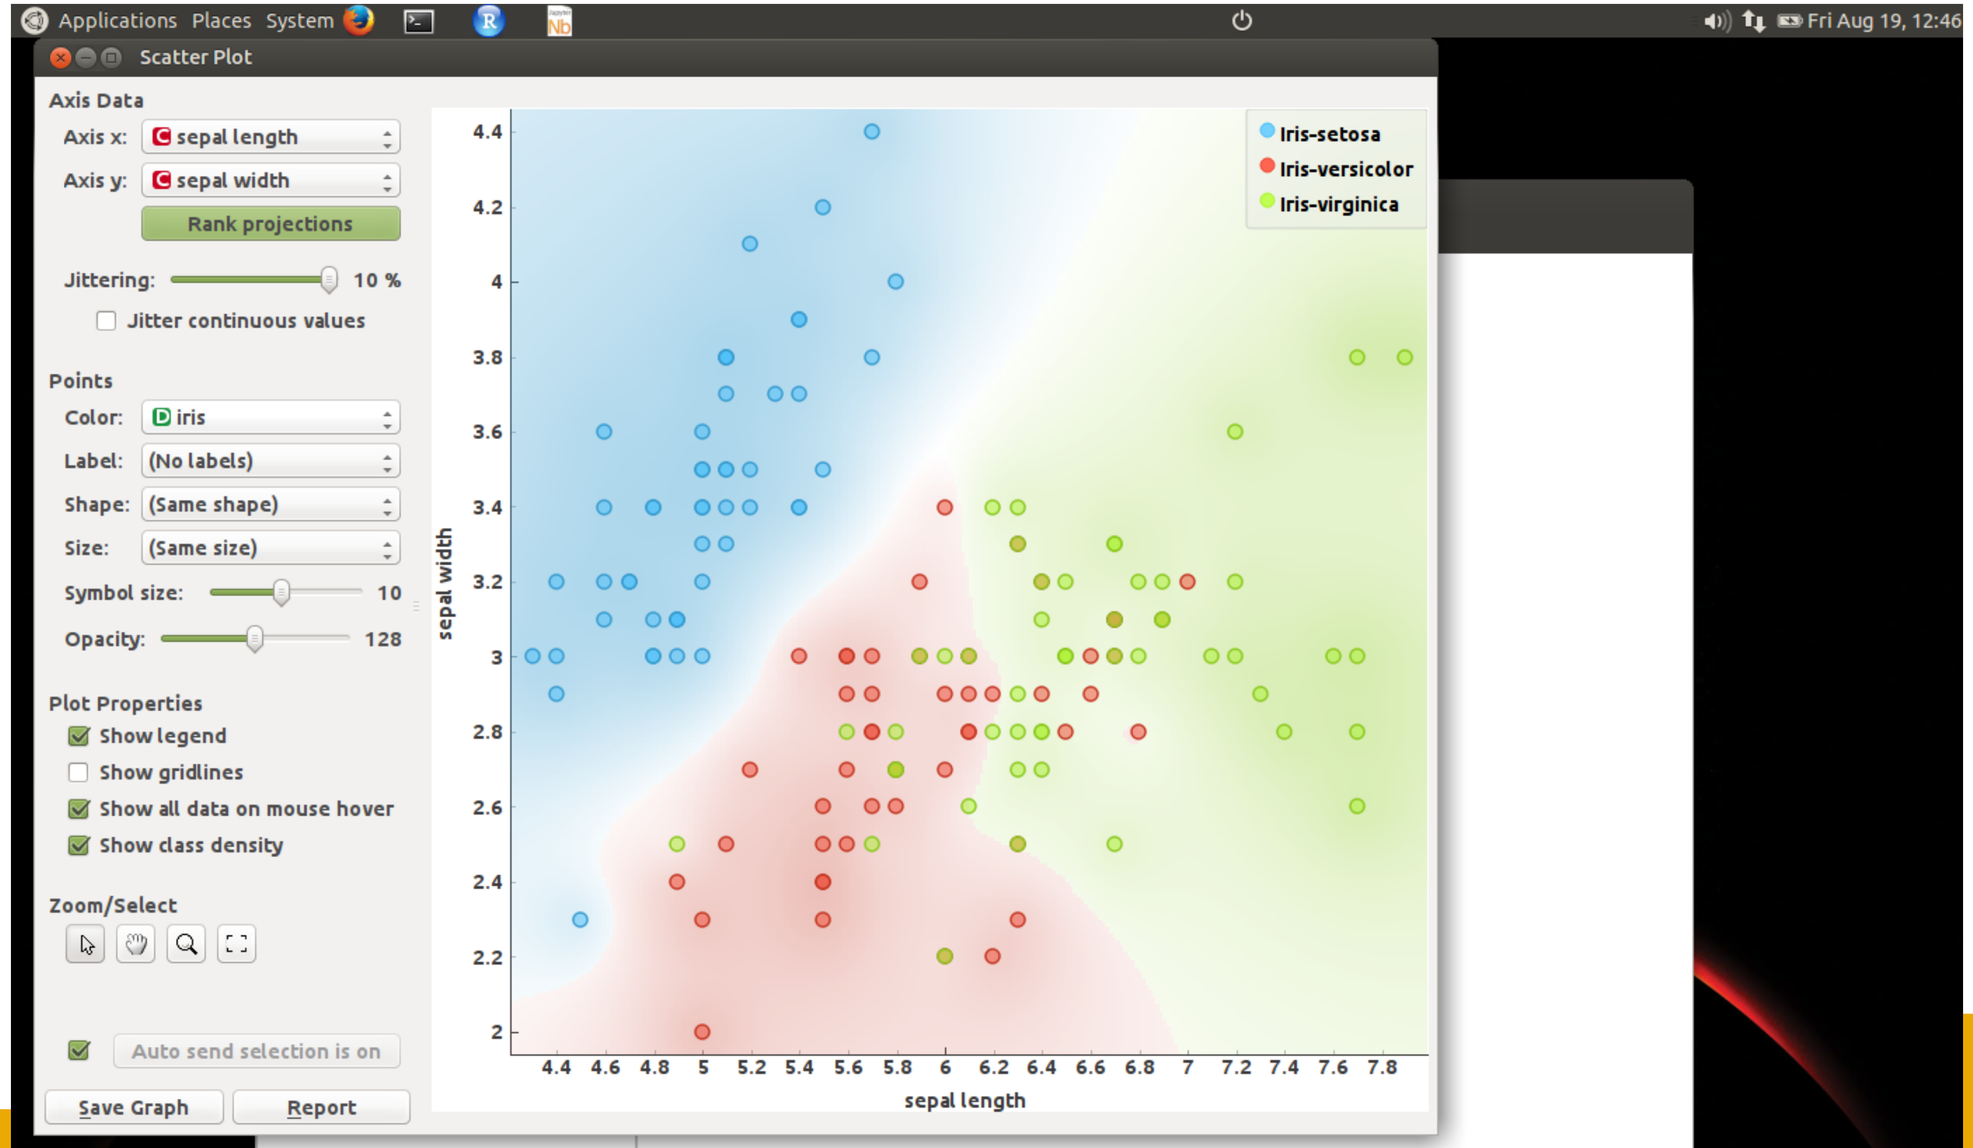

orange

DATA MINING
FRUITFUL&FUN

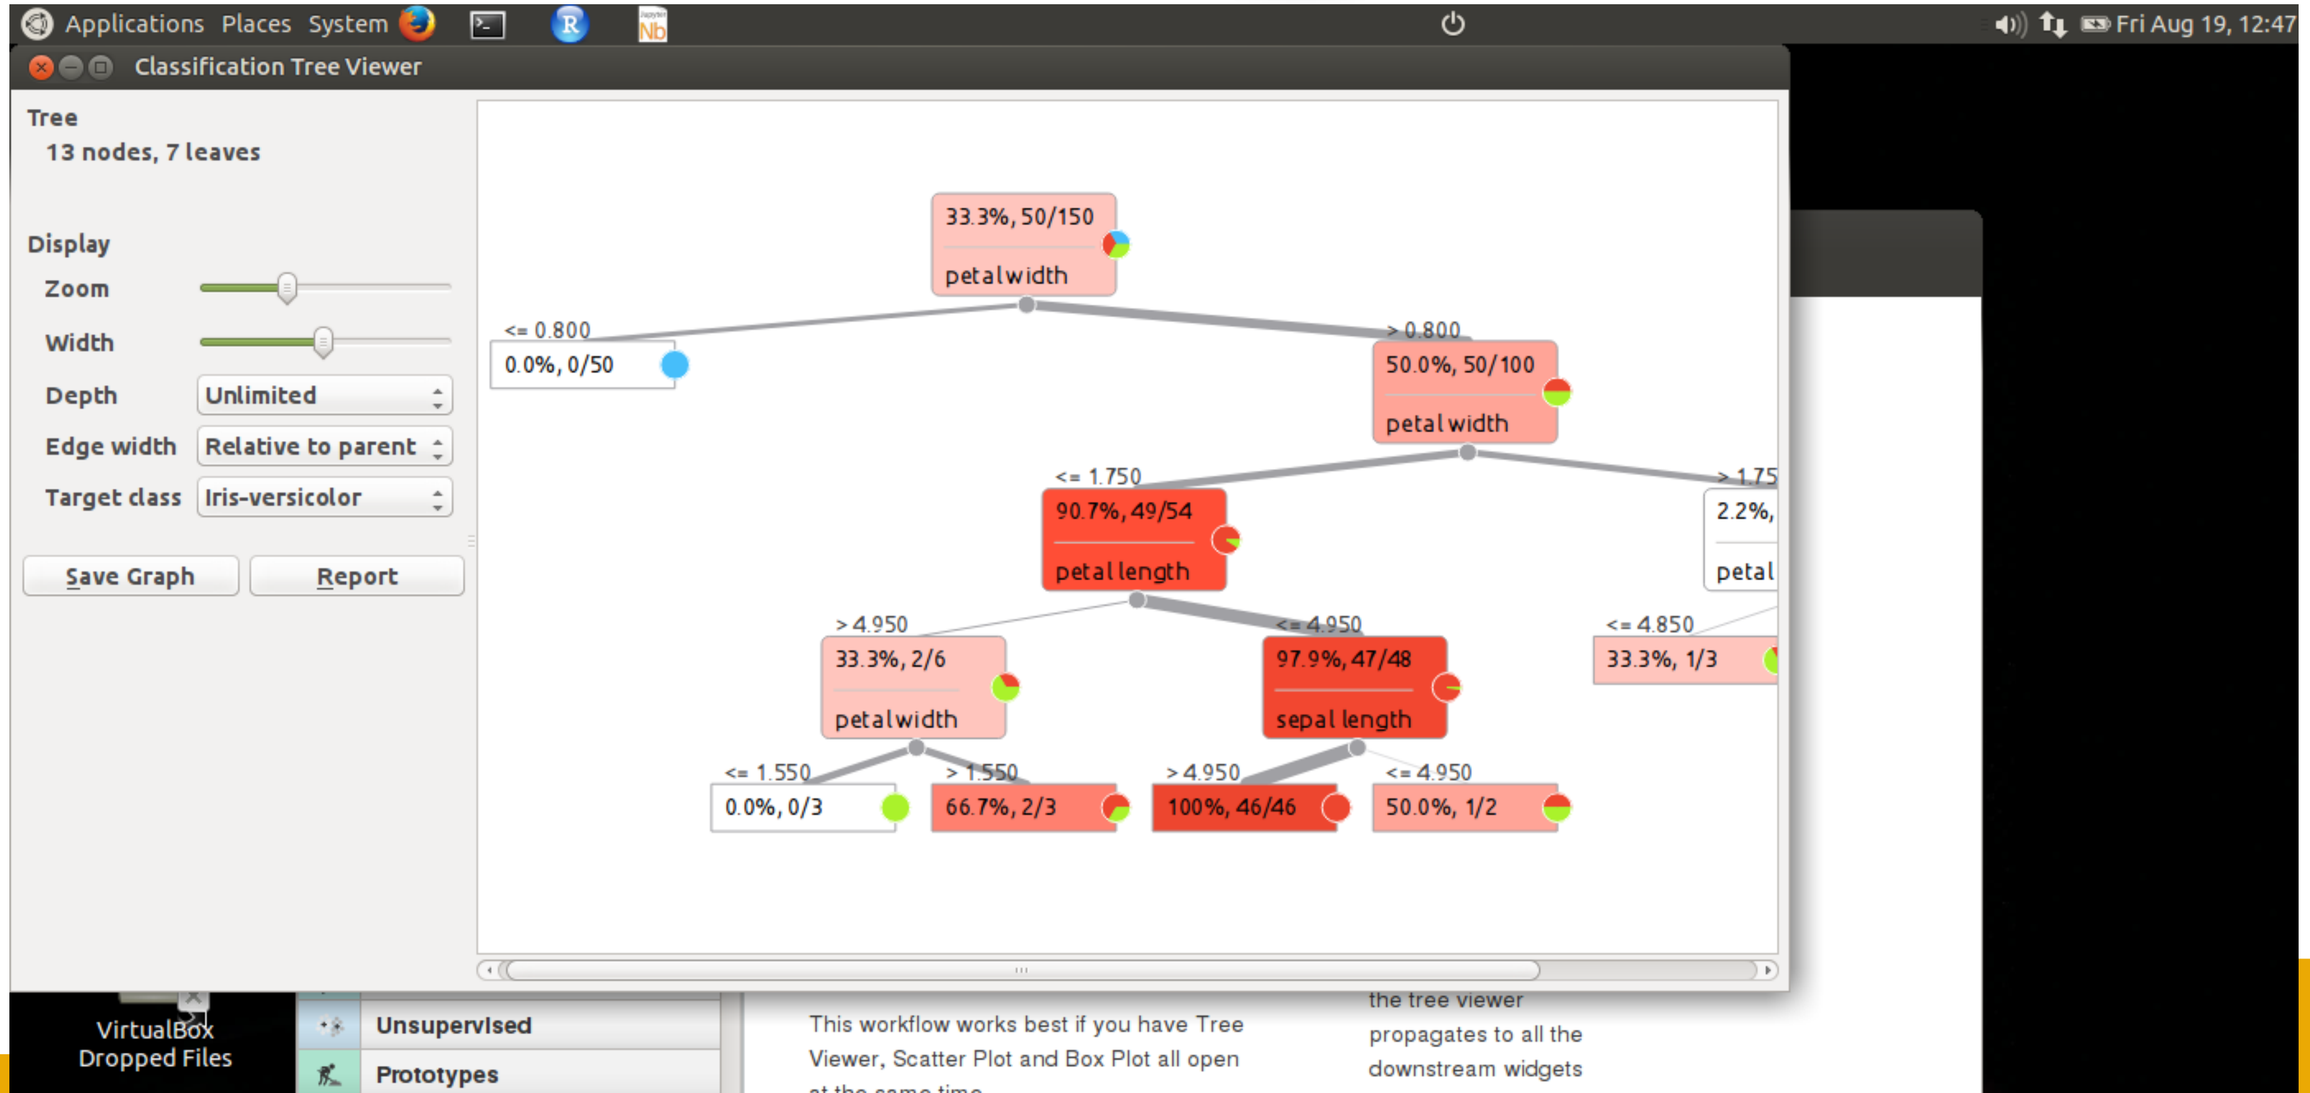

Questions?

Using Jupyter Notebook

Using Jupyter Notebook

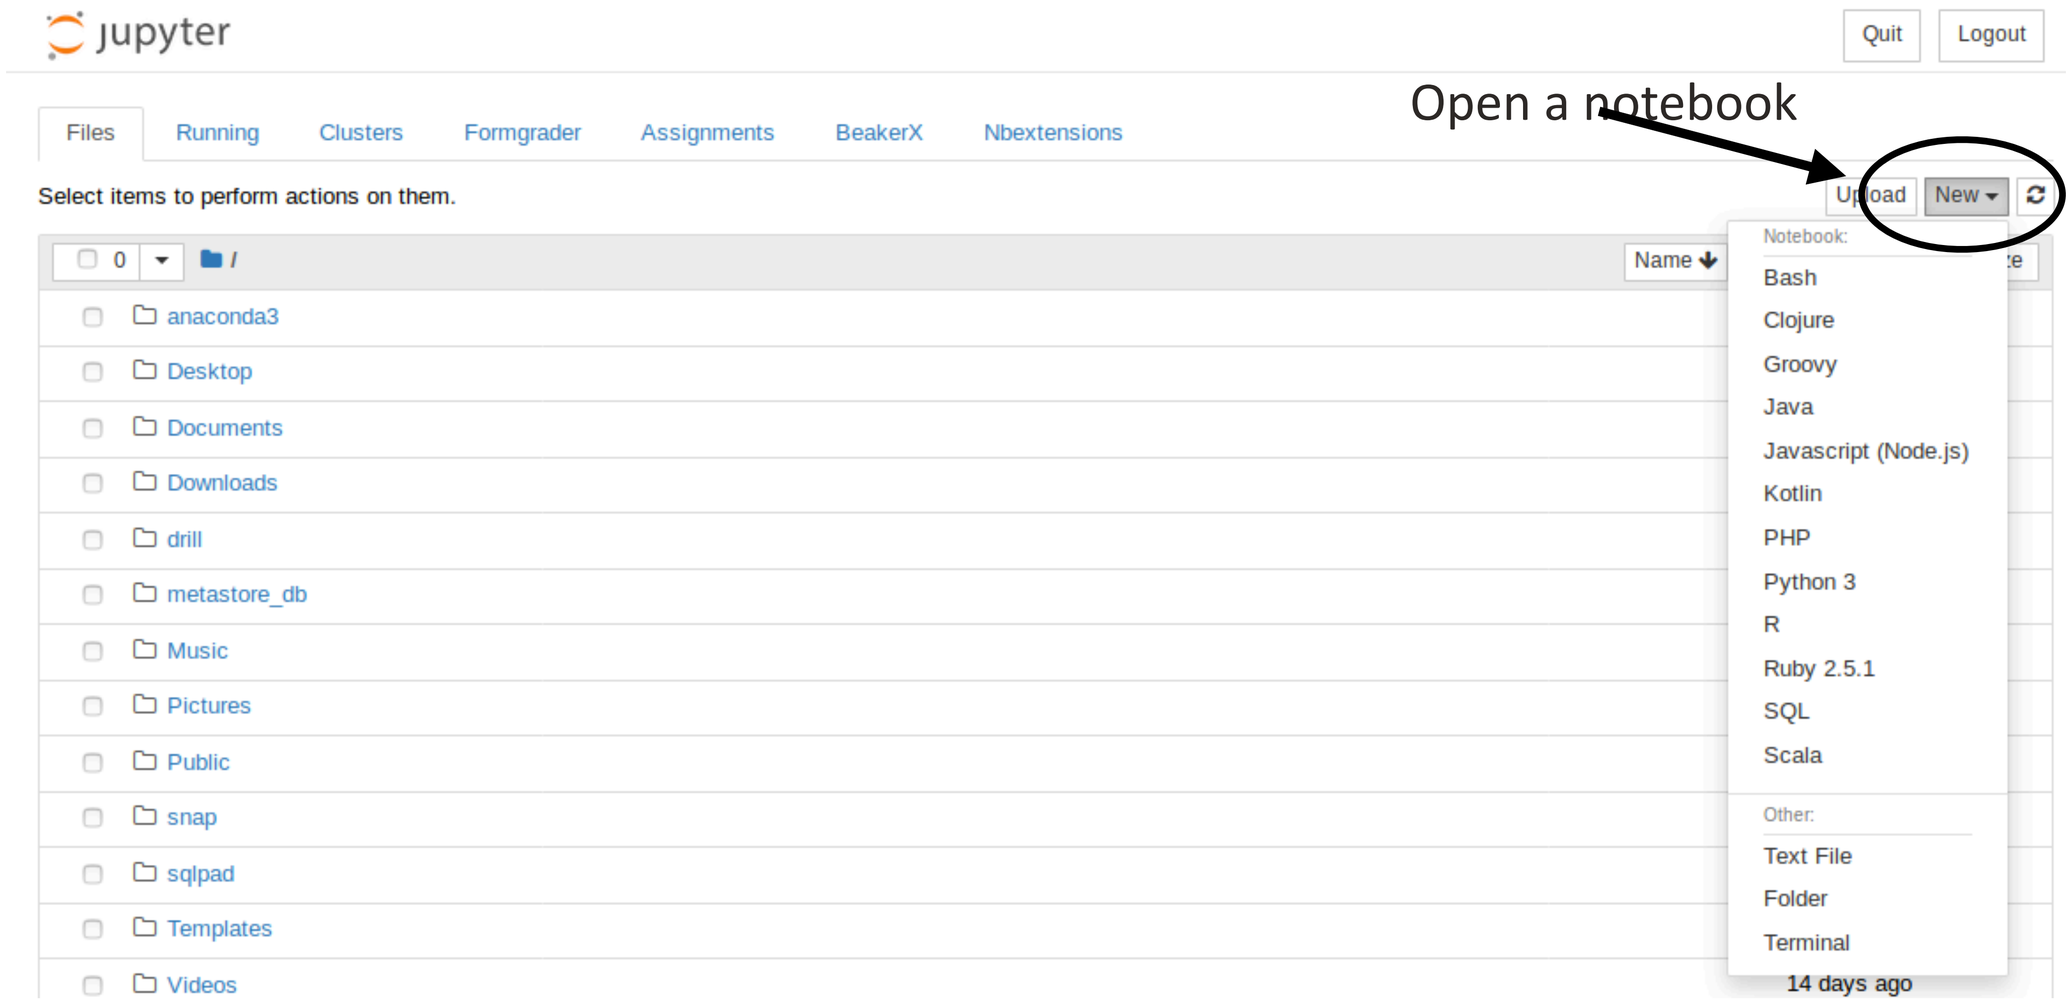

Using Jupyter Notebook

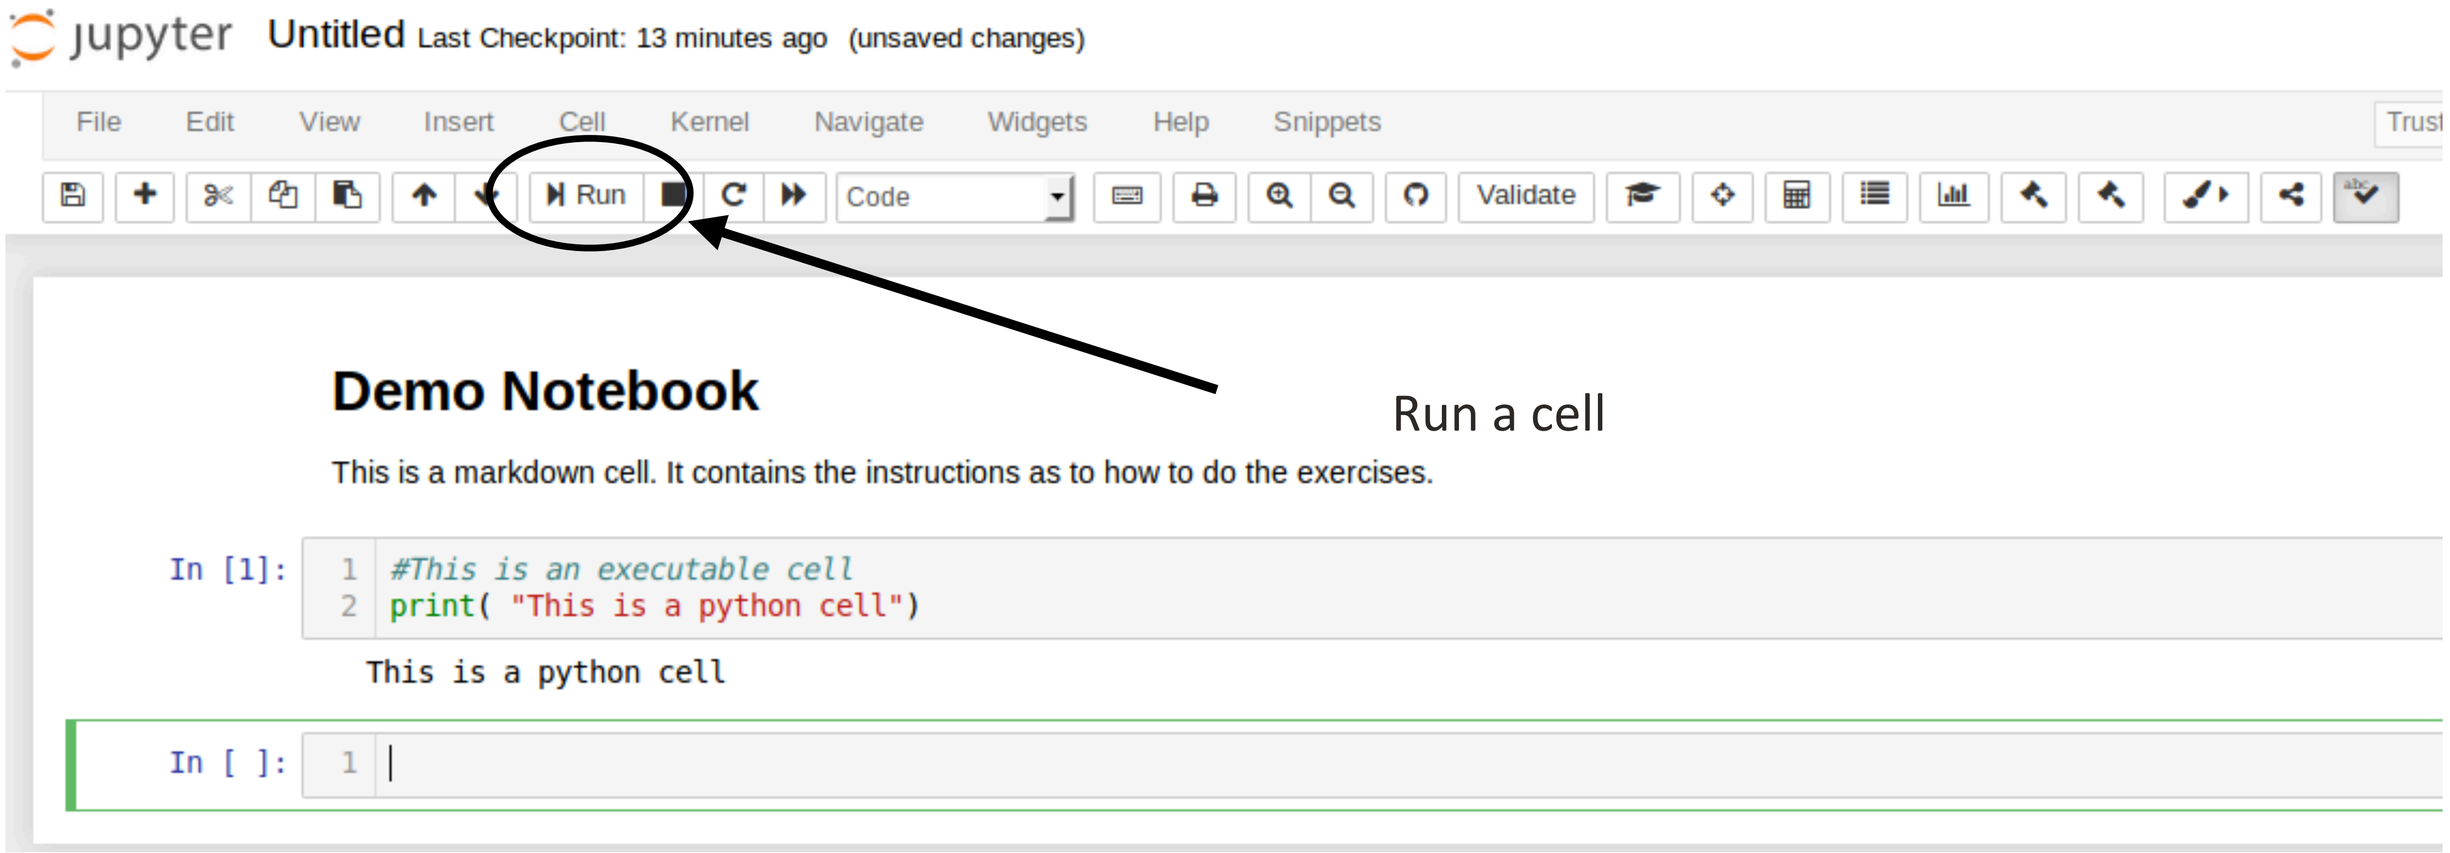

Using Jupyter Notebook

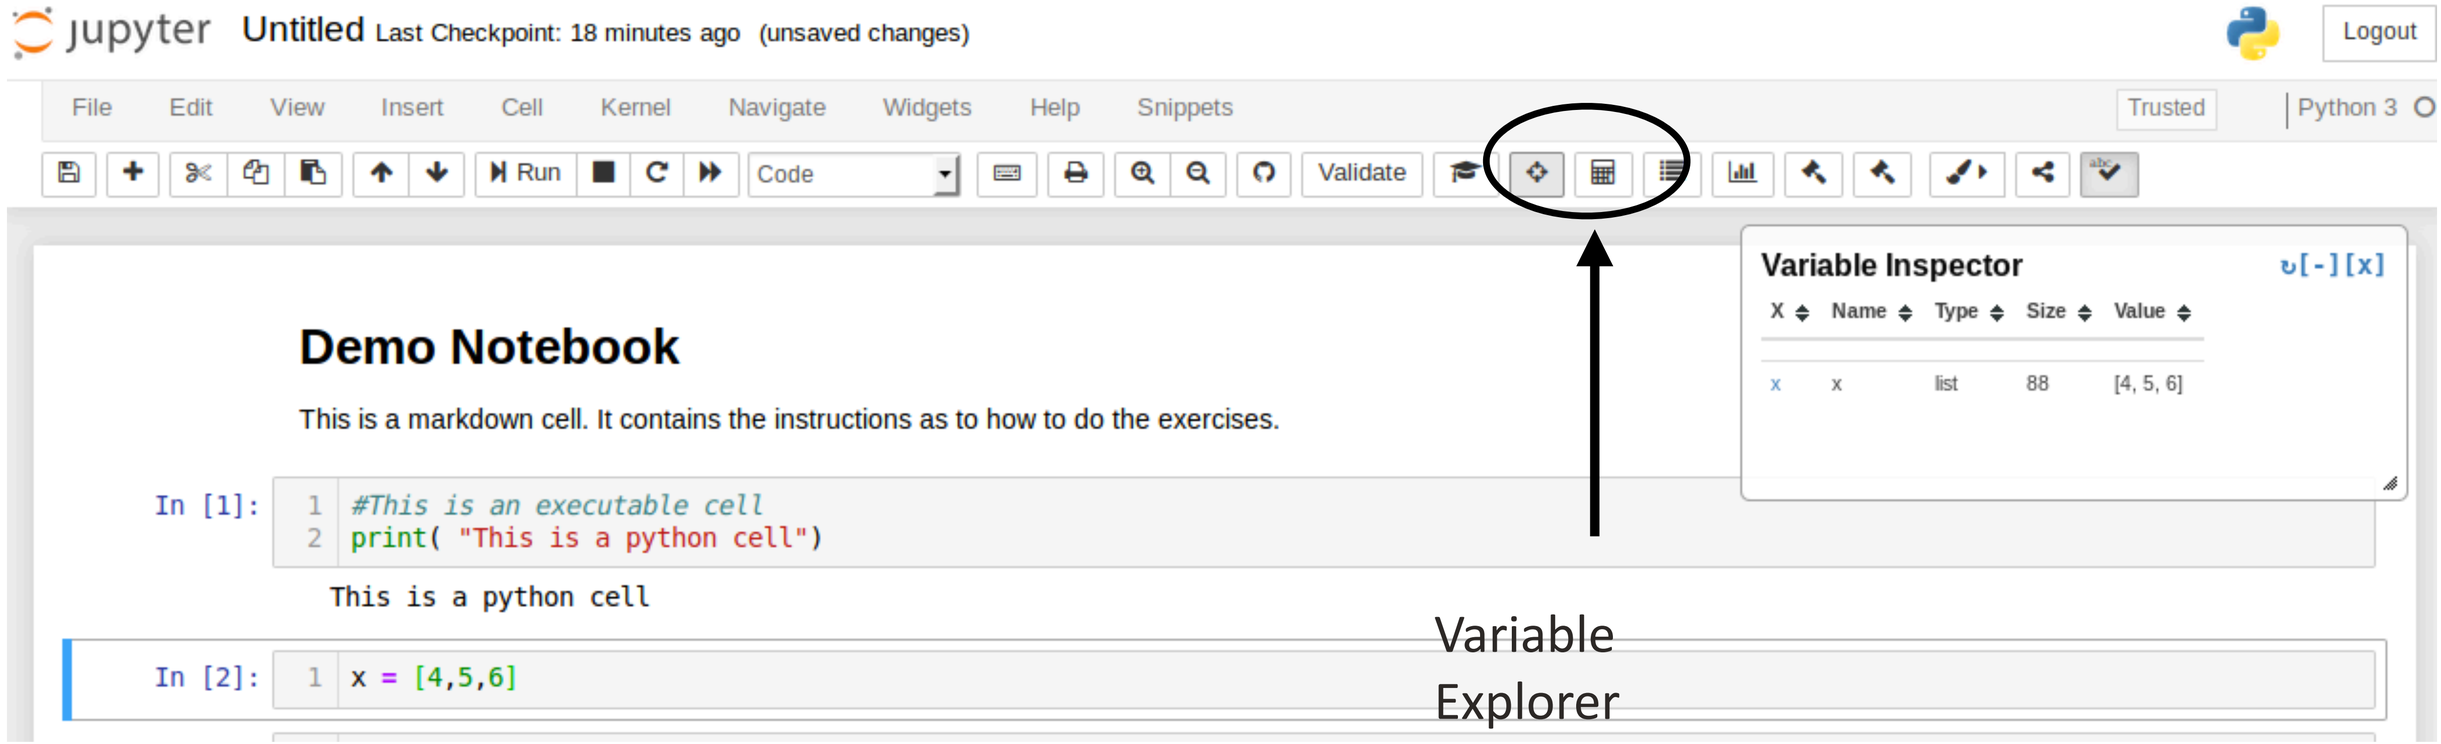

Questions?